

Package ‘pomp’

January 5, 2012

Type Package

Title Statistical inference for partially observed Markov processes

Version 0.40-1

Date 2011-12-24

Author Aaron A. King, Edward L. Ionides, Carles Breto, Steve Ellner, Bruce Kendall, Helen Wear-
ing, Matthew J. Ferrari, Michael Lavine, Daniel C. Reuman

Maintainer Aaron A. King <kingaa@umich.edu>

URL <http://pomp.r-forge.r-project.org>

Description Inference methods for partially-observed Markov processes

Depends R(>= 2.10.1), stats, methods, graphics, mvtnorm, subplex, deSolve

License GPL(>= 2)

LazyLoad true

LazyData false

Collate aaa.R version.R eulermultinom.R plugins.R
slice-design.R profile-design.R sobol.R bsplines.R sannbox.R
pomp-fun.R pomp.R pomp-methods.R rmeasure-pomp.R rprocess-pomp.R init-state-pomp.R
dmeasure-pomp.R dprocess-pomp.R skeleton-pomp.R simulate-pomp.R trajectory-pomp.R plot-
pomp.R pfilter.R pfilter-methods.R traj-match.R bsmc.R
mif-class.R particles-mif.R mif.R mif-methods.R compare-mif.R
pmcmc.R pmcmc-methods.R compare-pmcmc.R nlf-funcs.R nlf-guts.R nlf-objfun.R nlf.R
probe.R probe-match.R basic-probes.R spect.R spect-match.R

R topics documented:

pomp-package	2
B-splines	3
basic.probes	5

blowflies	7
bsmc	7
dacca	9
eulermultinom	10
gompertz	12
LondonYorke	12
mif	13
mif-methods	16
nlf	18
ou2	20
pfilter	21
pfilter-methods	23
plugins	24
pmcmc	27
pmcmc-methods	29
pomp	30
pomp-methods	37
probe	40
probed.pomp-methods	43
profileDesign	45
ricker	46
rw2	46
simulate-pomp	47
sir	48
sliceDesign	49
sobol	50
spect	51
traj.match	54
trajectory	56
verhulst	57
Index	58

pomp-package

Partially-observed Markov processes

Description

The **pomp** package provides facilities for inference on time series data using partially-observed Markov processes (AKA state-space models or nonlinear stochastic dynamical systems). One can use **pomp** to fit nonlinear, non-Gaussian dynamic models to time-series data. The first step in using **pomp** is to encode one's model and data in an object of class `pomp`. One does this via a call to `pomp`, which involves specifying the process and measurement components of the model in one or more of a variety of ways. Details on this are given in the documentation for the `pomp` function and examples are given in the 'intro_to_pomp' vignette.

Currently, **pomp** provides algorithms for (i) simulation of stochastic dynamical systems (see `simulate`), (ii) particle filtering (AKA sequential Monte Carlo or sequential importance sampling), see `pfilter`),

(iii) the iterated filtering method of Ionides et al. (2006), see `mif`), (iv) the nonlinear forecasting algorithm of Kendall et al. (2005), see `nlf`), (v) the particle MCMC approach of Andrieu et al. (2010), see `pmcmc`, (vi) basic trajectory matching, see `traj.match`, (vi) the probe-matching method of Wood (2010) and Kendall et al. (1999), see `probe.match`, (vii) a spectral probe-matching method (Reuman et al., 2006), see `spect.match`. See the package website <http://pomp.r-forge.r-project.org> for these references. The package also provides various tools for plotting and extracting information on models and data as well as an API for algorithm development. Future support for additional algorithms is envisioned, and implementations of the Bayesian sequential Monte Carlo approach of Liu & West. Much of the work in **pomp** has been done under the auspices of a working group of the National Center for Ecological Analysis and Synthesis (NCEAS), "Inference for Mechanistic Models".

The package is provided under the GNU Public License (GPL). Contributions are welcome, as are comments, suggestions for improvements, and bug reports. See the package website <http://pomp.r-forge.r-project.org> for more information, access to the package mailing list, links to the authors' websites, and references to the literature.

Classes

pomp makes extensive use of S4 classes. The basic class, `pomp`, encodes a partially-observed Markov process together with a uni- or multi-variate data set and (possibly) parameters.

Vignettes

The vignette 'Introduction to pomp' illustrates the facilities of the package using familiar stochastic processes. Run `vignette("intro_to_pomp")` or look at the HTML documentation to view the vignette. Methods for accelerating your codes are discussed in the 'Advanced topics in pomp' vignette; run `vignette("advanced_topics_in_pomp")` to view it.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

`pomp`, `pfilter`, `simulate`, `trajectory`, `mif`, `nlf`, `probe.match`, `traj.match`, `bsmc`, `pmcmc`

B-splines

B-spline bases

Description

These functions generate B-spline basis functions. `bspline.basis` gives a basis of spline functions. `periodic.bspline.basis` gives a basis of periodic spline functions.

Usage

```
bspline.basis(x, nbasis, degree = 3, names = NULL)
periodic.bspline.basis(x, nbasis, degree = 3, period = 1, names = NULL)
```

Arguments

<code>x</code>	Vector at which the spline functions are to be evaluated.
<code>nbasis</code>	The number of basis functions to return.
<code>degree</code>	Degree of requested B-splines.
<code>period</code>	The period of the requested periodic B-splines.
<code>names</code>	optional; the names to be given to the basis functions. These will be the column-names of the matrix returned. If the names are specified as a format string (e.g., "basis%d"), <code>sprintf</code> will be used to generate the names from the column number. If a single non-format string is specified, the names will be generated by <code>paste</code> -ing name to the column number. One can also specify each column name explicitly by giving a length-nbasis string vector. By default, no column-names are given.

Details

Direct access to the underlying C routines is available. See the header file “pomp.h” for details.

Value

<code>bspline.basis</code>	Returns a matrix with <code>length(x)</code> rows and <code>nbasis</code> columns. Each column contains the values one of the spline basis functions.
<code>periodic.bspline.basis</code>	Returns a matrix with <code>length(x)</code> rows and <code>nbasis</code> columns. The basis functions returned are periodic with period <code>period</code> .

Author(s)

Aaron A. King <kingaa at umich dot edu>

Examples

```
x <- seq(0,2,by=0.01)
y <- bspline.basis(x,degree=3,nbasis=9,names="basis")
matplot(x,y,type='l',ylim=c(0,1.1))
lines(x,apply(y,1,sum),lwd=2)

x <- seq(-1,2,by=0.01)
y <- periodic.bspline.basis(x,nbasis=5,names="spline%d")
matplot(x,y,type='l')
```

Description

Several simple and configurable probes are provided in the package. These can be used directly and as examples for building custom probes.

Usage

```
probe.mean(var, trim = 0, transform = identity, na.rm = TRUE)
probe.median(var, na.rm = TRUE)
probe.var(var, transform = identity, na.rm = TRUE)
probe.sd(var, transform = identity, na.rm = TRUE)
probe.marginal(var, ref, order = 3, diff = 1, transform = identity)
probe.nlar(var, lags, powers, transform = identity)
probe.acf(var, lags, type = c("covariance", "correlation"),
  transform = identity)
probe.ccf(vars, lags, type = c("covariance", "correlation"),
  transform = identity)
probe.period(var, kernel.width, transform = identity)
probe.quantile(var, prob, transform = identity)
```

Arguments

<code>var, vars</code>	character; the name(s) of the observed variable(s).
<code>trim</code>	the fraction of observations to be trimmed (see mean).
<code>transform</code>	transformation to be applied to the data before the probe is computed.
<code>na.rm</code>	if TRUE, remove all NA observations prior to computing the probe.
<code>kernel.width</code>	width of modified Daniell smoothing kernel to be used in power-spectrum computation: see kernel .
<code>prob</code>	a single probability; the quantile to compute: see quantile .
<code>lags</code>	In <code>probe.ccf</code> , a vector of lags between time series. Positive lags correspond to x advanced relative to y ; negative lags, to the reverse. In <code>probe.nlar</code> , a vector of lags present in the nonlinear autoregressive model that will be fit to the actual and simulated data. See Details, below, for a precise description.
<code>powers</code>	the powers of each term (corresponding to lags) in the the nonlinear autoregressive model that will be fit to the actual and simulated data. See Details, below, for a precise description.
<code>type</code>	Compute autocorrelation or autocovariance?
<code>ref</code>	empirical reference distribution. Simulated data will be regressed against the values of <code>ref</code> , sorted and, optionally, differenced. The resulting regression coefficients capture information about the shape of the marginal distribution. A good choice for <code>ref</code> is the data itself.

order	order of polynomial regression.
diff	order of differencing to perform.
...	Additional arguments to be passed through to the probe computation.

Details

Each of these functions is relatively simple. See the source code for a complete understanding of what each does.

`probe.mean`, `probe.median`, `probe.var`, `probe.sd` return functions that compute the mean, median, variance, and standard deviation of variable `var`, respectively.

`probe.period` returns a function that estimates the period of the Fourier component of the `var` series with largest power.

`probe.marginal` returns a function that regresses the marginal distribution of variable `var` against the reference distribution `ref`. If `diff>0`, the data and the reference distribution are first differenced `diff` times and centered. Polynomial regression of order `order` is used. This probe returns order regression coefficients (the intercept is zero).

`probe.nlar` returns a function that fit a nonlinear (polynomial) autoregressive model to the univariate series (variable `var`). Specifically, a model of the form $y_t = \sum \beta_k y_{t-\tau_k}^{p_k} + \epsilon_t$ will be fit, where τ_k are the lags and p_k are the powers. The data are first centered. This function returns the regression coefficients, β_k .

`probe.acf` returns a function that, if `type=="covariance"`, computes the autocovariance of variable `var` at lags `lags`; if `type=="correlation"`, computes the autocorrelation of variable `var` at lags `lags`.

`probe.ccf` returns a function that, if `type=="covariance"`, computes the cross covariance of the two variables named in `vars` at lags `lags`; if `type=="correlation"`, computes the cross correlation.

`probe.quantile` returns a function that estimates the `prob`-th quantile of variable `var`.

Value

A call to any one of these functions returns a probe function, suitable for use in `probe` or `probe.match`. That is, the function returned by each of these takes a data array (such as comes from a call to `obs`) as input and returns a single numerical value.

Author(s)

Daniel C. Reuman (d.reuman at imperial dot ac dot uk)

Aaron A. King (kingaa at umich dot edu)

References

B. E. Kendall, C. J. Briggs, W. M. Murdoch, P. Turchin, S. P. Ellner, E. McCauley, R. M. Nisbet, S. N. Wood Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches, *Ecology*, 80:1789–1805, 1999.

S. N. Wood Statistical inference for noisy nonlinear ecological dynamic systems, *Nature*, 466: 1102–1104, 2010.

See Also

[pomp-class](#), [pomp-methods](#)

blowflies

Model for Nicholson's blowflies.

Description

blowfly1 and blowfly2 are pomp objects encoding stochastic delay-difference models.

Usage

```
data(blowflies)
```

See Also

[pomp-class](#) and the vignettes

Examples

```
data(blowflies)
plot(blowflies1)
plot(blowflies2)
```

bsmc

Liu and West Bayesian Particle Filter

Description

Generates draws from the posterior distribution for the parameters using the Liu and West algorithm. bsmc gives draws from the posterior.

Usage

```
## S4 method for signature 'pomp'
bsmc(object, params, Np, est, smooth = 0.1,
      ntries = 1, tol = 1e-17, lower = -Inf, upper = Inf, seed = NULL,
      verbose = getOption("verbose"), max.fail = 0, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> or inheriting class <code>pomp</code> .
<code>params</code>	A $n_{\text{pars}} \times N_p$ matrix (with rownames) containing the parameters corresponding to the initial state values in <code>xstart</code> . N_p is the number of particles, i.e., each row should contain N_p draws from the prior distribution for that parameter. It is permissible to supply <code>params</code> as a named numeric vector, i.e., without a <code>dim</code> attribute. In this case, all particles will inherit the same parameter values, which is equivalent to a degenerate prior.
<code>Np</code>	If <code>params</code> is specified as a named vector, <code>Np</code> specifies the number of particles to use. If <code>params</code> is specified as a matrix, <code>Np</code> should not be specified; it is taken to be the number of columns of <code>params</code> .
<code>est</code>	Names of the rows of <code>params</code> that are to be estimated. No updates will be made to the other parameters. If <code>est</code> is not specified, all parameters for which there is variation in <code>params</code> will be estimated.
<code>smooth</code>	Kernel density smoothing parameters. The compensating shrinkage factor will be $\sqrt{1 - \text{smooth}^2}$. Thus, <code>smooth=0</code> means that no noise will be added to parameters. Generally, the value of <code>smooth</code> should be chosen close to 0 (i.e., $\text{shrink} \sim 0.1$).
<code>ntries</code>	Number of draws from <code>rprocess</code> per particle used to estimate the expected value of the state process at time $t+1$ given the state and parameters at time t .
<code>tol</code>	Particles with log likelihood below <code>tol</code> are considered to be “lost”. A filtering failure occurs when, at some time point, all particles are lost. When all particles are lost, the conditional log likelihood at that time point is set to be $\log(\text{tol})$.
<code>lower, upper</code>	optional; lower and upper bounds on the priors. This is useful in case there are box constraints satisfied by the priors. The posterior is guaranteed to lie within these bounds.
<code>seed</code>	optional; an object specifying if and how the random number generator should be initialized (“seeded”). If <code>seed</code> is an integer, it is passed to <code>set.seed</code> prior to any simulation and is returned as the “seed” element of the return list. By default, the state of the random number generator is not changed and the value of <code>.Random.seed</code> on the call is stored in the “seed” element of the return list.
<code>verbose</code>	logical; if TRUE, print diagnostic messages.
<code>max.fail</code>	The maximum number of filtering failures allowed. If the number of filtering failures exceeds this number, execution will terminate with an error.
<code>...</code>	currently ignored.

Value

A list with the following elements:

<code>post</code>	A matrix containing draws from the approximate posterior distribution.
<code>prior</code>	A matrix containing draws from the prior distribution (identical to <code>params</code> on call).
<code>eff.sample.size</code>	A vector containing the effective number of particles at each time point.

cond.loglik	A vector containing the conditional log likelihoods at each time point.
smooth	The smoothing parameter used (see above).
seed	The state of the random number generator at the time bsmc was called. If the argument seed was specified, this is a copy; if not, this is the internal state of the random number generator at the time of call.
nfail	The number of filtering failures encountered.
loglik	The estimated log-likelihood.
weights	The resampling weights for each particle.

Author(s)

Michael Lavine (lavine at math dot umass dot edu), Matthew Ferrari (mferrari at psu dot edu), Aaron A. King

See Also

[pomp-class](#)

Examples

```
## See the vignettes for examples.
```

dacca	<i>Model of cholera transmission for historic Bengal.</i>
-------	---

Description

dacca is a pomp object containing census and cholera mortality data from the Dacca district of the former British province of Bengal over the years 1891 to 1940 together with a stochastic differential equation transmission model. The model is that of King et al. (2008). The parameters are the MLE for the SIRS model with seasonal reservoir.

Data are provided courtesy of Dr. Menno J. Bouma, London School of Tropical Medicine and Hygiene.

Usage

```
data(dacca)
```

Details

dacca is a pomp object containing the model, data, and MLE parameters. Parameters that naturally range over the positive reals are log-transformed; parameters that range over the unit interval are logit-transformed; parameters that are naturally unbounded or take integer values are not transformed.

References

King, A. A., Ionides, E. L., Pascual, M., and Bouma, M. J. Inapparent infections and cholera dynamics. *Nature* 454:877-880 (2008)

See Also

[euler.sir](#), [pomp](#)

Examples

```
data(dacca)
plot(dacca)
#MLEs on the natural scale
coef(dacca,transform=TRUE)
plot(simulate(dacca))
# now change 'eps' and simulate again
coef(dacca,"eps",transform=TRUE) <- 1
plot(simulate(dacca))
```

eulermultinom

Euler-multinomial death process

Description

Density and random-deviate generation for the Euler-multinomial death process with parameters size, rate, and dt.

Usage

```
reulermultinom(n = 1, size, rate, dt)
deulermultinom(x, size, rate, dt, log = FALSE)
```

Arguments

n	integer; number of random variates to generate.
size	scalar integer; number of individuals at risk.
rate	numeric vector of hazard rates.
dt	numeric scalar; duration of Euler step.
x	matrix or vector containing number of individuals that have succumbed to each death process.
log	logical; if TRUE, return logarithm(s) of probabilities.

Details

If N individuals face constant hazards of death in k ways at rates r_1, r_2, \dots, r_k , then in an interval of duration Δt , the number of individuals remaining alive and dying in each way is multinomially distributed:

$$(N - \sum_{i=1}^k \Delta n_i, \Delta n_1, \dots, \Delta n_k) \sim \text{multinomial}(N; p_0, p_1, \dots, p_k),$$

where Δn_i is the number of individuals dying in way i over the interval, the probability of remaining alive is $p_0 = \exp(-\sum_i r_i \Delta t)$, and the probability of dying in way j is

$$p_j = \frac{r_j}{\sum_i r_i} (1 - \exp(-\sum_i r_i \Delta t)).$$

In this case, we can say that

$$(\Delta n_1, \dots, \Delta n_k) \sim \text{eulermultinom}(N, r, \Delta t),$$

where $r = (r_1, \dots, r_k)$. Draw m random samples from this distribution by doing

`reulermultinom(n=m, size=N, rate=r, dt=dt)`,

where r is the vector of rates. Evaluate the probability that $x = (x_1, \dots, x_k)$ are the numbers of individuals who have died in each of the k ways over the interval $\Delta t = dt$, by doing

`deulermultinom(x=x, size=N, rate=r, dt=dt)`.

Direct access to the underlying C routines is available: see the header file “pomp.h”, included with the package.

Value

`reulermultinom` Returns a `length(rate)` by `n` matrix. Each column is a different random draw. Each row contains the numbers of individuals succumbed to the corresponding process.

`deulermultinom` Returns a vector (of length equal to the number of columns of `x`) containing the probabilities of observing each column of `x` given the specified parameters (`size`, `rate`, `dt`).

Author(s)

Aaron A. King <kingaa at umich dot edu>

Examples

```
print(x <- reulermultinom(5, size=100, rate=c(a=1, b=2, c=3), dt=0.1))
deulermultinom(x, size=100, rate=c(1, 2, 3), dt=0.1)
```

gompertz

Gompertz model with log-normal observations.

Description

gompertz is a pomp object encoding a stochastic Gompertz population model with log-normal measurement error.

Usage

```
data(gompertz)
```

Details

The state process is $X_{t+1} = K^{(1-S)} X_t^S \varepsilon_t$, where $S = e^{-r}$ and the ε_t are i.i.d. lognormal random deviates with variance σ^2 . The observed variables Y_t are distributed as $\text{lognormal}(\log X_t, \tau)$. Parameters include the per-capita growth rate r , the carrying capacity K , the process noise s.d. *sigma*, the measurement error s.d. *tau*, and the initial condition X_0 . The model is parameterized internally by the logarithms of r , K , σ , and τ ; the initial condition is parameterized directly. The pomp object includes parameter transformations to and from this internal parameterization.

See Also

[pomp-class](#) and the introductory vignette `vignette("intro_to_pomp")`.

Examples

```
data(gompertz)
plot(gompertz)
coef(gompertz)
coef(gompertz, transform=TRUE)
```

LondonYorke

Historical childhood disease incidence data

Description

LondonYorke is a data-frame containing the monthly number of reported cases of chickenpox, measles, and mumps from two American cities (Baltimore and New York) in the mid-20th century (1928–1972).

Usage

```
data(LondonYorke)
```

References

W. P. London and J. A. Yorke, Recurrent Outbreaks of Measles, Chickenpox and Mumps: I. Seasonal Variation in Contact Rates, American Journal of Epidemiology, 98:453–468, 1973.

See Also

[pomp-class](#) and the vignettes

Examples

```
data(LondonYorke)

plot(cases~time,data=LondonYorke,subset=disease=="measles",type='n',main="measles",bty='l')
lines(cases~time,data=LondonYorke,subset=disease=="measles"&town=="Baltimore",col="red")
lines(cases~time,data=LondonYorke,subset=disease=="measles"&town=="New York",col="blue")
legend("topright",legend=c("Baltimore","New York"),lty=1,col=c("red","blue"),bty='n')

plot(
  cases~time,
  data=LondonYorke,
  subset=disease=="chickenpox"&town=="New York",
  type='l',col="blue",main="chickenpox, New York",
  bty='l'
)

plot(
  cases~time,
  data=LondonYorke,
  subset=disease=="mumps"&town=="New York",
  type='l',col="blue",main="mumps, New York",
  bty='l'
)
```

mif

The MIF algorithm

Description

The MIF algorithm for estimating the parameters of a partially-observed Markov process.

Usage

```
mif(object, ...)
## S4 method for signature 'pomp'
mif(object, Nmif = 1, start, pars, ivps = character(0),
     particles, rw.sd, Np, ic.lag, var.factor, cooling.factor,
     weighted = TRUE, tol = 1e-17, max.fail = 0,
     verbose = getOption("verbose"), ...)
## S4 method for signature 'pfilterd.pomp'
```

```

mif(object, Nmif = 1, start, pars, ivps = character(0),
    particles, rw.sd, Np, ic.lag, var.factor, cooling.factor,
    weighted = TRUE, tol, max.fail = 0,
    verbose = getOption("verbose"), ...)
## S4 method for signature 'mif'
mif(object, Nmif, start, pars, ivps,
    particles, rw.sd, Np, ic.lag, var.factor, cooling.factor,
    weighted = TRUE, tol, max.fail = 0,
    verbose = getOption("verbose"), ...)
## S4 method for signature 'mif'
continue(object, Nmif = 1, start, pars, ivps,
    particles, rw.sd, Np, ic.lag, var.factor, cooling.factor,
    weighted = TRUE, tol, max.fail = 0,
    verbose = getOption("verbose"), ...)

```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>Nmif</code>	The number of MIF iterations to perform.
<code>start</code>	named numerical vector; the starting guess of the parameters.
<code>pars</code>	optional character vector naming the ordinary parameters to be estimated. Every parameter named in <code>pars</code> must have a positive random-walk standard deviation specified in <code>rw.sd</code> . Leaving <code>pars</code> unspecified is equivalent to setting it equal to the names of all parameters with a positive value of <code>rw.sd</code> that are not <code>ivps</code> .
<code>ivps</code>	optional character vector naming the initial-value parameters (IVPs) to be estimated. Every parameter named in <code>ivps</code> must have a positive random-walk standard deviation specified in <code>rw.sd</code> . If <code>pars</code> is empty, i.e., only IVPs are to be estimated, see below ““Using MIF to estimate initial-value parameters only””.
<code>particles</code>	Function of prototype <code>particles(Np, center, sd, ...)</code> which sets up the starting particle matrix by drawing a sample of size <code>Np</code> from the starting particle distribution centered at <code>center</code> and of width <code>sd</code> . If <code>particles</code> is not supplied by the user, the default behavior is to draw the particles from a multivariate normal distribution with mean <code>center</code> and standard deviation <code>sd</code> .
<code>rw.sd</code>	numeric vector with names; the intensity of the random walk to be applied to parameters. The random walk is only applied to parameters named in <code>pars</code> (i.e., not to those named in <code>ivps</code>). The algorithm requires that the random walk be nontrivial, so each element in <code>rw.sd[pars]</code> must be positive. <code>rw.sd</code> is also used to scale the initial-value parameters (via the <code>particles</code> function). Therefore, each element of <code>rw.sd[ivps]</code> must be positive. The following must be satisfied: <code>names(rw.sd)</code> must be a subset of <code>names(start)</code> , <code>rw.sd</code> must be non-negative (zeros are simply ignored), the name of every positive element of <code>rw.sd</code> must be in either <code>pars</code> or <code>ivps</code> .
<code>Np</code>	the number of particles to use in filtering. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across

timestep, one may specify `Np` either as a vector of positive integers (of length `length(time(object, t0=TRUE))`) or as a function taking a positive integer argument. In the latter case, `Np(k)` must be a single positive integer, representing the number of particles to be used at the k -th timestep: `Np(0)` is the number of particles to use going from `timezero(object)` to `time(object)[1]`, `Np(1)`, from `timezero(object)` to `time(object)[1]`, and so on, while when `T=length(time(object, t0=TRUE))`, `Np(T)` is the number of particles to sample at the end of the time-series.

<code>ic.lag</code>	a positive integer; the timepoint for fixed-lag smoothing of initial-value parameters. The <code>mif</code> update for initial-value parameters consists of replacing them by their filtering mean at time <code>times[ic.lag]</code> , where <code>times=time(object)</code> . It makes no sense to set <code>ic.lag>length(times)</code> ; if it is so set, <code>ic.lag</code> is set to <code>length(times)</code> with a warning.
<code>var.factor</code>	a positive number; the scaling coefficient relating the width of the starting particle distribution to <code>rw.sd</code> . In particular, the width of the distribution of particles at the start of the first MIF iteration will be <code>random.walk.sd*var.factor</code> .
<code>cooling.factor</code>	a positive number not greater than 1; the exponential cooling factor, α .
<code>weighted</code>	logical; if <code>TRUE</code> , the MIF update (a weighted average) is used. If <code>FALSE</code> , the MIF update is not used; instead, an unweighed average of the filtering means is used for the update.
<code>tol</code>	numeric scalar; particles with log likelihood below <code>tol</code> are considered to be “lost”. A filtering failure occurs when, at some time point, all particles are lost.
<code>max.fail</code>	integer; maximum number of filtering failures permitted. If the number of failures exceeds this number, execution will terminate with an error.
<code>verbose</code>	logical; if <code>TRUE</code> , print progress reports.
<code>...</code>	additional arguments. Currently, these are ignored.

Re-running MIF Iterations

To re-run a sequence of MIF iterations, one can use the `mif` method on a `mif` object. By default, the same parameters used for the original MIF run are re-used (except for `weighted`, `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing MIF Iterations

One can resume a series of MIF iterations from where one left off using the `continue` method. A call to `mif` to perform `Nmif=m` iterations followed by a call to `continue` to perform `Nmif=n` iterations will produce precisely the same effect as a single call to `mif` to perform `Nmif=m+n` iterations. By default, all the algorithmic parameters are the same as used in the original call to `mif`. Additional arguments will override the defaults.

Using MIF to estimate initial-value parameters only

One can use MIF’s fixed-lag smoothing to estimate only initial value parameters (IVPs). In this case, `pars` is left empty and the IVPs to be estimated are named in `ivps`. If `theta` is the current parameter vector, then at each MIF iteration, `Np` particles are drawn from a distribution centered

at `theta` and with width proportional to `var.factor*rw.sd`, a particle filtering operation is performed, and `theta` is replaced by the filtering mean at `time(object)[ic.lag]`. Note the implication that, when `mif` is used in this way on a time series any longer than `ic.lag`, unnecessary work is done. If the time series in `object` is longer than `ic.lag`, consider replacing `object` with `window(object, end=ic.lag)`.

Details

If `particles` is not specified, the default behavior is to draw the particles from a multivariate normal distribution. **It is the user's responsibility to ensure that, if the optional `particles` argument is given, that the `particles` function satisfies the following conditions:**

`particles` has at least the following arguments: `Np`, `center`, `sd`, and `.` `Np` may be assumed to be a positive integer; `center` and `sd` will be named vectors of the same length. Additional arguments may be specified; these will be filled with the elements of the `userdata` slot of the underlying `pomp` object (see [pomp-class](#)).

`particles` returns a `length(center) x Np` matrix with rownames matching the names of `center` and `sd`. Each column represents a distinct particle.

The center of the particle distribution returned by `particles` should be `center`. The width of the particle distribution should vary monotonically with `sd`. In particular, when `sd=0`, the `particles` should return matrices with `Np` identical columns, each given by the parameters specified in `center`.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

E. L. Ionides, C. Bret\'o, & A. A. King, Inference for nonlinear dynamical systems, *Proc. Natl. Acad. Sci. U.S.A.*, 103:18438–18443, 2006.

A. A. King, E. L. Ionides, M. Pascual, and M. J. Bouma, Inapparent infections and cholera dynamics, *Nature*, 454:877–880, 2008.

See Also

[mif-methods](#), [pomp](#), [pomp-class](#), [pfilter](#). See the “`intro_to_pomp`” vignette for examples.

Description

Methods of the "mif" class.

Usage

```
## S4 method for signature 'mif'
logLik(object, ...)
## S4 method for signature 'mif'
conv.rec(object, pars, transform = FALSE, ...)
## S4 method for signature 'mif'
plot(x, y = NULL, ...)
compare.mif(z)
```

Arguments

<code>object</code>	The mif object.
<code>pars</code>	Names of parameters.
<code>x</code>	The mif object.
<code>y</code>	Ignored.
<code>z</code>	A mif object or list of mif objects.
<code>transform</code>	optional logical; should the parameter transformations be applied? See coef for details.
<code>...</code>	Further arguments (either ignored or passed to underlying functions).

Methods

conv.rec `conv.rec(object, pars = NULL)` returns the columns of the convergence-record matrix corresponding to the names in `pars`. By default, all rows are returned.

logLik Returns the value in the `loglik` slot.

mif Re-runs the MIF iterations. See the documentation for [mif](#).

compare.mif Given a mif object or a list of mif objects, `compare.mif` produces a set of diagnostic plots.

plot Plots a series of diagnostic plots. When `x` is a mif object, `plot(x)` is equivalent to `compare.mif(list(x))`.

predvarplot `predvarplot(object, pars = NULL, mean = FALSE, ...)` produces a plot of the scaled prediction variances for each parameter. This can be used to diagnose a good value of the mif parameters `var.factor` and `ic.lag`. If used in this way, one should run `mif` with `Nmif=1` first. Additional arguments in `...` will be passed to the actual plotting function.

print Prints a summary of the mif object.

show Displays the mif object.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

E. L. Ionides, C. Bret\'o, & A. A. King, Inference for nonlinear dynamical systems, *Proc. Natl. Acad. Sci. U.S.A.*, 103:18438–18443, 2006.

A. A. King, E. L. Ionides, M. Pascual, and M. J. Bouma, Inapparent infections and cholera dynamics, *Nature*, 454:877–880, 2008.

See Also

[mif](#), [pomp](#), [pomp-class](#), [pfilter](#)

nlf

Fit Model to Data Using Nonlinear Forecasting (NLF)

Description

Calls an optimizer to maximize the nonlinear forecasting (NLF) goodness of fit, by simulating data from a model, fitting a nonlinear autoregressive model to the simulated time series (which may be multivariate) and using the fitted model to predict some or all variables in the data time series. NLF is an ‘indirect inference’ method using a quasi-likelihood as the objective function.

Usage

```
nlf(object, start, est, lags, period = NA, tensor = FALSE,
     nconverge=1000, nasymp=1000, seed = 1066,
     transform = function (x) x,
     nrbf = 4, method = "subplex", skip.se = FALSE,
     verbose = FALSE, gr = NULL,
     bootstrap=FALSE, bootsamp = NULL,
     lql.frac = 0.1, se.par.frac = 0.1, eval.only = FALSE, ...)
```

Arguments

object	A pomp object, with the data and model to fit to it.
start	Named numeric vector with guessed parameters.
est	Vector containing the names or indices of parameters to be estimated.
lags	A vector specifying the lags to use when constructing the nonlinear autoregressive prediction model. The first lag is the prediction interval.
period	numeric; period=NA means the model is nonseasonal. period>0 is the period of seasonal forcing in ‘real time’.
tensor	logical; if FALSE, the fitted model is a generalized additive model with time mod period as one of the predictors, i.e., a gam with time-varying intercept. If TRUE, the fitted model is a gam with lagged state variables as predictors and time-periodic coefficients, constructed using tensor products of basis functions of state variables with basis functions of time.
nconverge	Number of convergence timesteps to be discarded from the model simulation.
nasymp	Number of asymptotic timesteps to be recorded from the model simulation.
seed	Integer specifying the random number seed to use. When fitting, it is usually best to always run the simulations with the same sequence of random numbers, which is accomplished by setting seed to an integer. If you want a truly random simulation, set seed=NULL.

transform	optional function. If specified, forecasting is performed using data and model simulations transformed by this function. By default, transform is the identity function. The main purpose of transform is to achieve approximately multivariate normal forecasting errors. If data are univariate, transform should take a scalar and return a scalar. If data are multivariate, transform should assume a vector input and return a vector of the same length.
nrbf	A scalar specifying the number of radial basis functions to be used at each lag.
method	Optimization method. Choices are <code>subplex</code> and any of the methods used by <code>optim</code> .
skip.se	Logical; if TRUE, skip the computation of standard errors.
verbose	Logical; if TRUE, the negative log quasiliikelihood and parameter values are printed at each iteration of the optimizer.
gr	optional; passed to <code>optim</code> if <code>optim</code> is used.
bootstrap	Logical; if TRUE the indices in <code>bootsamp</code> will determine which of the conditional likelihood values be used in computing the quasi-loglikelihood.
bootsamp	Vector of integers; used to have the quasi-loglikelihood evaluated using a bootstrap re-sampling of the data set.
lql.frac	target fractional change in log quasi-likelihood for quadratic standard error estimate
se.par.frac	initial parameter-change fraction for quadratic standard error estimate
eval.only	logical; if TRUE, no optimization is attempted and the quasi-loglikelihood value is evaluated at the <code>start</code> parameters.
...	Arguments that will be passed to <code>optim</code> or <code>subplex</code> in the control list.

Details

This is functionally a wrapper for `nlf.objfun`, which does the statistical heavy lifting and should be consulted for details.

Value

A list corresponding to the output from the optimizer, except that the full parameter vector is returned (not just the ones fitted), the log quasiliikelihood (LQL) (*not* -LQL) is reported, `xstart` is included, and asymptotic Wald standard errors based on M-estimator theory are returned for each fitted parameter.

Author(s)

Stephen P. Ellner <spe2 at cornell dot edu> and Bruce E. Kendall <kendall at bren dot ucsb dot edu>

References

The following papers describe and motivate the NLF approach to model fitting:

Ellner, S. P., Bailey, B. A., Bobashev, G. V., Gallant, A. R., Grenfell, B. T. and Nychka D. W. (1998) Noise and nonlinearity in measles epidemics: combining mechanistic and statistical approaches to population modeling. *American Naturalist* **151**, 425–440.

Kendall, B. E., Briggs, C. J., Murdoch, W. W., Turchin, P., Ellner, S. P., McCauley, E., Nisbet, R. M. and Wood S. N. (1999) Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches. *Ecology* **80**, 1789–1805. Available online at <http://www2.bren.ucsb.edu/~kendall/pubs/1999Ecology.pdf>

Kendall, B. E., Ellner, S. P., McCauley, E., Wood, S. N., Briggs, C. J., Murdoch, W. W. and Turchin, P. (2005) Population cycles in the pine looper moth (*Bupalus piniarius*): dynamical tests of mechanistic hypotheses. *Ecological Monographs* **75**, 259–276. Available online at <http://repositories.cdlib.org/postprints/818/>

ou2

Two-dimensional discrete-time Ornstein-Uhlenbeck process

Description

ou2 is a pomp object encoding a bivariate discrete-time Ornstein-Uhlenbeck process.

Usage

```
data(ou2)
```

Details

If the state process is $X(t) = (x_1(t), x_2(t))$, then

$$X(t+1) = \alpha X(t) + \sigma \epsilon(t),$$

where α and σ are 2x2 matrices, σ is lower-triangular, and $\epsilon(t)$ is standard bivariate normal. The observation process is $Y(t) = (y_1(t), y_2(t))$, where $y_i(t) \sim \text{normal}(x_i(t), \tau)$. The functions `rprocess`, `dprocess`, `rmeasure`, `dmeasure`, and `skeleton` are implemented using compiled C code for computational speed: see the source code for details. This object is demonstrated in the vignette "Advanced topics in pomp".

See Also

[pomp](#) and the vignettes

Examples

```
data(ou2)
plot(ou2)
coef(ou2)
x <- simulate(ou2)
plot(x)
pf <- pfilter(ou2, Np=1000)
logLik(pf)
```

pfilter

*Particle filter***Description**

Run a plain vanilla particle filter. Resampling is performed at each observation.

Usage

```
## S4 method for signature 'pomp'
pfilter(object, params, Np, tol = 1e-17,
        max.fail = 0, pred.mean = FALSE, pred.var = FALSE,
        filter.mean = FALSE, save.states = FALSE,
        save.params = FALSE, seed = NULL,
        verbose = getOption("verbose"), ...)
## S4 method for signature 'pfilterd.pomp'
pfilter(object, params, Np, tol,
        max.fail = 0, pred.mean = FALSE, pred.var = FALSE,
        filter.mean = FALSE, save.states = FALSE,
        save.params = FALSE, seed = NULL,
        verbose = getOption("verbose"), ...)
```

Arguments

object	An object of class pomp or inheriting class pomp.
params	A npars x Np numeric matrix containing the parameters corresponding to the initial state values in xstart. This must have a 'rownames' attribute. If it desired that all particles should share the same parameter values, one may supply params as a named numeric vector.
Np	the number of particles to use. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timesteps, one may specify Np either as a vector of positive integers (of length length(time(object, t0=TRUE))) or as a function taking a positive integer argument. In the latter case, Np(k) must be a single positive integer, representing the number of particles to be used at the k-th timestep: Np(0) is the number of particles to use going from timezero(object) to time(object)[1], Np(1), from timezero(object) to time(object)[1], and so on, while when T=length(time(object, t0=TRUE)), Np(T) is the number of particles to sample at the end of the time-series. When object is of class mif, this is by default the same number of particles used in the mif iterations.
tol	positive numeric scalar; particles with log likelihood below tol are considered to be "lost". A filtering failure occurs when, at some time point, all particles are lost. When all particles are lost, the conditional log likelihood at that time point is set to be log(tol).

<code>max.fail</code>	integer; the maximum number of filtering failures allowed. If the number of filtering failures exceeds this number, execution will terminate with an error.
<code>pred.mean</code>	logical; if TRUE, the prediction means are calculated for the state variables and parameters.
<code>pred.var</code>	logical; if TRUE, the prediction variances are calculated for the state variables and parameters.
<code>filter.mean</code>	logical; if TRUE, the filtering means are calculated for the state variables and parameters.
<code>save.states, save.params</code>	logical. If <code>save.states=TRUE</code> , the state-vector for each particle at each time is saved in the <code>saved.states</code> slot of the returned <code>pfilterd.pomp</code> object. If <code>save.params=TRUE</code> , the parameter-vector for each particle at each time is saved in the <code>saved.params</code> slot of the returned <code>pfilterd.pomp</code> object.
<code>seed</code>	optional; an object specifying if and how the random number generator should be initialized ('seeded'). If <code>seed</code> is an integer, it is passed to <code>set.seed</code> prior to any simulation and is returned as the "seed" element of the return list. By default, the state of the random number generator is not changed and the value of <code>.Random.seed</code> on the call is stored in the "seed" element of the return list.
<code>verbose</code>	logical; if TRUE, progress information is reported as <code>pfilter</code> works.
<code>...</code>	Additional arguments unused at present.

Value

An object of class `pfilterd.pomp`. This class inherits from class `pomp` and contains the following additional slots:

pred.mean, pred.var, filter.mean matrices of prediction means, variances, and filter means, respectively. In each of these, the rows correspond to states and parameters (if appropriate), in that order, the columns to successive observations in the time series contained in object.

eff.sample.size numeric vector containing the effective number of particles at each time point.

cond.loglik numeric vector containing the conditional log likelihoods at each time point.

saved.states If `pfilter` was called with `save.states=TRUE`, this is the list of state-vectors at each time point, for each particle. It is a length-`ntimes` list of `nvars`-by-`Np` arrays. In particular, `saved.states[[t]][,i]` can be considered a sample from $f[X_t|y_{1:t}]$.

saved.params If `pfilter` was called with `save.params=TRUE`, this is the list of parameter-vectors at each time point, for each particle. It is a length-`ntimes` list of `npars`-by-`Np` arrays. In particular, `saved.params[[t]][,i]` is the parameter portion of the i -th particle at time t .

seed the state of the random number generator at the time `pfilter` was called. If the argument `seed` was specified, this is a copy; if not, this is the internal state of the random number generator at the time of call.

Np, tol, nfail the number of particles used, failure tolerance, and number of filtering failures, respectively.

loglik the estimated log-likelihood.

These can be accessed using the `$` operator as if the returned object were a list. In addition, `logLik` returns the log likelihood. Note that if the argument `params` is a named vector, then these parameters are included in the `params` slot of the returned `pfilterd.pomp` object. That is `coef(pfilter(obj, params=theta))==theta` if `theta` is a named vector of parameters.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

M. S. Arulampalam, S. Maskell, N. Gordon, & T. Clapp. A Tutorial on Particle Filters for Online Nonlinear, Non-Gaussian Bayesian Tracking. IEEE Trans. Sig. Proc. 50:174–188, 2002.

See Also

[pomp-class](#)

Examples

```
## See the vignettes for examples.
```

pfilter-methods

Methods of the "pfilterd.pomp" class

Description

Methods of the "pfilterd.pomp" class.

Usage

```
## S4 method for signature 'pfilterd.pomp'
logLik(object, ...)
## S4 method for signature 'pfilterd.pomp'
pred.mean(object, pars, ...)
## S4 method for signature 'pfilterd.pomp'
pred.var(object, pars, ...)
## S4 method for signature 'pfilterd.pomp'
filter.mean(object, pars, ...)
```

Arguments

<code>object</code>	An object of class <code>pfilterd.pomp</code> or inheriting class <code>pfilterd.pomp</code> .
<code>pars</code>	Names of parameters.
<code>...</code>	Additional arguments unused at present.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pfilter](#), [pomp-class](#)

plugins

Plug-ins for dynamical models based on stochastic Euler algorithms

Description

Plug-in facilities for implementing discrete-time Markov processes and continuous-time Markov processes using the Euler algorithm. These can be used in the `rprocess` and `dprocess` slots of `pomp`.

Usage

```
onestep.sim(step.fun, PACKAGE)
euler.sim(step.fun, delta.t, PACKAGE)
discrete.time.sim(step.fun, delta.t = 1, PACKAGE)
gillespie.sim(rate.fun, v, d, PACKAGE)
onestep.dens(dens.fun, PACKAGE)
```

Arguments

<code>step.fun</code>	This can be either an R function or the name of a compiled, dynamically loaded native function containing the model simulator. It should be written to take a single Euler step from a single point in state space. If it is an R function, it should be of the form <code>step.fun(x,t,params,delta.t,...)</code> . Here, <code>x</code> is a named numeric vector containing the value of the state process at time <code>t</code> , <code>params</code> is a named numeric vector containing parameters, and <code>delta.t</code> is the length of the Euler time-step. If <code>step.fun</code> is the name of a native function, it must be of type “ <code>pomp_onestep_sim</code> ” as defined in the header “ <code>pomp.h</code> ”, which is included with the pomp package. For details on how to write such codes, see Details .
<code>rate.fun</code>	This can be either an R function or the name of a compiled, dynamically loaded native function that computes the transition rates. If it is an R function, it should be of the form <code>rate.fun(j,x,t,params,...)</code> . Here, <code>j</code> is the number of the event, <code>x</code> is a named numeric vector containing the value of the state process at time <code>t</code> and <code>params</code> is a named numeric vector containing parameters. If <code>rate.fun</code> is a native function, it must be of type “ <code>pomp_ssa_rate_fn</code> ” as defined in the header “ <code>pomp.h</code> ”, which is included with the package. For details on how to write such codes, see Details .
<code>v, d</code>	Matrices that specify the continuous-time Markov process in terms of its elementary events. Each should have dimensions <code>nvar</code> x <code>nevent</code> , where <code>nvar</code> is the number of state variables and <code>nevent</code> is the number of elementary events. <code>v</code>

	describes the changes that occur in each elementary event: it will usually comprise the values 1, -1, and 0 according to whether a state variable is incremented, decremented, or unchanged in an elementary event. <code>d</code> is a binary matrix that describes the dependencies of elementary event rates on state variables: <code>d[i, j]</code> will have value 1 if event rate <code>j</code> must be updated as a result of a change in state variable <code>i</code> and 0 otherwise
<code>dens.fun</code>	This can be either an R function or a compiled, dynamically loaded native function containing the model transition log probability density function. If it is an R function, it should be of the form <code>dens.fun(x1, x2, t1, t2, params, ...)</code> . Here, <code>x1</code> and <code>x2</code> are named numeric vectors containing the values of the state process at times <code>t1</code> and <code>t2</code> , <code>params</code> is a named numeric vector containing parameters. If <code>dens.fun</code> is the name of a native function, it should be of type “ <code>pomp_onestep_pdf</code> ” as defined in the header “ <code>pomp.h</code> ”, which is included with the pomp package. This function should return the log likelihood of a transition from <code>x1</code> at time <code>t1</code> to <code>x2</code> at time <code>t2</code> , assuming that no intervening transitions have occurred. For details on how to write such codes, see Details.
<code>delta.t</code>	Size of Euler time-steps.
<code>PACKAGE</code>	an optional argument that specifies to which dynamically loaded library we restrict the search for the native routines. If this is “ <code>base</code> ”, we search in the R executable itself.

Details

`onestep.sim` is the appropriate choice when it is possible to simulate the change in state from one time to another, regardless of how large the interval between them is. To use `onestep.sim`, you must write a function `step.fun` that will advance the state process from one arbitrary time to another. `euler.sim` is appropriate when one cannot do this but can compute the change in state via a sequence of smaller steps. This is desirable, for example, if one is simulating a continuous time process but is willing to approximate it using an Euler approach. `discrete.time.sim` is appropriate when the process evolves in discrete time. In this case, by default, the intervals between observations are integers.

To use `euler.sim` or `discrete.time.sim`, you must write a function `step.fun` that will take a single Euler step, of size at most `delta.t`. `euler.sim` and `discrete.time.sim` will create simulators that take as many steps as needed to get from one time to another. See below for information on how `euler.sim` chooses the actual step size it uses.

`gillespie.sim` allows exact simulation of a continuous-time, discrete-state Markov process using Gillespie’s algorithm. This is an “event-driven” approach: correspondingly, to use `gillespie.sim`, you must write a function `rate.fun` that computes the rates of each elementary event and specify two matrices (`d`, `v`) that describe, respectively, the dependencies of each rate and the consequences of each event.

`onestep.dens` will generate a suitable `dprocess` function when one can compute the likelihood of a given state transition simply by knowing the states at two times under the assumption that the state has not changed between the times. This is typically possible, for instance, when the `rprocess` function is implemented using `onestep.sim`, `euler.sim`, or `discrete.time.sim`. [NB: currently, there are no high-level algorithms in **pomp** that use `dprocess`. This function is provided for completeness only, and with an eye toward future development.]

If `step.fun` is written as an R function, it must have at least the arguments `x`, `t`, `params`, `delta.t`, and `...`. On a call to this function, `x` will be a named vector of state variables, `t` a scalar time, and `params` a named vector of parameters. The length of the Euler step will be `delta.t`. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t`, of the covariates. This is accomplished via interpolation of the covariate table. Additional arguments may be given: these will be filled by the correspondingly-named elements in the `userdata` slot of the `pomp` object (see [pomp](#)). If `step.fun` is written in a native language, it must be a function of type “`pomp_onestep_sim`” as specified in the header “`pomp.h`” included with the package (see the directory “`include`” in the installed package directory).

If `rate.fun` is written as an R function, it must have at least the arguments `j`, `x`, `t`, `params`, and `...`. Here, `j` is the an integer that indicates which specific elementary event we desire the rate of. `x` is a named vector containing the value of the state process at time `t`, and `params` is a named vector containing parameters. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t`, of the covariates. This is accomplished via interpolation of the covariate table. If `rate.fun` is a native function, it must be of type “`pomp_ssa_rate_fn`” as defined in the header “`pomp.h`”, which is included with the package.

In writing `dens.fun`, you must assume that no state transitions have occurred between `t1` and `t2`. If `dens.fun` is written as an R function, it must have at least the arguments `x1`, `x2`, `t1`, `t2`, `params`, and `...`. On a call to this function, `x1` and `x2` will be named vectors of state variables at times `t1` and `t2`, respectively. The named vector `params` contains the parameters. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t1`, of the covariates. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t1`, of the covariates. This is accomplished via interpolation of the covariate table. As above, any additional arguments will be filled by the correspondingly-named elements in the `userdata` slot of the `pomp` object (see [pomp](#)). If `dens.fun` is written in a native language, it must be a function of type “`pomp_onestep_pdf`” as defined in the header “`pomp.h`” included with the package (see the directory “`include`” in the installed package directory).

Value

`onestep.sim`, `euler.sim`, `discrete.time.sim`, and `gillespie.sim` each return functions suitable for use as the argument `rprocess` argument in [pomp](#).

`onestep.dens` returns a function suitable for use as the argument `dprocess` in [pomp](#).

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[eulermultinom](#), [pomp](#)

Examples

```
## examples showing how to use these functions
## are provided in the vignette "intro_to_pomp"
```

```
## Not run:
vignette("intro_to_pomp")

## End(Not run)
```

pmcmc

The PMCMC algorithm

Description

The Particle MCMC algorithm for estimating the parameters of a partially-observed Markov process.

Usage

```
## S4 method for signature 'pomp'
pmcmc(object, Nmcmc = 1, start, pars,
      rw.sd, dprior, Np, hyperparams, tol = 1e-17, max.fail = 0,
      verbose = getOption("verbose"), ...)
## S4 method for signature 'pfilterd.pomp'
pmcmc(object, Nmcmc = 1, start, pars,
      rw.sd, dprior, Np, hyperparams, tol, max.fail = 0,
      verbose = getOption("verbose"), ...)
## S4 method for signature 'pmcmc'
pmcmc(object, Nmcmc, start, pars,
      rw.sd, dprior, Np, hyperparams, tol, max.fail = 0,
      verbose = getOption("verbose"), ...)
## S4 method for signature 'pmcmc'
continue(object, Nmcmc = 1, start, pars,
      rw.sd, dprior, Np, hyperparams, tol, max.fail = 0,
      verbose = getOption("verbose"), ...)
```

Arguments

object	An object of class <code>pomp</code> .
Nmcmc	The number of PMCMC iterations to perform.
start	named numeric vector; the starting guess of the parameters.
pars	optional character vector naming the ordinary parameters to be estimated. Every parameter named in <code>pars</code> must have a positive random-walk standard deviation specified in <code>rw.sd</code> . Leaving <code>pars</code> unspecified is equivalent to setting it equal to the names of all parameters with a positive value of <code>rw.sd</code> .
dprior	Function of prototype <code>dprior(params, hyperparams, ..., log)</code> that evaluates the prior density. This defaults to an improper uniform prior.

<code>rw.sd</code>	numeric vector with names; used to parameterize a Gaussian random walk MCMC proposal. The random walk is only applied to parameters named in <code>pars</code> . The algorithm requires that the random walk be nontrivial, so each element in <code>rw.sd[pars]</code> must be positive. The following must be satisfied: <code>names(rw.sd)</code> must be a subset of <code>names(start)</code> , <code>rw.sd</code> must be non-negative (zeros are simply ignored), the name of every positive element of <code>rw.sd</code> must be in <code>pars</code> .
<code>Np</code>	a positive integer; the number of particles to use in each filtering operation.
<code>hyperparams</code>	optional list; parameters to be passed to <code>dprior</code> .
<code>tol</code>	numeric scalar; particles with log likelihood below <code>tol</code> are considered to be “lost”. A filtering failure occurs when, at some time point, all particles are lost.
<code>max.fail</code>	integer; maximum number of filtering failures permitted. If the number of failures exceeds this number, execution will terminate with an error.
<code>verbose</code>	logical; if TRUE, print progress reports.
<code>...</code>	Additional arguments. These are currently ignored.

Value

An object of class `pmcmc`. This class inherits from class `pfilterd.pomp` and contains the following additional slots:

`pars`, `Nmcmc`, `dprior`, `hyperparams` These slots hold the values of the corresponding arguments of the call to `pmcmc`.

`random.walk.sd` a named numeric vector containing the random-walk variances used to parameterize a Gaussian random walk MCMC proposal.

`log.prior` a numeric value containing the log of the prior density evaluated at the parameter vector in the `params` slot.

Re-running PMCMC Iterations

To re-run a sequence of PMCMC iterations, one can use the `pmcmc` method on a `pmcmc` object. By default, the same parameters used for the original PMCMC run are re-used (except for `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing PMCMC Iterations

One can continue a series of PMCMC iterations from where one left off using the `continue` method. A call to `pmcmc` to perform `Nmcmc=m` iterations followed by a call to `continue` to perform `Nmcmc=n` iterations will produce precisely the same effect as a single call to `pmcmc` to perform `Nmcmc=m+n` iterations. By default, all the algorithmic parameters are the same as used in the original call to `pmcmc`. Additional arguments will override the defaults.

Details

`pmcmc` implements an MCMC algorithm in which the true likelihood of the data is replaced by an unbiased estimate computed by a particle filter. This gives an asymptotically correct Bayesian procedure for parameter estimation (Andrieu and Roberts, 2009). An extension to give a correct Bayesian posterior distribution of unobserved state variables (Andrieu et al, 2010) has not yet been implemented.

Author(s)

Edward L. Ionides <ionides at umich dot edu>, Aaron A. King <kingaa at umich dot edu>

References

C. Andrieu, A. Doucet and R. Holenstein, Particle Markov chain Monte Carlo methods, J. R. Stat. Soc. B, to appear, 2010.

C. Andrieu and G.O. Roberts, The pseudo-marginal approach for efficient computation, Ann. Stat. 37:697-725, 2009.

See Also

[pmcmc-class](#), [pmcmc-methods](#), [pomp](#), [pomp-class](#), [pfilter](#). See the “intro_to_pomp” vignette for an example [CURRENTLY, ONLY DEMONSTRATING THE MIF ALGORITHM, WHICH IS ALGORITHMICALLY VERY SIMILAR TO PMCMC SINCE THEY BOTH DEPEND CRITICALLY ON A PARTICLE FILTERING STEP].

pmcmc-methods

Methods of the "pmcmc" class

Description

Methods of the "pmcmc" class.

Usage

```
## S4 method for signature 'pmcmc'
logLik(object, ...)
## S4 method for signature 'pmcmc'
conv.rec(object, pars, ...)
## S4 method for signature 'pmcmc'
filter.mean(object, pars, ...)
## S4 method for signature 'pmcmc'
plot(x, y = NULL, ...)
## S4 method for signature 'pmcmc'
dprior(object, params, log = FALSE, ...)
compare.pmcmc(z)
```

Arguments

object, x	The pmcmc object.
pars	Names of parameters.
y	Ignored.
z	A pmcmc object or list of pmcmc objects.
params	Named vector of parameters.
log	if TRUE, log probabilities are returned.
...	Further arguments (either ignored or passed to underlying functions).

Methods

conv.rec `conv.rec(object, pars = NULL)` returns the columns of the convergence-record matrix corresponding to the names in `pars`. By default, all rows are returned.

logLik Returns the value in the `loglik` slot.

dprior `dprior(object, params, log)` evaluates the prior density at `params` with values of the hyperparameters given by `object@hyperparams`.

pmcmc Re-runs the PMCMC iterations. See the documentation for [pmcmc](#).

compare.pmcmc Given a `pmcmc` object or a list of `pmcmc` objects, `compare.pmcmc` produces a set of diagnostic plots.

plot Plots a series of diagnostic plots. When `x` is a `pmcmc` object, `plot(x)` is equivalent to `compare.pmcmc(list(x))`.

filter.mean `filter.mean(object, pars = NULL)` returns the rows of the filtering-mean matrix corresponding to the names in `pars`. By default, all rows are returned.

print Prints a summary of the `pmcmc` object.

show Displays the `pmcmc` object.

pfilter See [pfilter](#).

Author(s)

Edward L. Ionides <ionides at umich dot edu>, Aaron A. King <kingaa at umich dot edu>

References

C. Andrieu, A. Doucet and R. Holenstein, Particle Markov chain Monte Carlo methods, J. Roy. Stat. Soc B, to appear, 2010.

C. Andrieu and G.O. Roberts, The pseudo-marginal approach for efficient computation, Ann Stat 37:697-725, 2009.

See Also

[pmcmc](#), [pomp](#), [pomp-class](#), [pfilter](#)

pomp

Partially-observed Markov process object.

Description

Create a new `pomp` object to hold a partially-observed Markov process model together with a uni- or multi-variate time series.

Usage

```
## S4 method for signature 'data.frame'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model,
      skeleton = NULL, skeleton.type = c("map", "vectorfield"),
      initializer, covar, tcovar,
      obsnames, statenames, paramnames, covarnames,
      PACKAGE, parameter.transform, parameter.inv.transform)
## S4 method for signature 'numeric'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model,
      skeleton = NULL, skeleton.type = c("map", "vectorfield"),
      initializer, covar, tcovar,
      obsnames, statenames, paramnames, covarnames,
      PACKAGE, parameter.transform, parameter.inv.transform)
## S4 method for signature 'matrix'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model,
      skeleton = NULL, skeleton.type = c("map", "vectorfield"),
      initializer, covar, tcovar,
      obsnames, statenames, paramnames, covarnames,
      PACKAGE, parameter.transform, parameter.inv.transform)
## S4 method for signature 'pomp'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model, skeleton, skeleton.type,
      initializer, covar, tcovar,
      obsnames, statenames, paramnames, covarnames,
      PACKAGE, parameter.transform, parameter.inv.transform)
```

Arguments

<code>data, times</code>	The time series data and times at which observations are made. <code>data</code> can be specified as a vector, a matrix, a data-frame, or a <code>pomp</code> object.. If <code>data</code> is a numeric vector, <code>times</code> must be a numeric vector of the same length. If <code>data</code> is a matrix, it should have dimensions <code>nobs</code> x <code>ntimes</code> , where <code>nobs</code> is the number of observed variables and <code>ntimes</code> is the number of times at which observations were made (i.e., each column is a distinct observation of the <code>nobs</code> variables). In this case, <code>times</code> must be given as a numeric vector (of length <code>ntimes</code>). If <code>data</code> is a data-frame, <code>times</code> must name of the column of observation times. Note that, in this case, <code>data</code> is a data-frame, it will be internally coerced to an array with storage-mode 'double'. Note that the times must be numeric and strictly increasing.
<code>t0</code>	The zero-time. This must be no later than the time of the first observation, <code>times[1]</code> . The stochastic dynamical system is initialized at time <code>t0</code> .
<code>rprocess</code>	optional function; a function of prototype <code>rprocess(xstart, times, params, ...)</code> that simulates from the unobserved process. The easiest way to specify <code>rprocess</code> is to use one of the plugins provided as part of the pomp package. See below for details.

dprocess	optional function; a function of prototype <code>dprocess(x, t, times, params, log, ...)</code> that evaluates the likelihood of a sequence of consecutive state transitions. The easiest way to specify dprocess is to use one of the plugins provided as part of the pomp package. It is not typically necessary (or even feasible) to define dprocess. See below for details.
rmeasure	optional; the measurement model simulator. This can be specified in one of three ways: (1) as a function of prototype <code>rmeasure(x, t, params, ...)</code> that makes a draw from the observation process given states <code>x</code> , time <code>t</code> , and parameters <code>params</code> . (2) as the name of a native (compiled) routine with prototype “ <code>pomp_measure_model_simulator</code> ” as defined in the header file “ <code>examples/pomp.h</code> ”. In the above cases, if the measurement model depends on covariates, the optional argument <code>covars</code> will be filled with interpolated values at each call. (3) using the formula-based <code>measurement.model</code> facility (see below).
dmeasure	optional; the measurement model probability density function. This can be specified in one of three ways: (1) as a function of prototype <code>dmeasure(y, x, t, params, log, ...)</code> that computes the p.d.f. of <code>y</code> given <code>x</code> , <code>t</code> , and <code>params</code> . (2) as the name of a native (compiled) routine with prototype “ <code>pomp_measure_model_density</code> ” as defined in the header file “ <code>examples/pomp.h</code> ”. In the above cases, if the measurement model depends on covariates, the optional argument <code>covars</code> will be filled with interpolated values at each call. (3) using the formula-based <code>measurement.model</code> facility (see below). As might be expected, if <code>log=TRUE</code> , this function should return the log likelihood.
<code>measurement.model</code>	optional; a formula or list of formulae, specifying the measurement model. These formulae are parsed internally and used to generate <code>rmeasure</code> and <code>dmeasure</code> functions. If <code>measurement.model</code> is given it overrides any specification of <code>rmeasure</code> or <code>dmeasure</code> . See below for an example. NB: it will typically be possible to accelerate measurement model computations by writing <code>dmeasure</code> and/or <code>rmeasure</code> functions directly.
<code>skeleton, skeleton.type</code>	The function <code>skeleton</code> specifies the deterministic skeleton of the unobserved Markov process. If we are dealing with a discrete-time Markov process, its deterministic skeleton is a map: indicate this by specifying <code>skeleton.type="map"</code> . If we are dealing with a continuous-time Markov process, its deterministic skeleton is a vectorfield: indicate this by specifying <code>skeleton.type="vectorfield"</code> . The skeleton function can be specified in one of two ways: (1) as an R function of prototype <code>skeleton(x, t, params, ...)</code> that evaluates the deterministic skeleton at state <code>x</code> and time <code>t</code> given the parameters <code>params</code> , or (2) as the name of a native (compiled) routine with prototype “ <code>pomp_skeleton</code> ” as defined in the header file “ <code>pomp.h</code> ”. If the deterministic skeleton depends on covariates, the optional argument <code>covars</code> will be filled with interpolated values of the covariates at the time <code>t</code> .
initializer	optional function of prototype <code>initializer(params, t0, ...)</code> that yields initial conditions for the state process when given a vector, <code>params</code> , of parameters. By default (i.e., if it is unspecified when <code>pomp</code> is called), the initializer assumes any parameters in <code>params</code> the names of which end in “.0” are initial values. These are simply copied over as initial conditions when <code>init.state</code> is called

	(see <code>init.state-pomp</code>). The names of the state variables are the same as the corresponding initial value parameters, but with the “.0” dropped.
<code>covar, tcovar</code>	An optional table of covariates: <code>covar</code> is the table (with one column per variable) and <code>tcovar</code> the corresponding times (one entry per row of <code>covar</code>). <code>covar</code> can be specified as either a matrix or a data frame. In either case the columns are taken to be distinct covariates. If <code>covar</code> is a data frame, <code>tcovar</code> can be either the name or the index of the time variable. If a covariate table is supplied, then the value of each of the covariates is interpolated as needed, i.e., whenever <code>rprocess</code> , <code>dprocess</code> , <code>rmeasure</code> , <code>dmeasure</code> , or <code>init.state</code> is evaluated. The resulting interpolated values are passed to the corresponding functions as a numeric vector named <code>covars</code> .
<code>obsnames, statenames, paramnames, covarnames</code>	Optional character vectors specifying the names of observables, state variables, parameters, or covariates, respectively. These are only used in the event that one or more of the basic functions (<code>rprocess</code> , <code>dprocess</code> , <code>rmeasure</code> , <code>dmeasure</code> , <code>skeleton</code>) are defined using native routines. In that case, these name vectors are matched against the corresponding names and the indices of the names are passed to the native routines. Using this facility allows one to write one or more of <code>rprocess</code> , <code>dprocess</code> , <code>rmeasure</code> , <code>dmeasure</code> , <code>skeleton</code> in native code in a way that does not depend on the order of states, parameters, and covariates at run time. See the “advanced_topic_in_pomp” vignette for more on this topic and examples.
<code>PACKAGE</code>	An optional string giving the name of the dynamically loaded library in which any native routines are to be found.
<code>parameter.transform, parameter.inv.transform</code>	Optional functions specifying parameter transformations. These functions must have arguments <code>params</code> and <code>...</code> . <code>parameter.transform</code> should transform parameters from the user’s scale to the scale that <code>rprocess</code> , <code>dprocess</code> , <code>rmeasure</code> , <code>dmeasure</code> , <code>skeleton</code> , and <code>initializer</code> will use internally. <code>parameter.inv.transform</code> should be the inverse of <code>parameter.transform</code> . Note that it is the user’s responsibility to make sure this holds. If <code>obj</code> is the constructed <code>pomp</code> object, and <code>coef(obj)</code> is non-empty, a simple check of this is <code>x <- coef(obj, transform=TRUE); obj1 <- obj; coef(obj1, transform=TRUE) <- x; identical(coef(obj), coef(obj1))</code> and <code>identical(coef(obj1, transform=TRUE), x)</code> . By default, both functions are the identity transformation. See the “introduction_to_pomp” vignette for an example.
<code>...</code>	Any additional arguments given to <code>pomp</code> will be stored in the <code>pomp</code> object and passed as arguments to each of the functions <code>rprocess</code> , <code>dprocess</code> , <code>rmeasure</code> , <code>dmeasure</code> , and <code>initializer</code> whenever they are evaluated.

Details

It is not typically necessary (or desirable, or even feasible) to define all of the functions `rprocess`, `dprocess`, `rmeasure`, `dmeasure`, and `skeleton` in any given problem. Each algorithm makes use of a different subset of these functions. In general, the specification of process-model codes `rprocess` and/or `dprocess` can be somewhat nontrivial: for this reason, `plugins` have been developed to streamline this process for the user. Currently, if one’s process model evolves in discrete time or one is willing to make such an approximation (e.g., via an Euler approximation),

then the `euler.sim` or `onestep.sim` plugin for `rprocess` and `onestep.dens` plugin for `dprocess` are available. For exact simulation of certain continuous-time Markov chains, an implementation of Gillespie’s algorithm is available (see `gillespie.sim`). To use the plugins, consult the help documentation (`?plugins`) and the vignettes.

It is anticipated that, in specific cases, it will be possible to obtain increased computational efficiency by writing custom versions of `rprocess` and/or `dprocess`. See the “Advanced Topics in `pomp`” vignette for a discussion of this. If such custom versions are desired, the following describes how each of these functions should be written in this case.

rprocess In general, the specification of `rprocess` can be somewhat nontrivial: for this reason, `plugins` have been developed to streamline this process for the user. Currently, if one’s process model evolves in discrete time or one is willing to make such an approximation (e.g., via an Euler approximation), then the `euler.sim` or `onestep.sim` plugin is available. For exact simulation of certain continuous-time Markov chains, an implementation of Gillespie’s algorithm is available (see `gillespie.sim`). To use the plugins, consult the help documentation (`?plugins`) and the vignettes.

If the plugins are not used `rprocess` must have at least the following arguments: `xstart`, `times`, `params`, and `...`. It can also take additional arguments. It is guaranteed that these will be filled with the corresponding elements the user has included as additional arguments in the construction of the `pomp` object.

In calls to `rprocess`, `xstart` can be assumed to be a rank-2 array (matrix) with rows corresponding to state variables and columns corresponding to independent realizations of the process. `params` will similarly be a rank-2 array with rows corresponding to parameters and columns corresponding to independent realizations. The columns of `params` correspond to those of `xstart`; in particular, they will agree in number. Both `xstart` and `params` will have `rownames`, which are available for use by the user.

`rprocess` must return a rank-3 array with `rownames`. Suppose `x` is the array returned. Then `dim(x)=c(nvars,nreps,ntimes)`, where `nvars` (`=nrow(xstart)`) is the number of state variables, `nreps` (`=ncol(xstart)`) is the number of independent realizations simulated, and `ntimes` is the length of the vector `times`. `x[,j,k]` is the value of the state process in the `j`-th realization at time `times[k]`. In particular, `x[, ,1]` must be identical to `xstart`. The `rownames` of `x` must correspond to those of `xstart`.

At present, the following methods make use of `rprocess`:

- `simulate`
- `pfilter`
- `mif`
- `nlf`
- `probe`
- `probe.match`

dprocess In general, the specification of `dprocess` can be somewhat nontrivial: for this reason, `plugins` have been developed to streamline this process for the user. Currently, if one’s process model evolves in discrete time or one is willing to make such an approximation (e.g., via an Euler approximation), then the `onestep.dens` plugin for `dprocess` is available. To use the plugins, consult the help documentation (`?plugins`) and the vignettes.

If the plugins are not used, `dprocess` must have at least the following arguments: `x`, `times`, `params`, `log`, and `...`. It may take additional arguments. It is guaranteed that these will be

filled with the corresponding elements the user has included as additional arguments in the construction of the `pomp` object.

In calls to `dprocess`, `x` may be assumed to be an `nvars` x `nreps` x `ntimes` array, where these terms have the same meanings as above. `params` will be a rank-2 array with rows corresponding to individual parameters and columns corresponding to independent realizations. The columns of `params` correspond to those of `x`; in particular, they will agree in number. Both `x` and `params` will have rownames, available for use by the user.

`dprocess` must return a rank-2 array (matrix). Suppose `d` is the array returned. Then `dim(d)=c(nreps, ntimes-1)`. `d[j,k]` is the probability density of the transition from state `x[,j,k-1]` at time `times[k-1]` to state `x[,j,k]` at time `times[k]`. If `log=TRUE`, then the log of the pdf is returned.

In writing this function, you may assume that the transitions are consecutive. It should be quite clear that, but for this assumption, it would be quite difficult in general to write the transition probabilities. In fact, from one perspective, the algorithms in **pomp** are designed to overcome just this difficulty.

At present, no methods in pomp make use of dprocess.

The measurement-model, deterministic skeleton, and initializer components are easily specified without the use of plugins. The following is a guide to writing these components.

`rmeasure` if provided, must take at least the arguments `x`, `t`, `params`, and `...`. It may take additional arguments, which will be filled with user-specified data as above. `x` may be assumed to be a named numeric vector of length `nvars`, (which has the same meanings as above). `t` is a scalar quantity, the time at which the measurement is made. `params` may be assumed to be a named numeric vector of length `nparams`.

`rmeasure` must return a named numeric vector. If `y` is the returned vector, then `length(y)=nobs`, where `nobs` is the number of observable variables.

At present, the following methods make use of `rmeasure`:

- `simulate`
- `nlf`
- `probe`
- `probe.match`

`dmeasure` if provided, must take at least the arguments `y`, `x`, `t`, `params`, `log`, and `...`. `y` may be assumed to be a named numeric vector of length `nobs` containing (actual or simulated) values of the observed variables; `x` will be a named numeric vector of length `nvar` containing state variables; `params`, a named numeric vector containing parameters; and `t`, a scalar, the corresponding observation time. It may take additional arguments which will be filled with user-specified data as above. `dmeasure` must return a single numeric value, the pdf of `y` given `x` at time `t`. If `log=TRUE`, then the log of the pdf is returned.

At present, the following methods make use of `dmeasure`:

- `pfilter`
- `mif`

`skeleton` If `skeleton` is an R function, it must have at least the arguments `x`, `t`, `params`, and `...`.

`x` is a numeric vector containing the coordinates of a point in state space at which evaluation of the skeleton is desired. `t` is a numeric value giving the time at which evaluation of the skeleton is desired. Of course, these will be irrelevant in the case of an autonomous skeleton. `params` is a numeric vector holding the parameters. The optional argument `covars` is a numeric

vector containing the values of the covariates at the time t . `covars` will have one value for each column of the covariate table specified when the `pomp` object was created. `covars` is constructed from the covariate table (see `covar`, below) by interpolation. `skeleton` may take additional arguments, which will be filled, as above, with user-specified data. `skeleton` must return a numeric vector of the same length as `x`. The return value is interpreted as the vectorfield (if the dynamical system is continuous) or the value of the map (if the dynamical system is discrete), at the point `x` at time `t`.

If `skeleton` is the name of a native routine, this routine must be of prototype “`pomp_skeleton`” as defined in the header “`pomp.h`” (see the “`include`” directory in the installed package directory).

At present, the following methods make use of `skeleton`:

- `trajectory`
- `traj.match`

`initializer` if provided, must have at least the arguments `params`, `t0`, and `...`. `params` is a named numeric vector of parameters. `t0` will be the time at which initial conditions are desired. `initializer` must return a named numeric vector of initial states.

Value

An object of class `pomp`. If `data` is an object of class `pomp`, then by default the returned `pomp` object is identical to `data`. If additional arguments are given, these override the defaults.

Warning

Some error checking is done by `pomp`, but complete error checking is impossible. If the user-specified functions do not conform to the above specifications (see Details), then the results may be invalid. In particular, if both `rmeasure` and `dmeasure` are specified, the user should verify that these two functions correspond to the same model and if `skeleton` is specified, the user is responsible for verifying that it corresponds to the true deterministic skeleton of the model. Each **pomp**-package algorithm uses some subset of the five basic components (`rprocess`, `dprocess`, `rmeasure`, `dmeasure`, `skeleton`). If an algorithm requires a component that was not given in the construction of the `pomp` object, an error is generated.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

`pomp-methods`, `plugins`, `time`, `time<-`, `timezero`, `timezero<-`, `coef`, `coef<-`, `obs`, `states`, `window`, `as.data.frame.pomp`

Examples

```
## For examples, see the vignettes, the data()-loadable
## example \code{pomp} objects, and the provided example files.
## Not run:
vignette("intro_to_pomp")
vignette("advanced_topics_in_pomp")
```

```

data(package="pomp")
pomp.home <- system.file("examples",package="pomp")
pomp.examples <- list.files(pomp.home)
file.show(
  file.path(pomp.home,pomp.examples),
  header=paste("=====",pomp.examples,"=====")
)

## End(Not run)

```

pomp-methods

Methods of the "pomp" class

Description

Methods of the pomp class.

Usage

```

## S4 method for signature 'pomp'
coef(object, pars, transform = FALSE, ...)
## S4 replacement method for signature 'pomp'
coef(object, pars, transform = FALSE, ...) <- value
## S4 method for signature 'pomp'
obs(object, vars, ...)
## S4 method for signature 'pomp'
data.array(object, vars, ...)
## S4 method for signature 'pomp'
states(object, vars, ...)
## S4 method for signature 'pomp'
time(x, t0 = FALSE, ...)
## S4 replacement method for signature 'pomp'
time(object, t0 = FALSE, ...) <- value
## S4 method for signature 'pomp'
timezero(object, ...)
## S4 replacement method for signature 'pomp'
timezero(object, ...) <- value
## S4 method for signature 'pomp'
window(x, start, end, ...)
## S4 method for signature 'pomp'
show(object)
## S4 method for signature 'pomp'
as(object, class)
## S3 method for class 'pomp'
as.data.frame(x, row.names, optional, ...)
## S4 method for signature 'pomp,data.frame'
coerce(from, to = "data.frame", strict = TRUE)
## S4 method for signature 'pomp'

```

```

print(x, ...)
## S4 method for signature 'pomp'
plot(x, y, variables, panel = lines,
      nc = NULL, yax.flip = FALSE,
      mar = c(0, 5.1, 0, if (yax.flip) 5.1 else 2.1),
      oma = c(6, 0, 5, 0), axes = TRUE, ...)

```

Arguments

<code>object, x</code>	The pomp object.
<code>pars</code>	optional character; names of parameters to be retrieved or set.
<code>vars</code>	optional character; names of observed variables to be retrieved.
<code>transform</code>	optional logical; should the parameter transformations be applied?
<code>value</code>	numeric; values to be assigned.
<code>t0</code>	logical; if TRUE on a call to <code>time</code> , the zero time is prepended to the time vector; if TRUE on a call to <code>time<=</code> , the first element in <code>value</code> is taken to be the initial time.
<code>start, end</code>	start and end times of the window.
<code>class</code>	character; name of the class to which object should be coerced.
<code>from, to</code>	the classes between which coercion should be performed.
<code>strict</code>	ignored.
<code>y</code>	ignored.
<code>variables</code>	optional character; names of variables to plot.
<code>panel</code>	a function of prototype <code>panel(x, col, bg, pch, type, ...)</code> which gives the action to be carried out in each panel of the display.
<code>nc</code>	the number of columns to use. Defaults to 1 for up to 4 series, otherwise to 2.
<code>yax.flip</code>	logical; if TRUE, the y-axis (ticks and numbering) should flip from side 2 (left) to 4 (right) from series to series.
<code>mar, oma</code>	the 'par' settings for 'mar' and 'oma' to use. Modify with care!
<code>axes</code>	logical; indicates if x- and y- axes should be drawn.
<code>row.names, optional</code>	ignored.
<code>...</code>	Further arguments (either ignored or passed to underlying functions).

Details

coef `coef(object)` returns the contents of the `params` slot of `object`. `coef(object, pars)` returns only those parameters named in `pars`. `coef(object, transform=TRUE)` returns `parameter.inv.transform(coef(object))` where `parameter.inv.transform` is the user parameter inverse transformation function specified when `object` was created. Likewise, `coef(object, pars, transform=TRUE)` returns `parameter.inv.transform(coef(object))[pars]`.

- coef<-** Assigns values to the `params` slot of the `pomp` object. `coef(object) <- value` has the effect of replacing the parameters of `object` with `value`. If `coef(object)` exists, then `coef(object, pars) <- value` replaces those parameters of `object` named in `pars` with the elements of `value`; the names of `value` are ignored. If some of the names in `pars` do not already name parameters in `coef(object)`, then they are concatenated. If `coef(object)` does not exist, then `coef(object, pars) <- value` assigns `value` to the parameters of `object`; in this case, the names of `object` will be `pars` and the names of `value` will be ignored. `coef(object, transform=TRUE) <- value` assigns `parameter.transform(value)` to the `params` slot of `object`. Here, `parameter.transform` is the parameter transformation function specified when `object` was created. `coef(object, pars, transform=TRUE) <- value` first, discards any names the `value` may have, sets `names(value) <- pars`, and then replaces the elements of `object`'s `params` slot `parameter.transform(value)`. In this case, if some of the names in `pars` do not already name parameters in `coef(object, transform=TRUE)`, then they are concatenated.
- obs, data.array** These functions are synonymous. `obs(object)` returns the array of observations. `obs(object, vars)` gives just the observations of variables named `vars`. `vars` may specify the variables by position or by name.
- states** `states(object)` returns the array of states. `states(object, vars)` gives just the state variables named in `vars`. `vars` may specify the variables by position or by name.
- time** `time(object)` returns the vector of observation times. `time(object, t0=TRUE)` returns the vector of observation times with the zero-time `t0` prepended.
- time<-** `time(object) <- value` replaces the observation times slot (`times`) of `object` with `value`. `time(object, t0=TRUE) <- value` has the same effect, but the first element in `value` is taken to be the initial time. The second and subsequent elements of `value` are taken to be the observation times. Those data and states (if they exist) corresponding to the new times are retained.
- timezero, timezero<-** `timezero(object)` returns the zero-time `t0`. `timezero(object) <- value` sets the zero-time to `value`.
- window** `window(x, start=t1, end=t2)` returns a new `pomp` object, identical to `x` but with only the data in the window between times `t1` and `t2` (inclusive). By default, `start` is the time of the first observation and `end` is the time of the last.
- show** Displays the `pomp` object.
- plot** Plots the data and state trajectories (if the latter exist). Additional arguments are passed to the low-level plotting routine.
- print** Prints the `pomp` object in a nice way.
- as, coerce** The `coerce` method should typically not be used directly. It is defined by `setAs` as a method to be used by `as`. A `pomp` object can be coerced to a data frame via `as(object, "data.frame")`. The data frame contains the times, the data, and the state trajectories, if they exist.
- rprocess** simulates the process model. See [rprocess-pomp](#).
- dprocess** evaluates the process model density. See [dprocess-pomp](#).
- rmeasure** simulates the measurement model. See [rmeasure-pomp](#).
- dmeasure** evaluates the measurement-model density. See [dmeasure-pomp](#).
- skeleton** evaluates the deterministic skeleton (be it a vector field or a map). See [skeleton-pomp](#).
- init.state** returns a vector of initial conditions. See [init.state-pomp](#).

simulate `simulate` can be used to simulate state and observation trajectories. See documentation under [simulate-pomp](#).

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pomp](#), [pomp-class](#), [rprocess](#), [dprocess](#), [rmeasure](#), [dmeasure](#), [init.state](#), [simulate](#)

probe

Probe a partially-observed Markov process.

Description

`probe` applies one or more “probes” to time series data and model simulations and compares the results. It can be used to diagnose goodness of fit and/or as the basis for “probe-matching”, a generalized method-of-moments approach to parameter estimation. `probe.match` calls an optimizer to adjust model parameters to do probe-matching, i.e., to minimize the discrepancy between simulated and actual data. This discrepancy is measured using the “synthetic likelihood” as defined by Wood (2010).

Usage

```
## S4 method for signature 'pomp'
probe(object, probes, params, nsim, seed = NULL, ...)
## S4 method for signature 'probed.pomp'
probe(object, probes, params, nsim, seed = NULL, ...)
## S4 method for signature 'pomp'
probe.match(object, start, est = character(0),
            probes, weights,
            nsim, seed = NULL,
            method = c("subplex", "Nelder-Mead", "SANN", "BFGS", "sannbox"),
            verbose = getOption("verbose"),
            eval.only = FALSE, fail.value = NA, ...)
## S4 method for signature 'probed.pomp'
probe.match(object, start, est = character(0),
            probes, weights,
            nsim, seed = NULL,
            method = c("subplex", "Nelder-Mead", "SANN", "BFGS", "sannbox"),
            verbose = getOption("verbose"),
            eval.only = FALSE, fail.value = NA, ...)
## S4 method for signature 'probe.matched.pomp'
probe.match(object, start, est,
            probes, weights,
            nsim, seed = NULL,
```



```
method = c("subplex", "Nelder-Mead", "SANN", "BFGS", "sannbox"),
verbose = getOption("verbose"),
eval.only = FALSE, fail.value, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>probes</code>	A single probe or a list of one or more probes. A probe is simply a scalar- or vector-valued function of one argument that can be applied to the data array of a <code>pomp</code> . A vector-valued probe must always return a vector of the same size. A number of basic examples are provided with the package (see basic.probes).
<code>params</code>	optional named numeric vector of model parameters. By default, <code>params=coef(object)</code> .
<code>nsim</code>	The number of model simulations to be computed.
<code>seed</code>	optional; if non-NULL, the random number generator will be initialized with this seed for simulations. See simulate-pomp .
<code>start</code>	named numeric vector; the initial guess of parameters.
<code>est</code>	character vector; the names of parameters to be estimated.
<code>weights</code>	optional numeric vector of relative weights. Must be of the same length as <code>probes</code> .
<code>method</code>	Optimization method. Choices are subplex and any of the methods used by optim .
<code>verbose</code>	logical; print diagnostic messages?
<code>eval.only</code>	logical; if TRUE, no optimization is attempted. Instead, the probe-mismatch value is simply evaluated at the <code>start</code> parameters.
<code>fail.value</code>	optional scalar; if non-NA, this value is substituted for non-finite values of the objective function.
<code>...</code>	Additional arguments. In the case of <code>probe</code> , these are currently ignored. In the case of <code>probe.match</code> , these are passed to <code>optim</code> or <code>subplex</code> in the control list.

Details

A call to `probe` results in the evaluation of the probe(s) in `probes` on the data. Additionally, `nsim` simulated data sets are generated (via a call to [simulate](#)) and the probe(s) are applied to each of these. The results of the probe computations on real and simulated data are stored in an object of class `probed.pomp`.

A call to `probe.match` results in an attempt to optimize the agreement between model and data, as measured by the specified probes, over the parameters named in `est`. The results, including coefficients of the fitted model and values of the probes for data and fitted-model simulations, are stored in an object of class [probe.matched.pomp](#).

Value

`probe` returns an object of class `probed.pomp`. `probed.pomp` is derived from the [pomp](#) class and therefore have all the slots of `pomp`. In addition, a `probed.pomp` class has the following slots:

probes list of the probes applied.

datvals, simvals values of each of the probes applied to the real and simulated data, respectively.

quantiles fraction of simulations with probe values less than the value of the probe of the data.

pvals two-sided p-values: fraction of the simvals that deviate more extremely from the mean of the simvals than does datavals.

synth.loglik the log synthetic likelihood (Wood 2010). This is the likelihood assuming that the probes are multivariate-normally distributed.

`probe.match` returns an object of class `probe.matched.pomp`, which is derived from class `probed.pomp`. `probe.matched.pomp` objects therefore have all the slots above plus the following:

est, weights, fail.value values of the corresponding arguments in the call to `spect.match`.

value value of the objective function.

evals number of function and gradient evaluations by the optimizer. See [optim](#).

convergence, msg Convergence code and message from the optimizer. See [optim](#).

Author(s)

Daniel C. Reuman (d.reuman at imperial dot ac dot uk)

Aaron A. King (kingaa at umich dot edu)

References

B. E. Kendall, C. J. Briggs, W. M. Murdoch, P. Turchin, S. P. Ellner, E. McCauley, R. M. Nisbet, S. N. Wood Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches, *Ecology*, 80:1789–1805, 1999.

S. N. Wood Statistical inference for noisy nonlinear ecological dynamic systems, *Nature*, 466: 1102–1104, 2010.

See Also

[pomp-class](#), [pomp-methods](#), [basic.probes](#), [probe.match](#)

Examples

```
data(ou2)
good <- probe(
  ou2,
  probes=list(
    y1.mean=probe.mean(var="y1"),
    y2.mean=probe.mean(var="y2"),
    y1.sd=probe.sd(var="y1"),
    y2.sd=probe.sd(var="y2"),
    extra=function(x)max(x["y1",]),
  ),
  nsim=500
)
summary(good)
plot(good)
```

```

bad <- probe(
  ou2,
  params=c(alpha.1=0.1,alpha.4=0.2,x1.0=0,x2.0=0,
            alpha.2=-0.5,alpha.3=0.3,
            sigma.1=3,sigma.2=-0.5,sigma.3=2,
            tau=1),
  probes=list(
    y1.mean=probe.mean(var="y1"),
    y2.mean=probe.mean(var="y2"),
    y1.sd=probe.sd(var="y1"),
    y2.sd=probe.sd(var="y2"),
    extra=function(x)range(x["y1",])
  ),
  nsim=500
)
summary(bad)
plot(bad)

```

probed.pomp-methods	<i>Methods of the "probed.pomp", "probe.matched.pomp", "spect.pomp", and "spect.matched.pomp" classes</i>
---------------------	---

Description

Methods of the `probed.pomp`, `probe.matched.pomp`, `spect.pomp`, and `spect.matched.pomp` classes

Usage

```

## S4 method for signature 'probed.pomp'
summary(object, ...)
## S4 method for signature 'probed.pomp'
plot(x, y, ...)
## S4 method for signature 'probe.matched.pomp'
summary(object, ...)
## S4 method for signature 'probe.matched.pomp'
plot(x, y, ...)
## S4 method for signature 'spect.pomp'
summary(object, ...)
## S4 method for signature 'probed.pomp'
logLik(object, ...)
## S4 method for signature 'spect.pomp'
plot(x, y, max.plots.per.page = 4,
     plot.data = TRUE,
     quantiles = c(.025, .25, .5, .75, .975),
     quantile.styles = list(lwd=1, lty=1, col="gray70"),
     data.styles = list(lwd=2, lty=2, col="black"))
## S4 method for signature 'spect.matched.pomp'
summary(object, ...)

```

```
## S4 method for signature 'spect.matched.pomp'
plot(x, y, ...)
## S4 method for signature 'probed.pomp'
as(object, class)
```

Arguments

<code>object, x</code>	the object to be summarized or plotted.
<code>y</code>	ignored.
<code>max.plots.per.page</code>	maximum number of plots per page
<code>plot.data</code>	plot the data spectrum?
<code>quantiles</code>	quantiles to plot
<code>quantile.styles</code>	plot style parameters for the quantiles
<code>data.styles</code>	plot style parameters for the data spectrum
<code>class</code>	character; name of the class to which object should be coerced.
<code>...</code>	Further arguments (either ignored or passed to underlying functions).

Methods

plot displays diagnostic plots.

summary displays summary information.

logLik returns the synthetic likelihood for the probes. NB: in general, this is not the same as the likelihood.

as when a 'probed.pomp' is coerced to a 'data.frame', the first row gives the probes applied to the data; the rest of the rows give the probes evaluated on simulated data. The rownames of the result can be used to distinguish these.

Author(s)

Daniel C. Reuman (d.reuman at imperial dot ac dot uk)

Aaron A. King (kingaa at umich dot edu)

See Also

[probe](#), [probed.pomp](#), [probe.matched.pomp](#), [probe.match](#)

profileDesign*Design matrices for likelihood profile calculations.*

Description

profileDesign generates a data-frame where each row can be used as the starting point for a profile likelihood calculation.

Usage

```
profileDesign(..., lower, upper, nprof)
```

Arguments

...	Specifies the parameters over which to profile.
lower, upper	Named numeric vectors, specifying the range over which the other parameters are to be sampled.
nprof	The number of starts per profile point.

Value

profileDesign returns a data frame with nprof points per profile point. The other parameters in vars are sampled using sobol.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[sobol](#)

Examples

```
## A one-parameter profile design:
x <- profileDesign(p=1:10, lower=c(a=0, b=0), upper=c(a=1, b=5), nprof=20)
dim(x)
plot(x)
## A two-parameter profile design:
x <- profileDesign(p=1:10, q=3:5, lower=c(a=0, b=0), upper=c(b=5, a=1), nprof=20)
dim(x)
plot(x)
```

ricker	<i>Ricker model with Poisson observations.</i>
--------	--

Description

ricker is a pomp object encoding a stochastic Ricker model with Poisson measurement error.

Usage

```
data(ricker)
```

Details

The state process is $N_{t+1} = rN_t \exp(-N_t + e_t)$, where the e_t are i.i.d. normal random deviates with variance σ^2 . The observed variables y_t are distributed as $\text{Poisson}(\phi N_t)$.

See Also

[pomp-class](#) and the vignettes

Examples

```
data(ricker)
plot(ricker)
coef(ricker)
```

rw2	<i>Two-dimensional random-walk process</i>
-----	--

Description

rw2 is a pomp object encoding a 2-D normal random walk.

Usage

```
data(rw2)
```

Details

The random-walk process is fully but noisily observed.

See Also

[pomp-class](#) and the vignettes

Examples

```
data(rw2)
plot(rw2)
x <- simulate(rw2, nsim=10, seed=20348585L, params=c(x1.0=0, x2.0=0, s1=1, s2=3, tau=1))
plot(x[[1]])
```

simulate-pomp

Running simulations of a partially-observed Markov process

Description

simulate can be used to generate simulated data sets and/or to simulate the state process.

Usage

```
## S4 method for signature 'pomp'
simulate(object, nsim = 1, seed = NULL, params,
         states = FALSE, obs = FALSE, times, t0, ...)
```

Arguments

object	An object of class pomp.
nsim	The number of simulations to perform. Note that the number of replicates will be nsim times ncol(xstart).
seed	optional; if set, the pseudorandom number generator (RNG) will be initialized with seed. the random seed to use. The RNG will be restored to its original state afterward.
params	either a named numeric vector or a numeric matrix with rownames. The parameters to use in simulating the model. If params is not given, then the contents of the params slot of object will be used, if they exist.
states	Do we want the state trajectories?
obs	Do we want data-frames of the simulated observations?
times, t0	times specifies the times at which simulated observations will be made. t0 specifies the start time (the time at which the initial conditions hold). The default for times is is times=time(object, t0=FALSE) and t0=timezero(object), respectively.
...	further arguments that are currently ignored.

Details

Simulation of the state process and of the measurement process are each accomplished by a single call to the user-supplied rprocess and rmeasure functions, respectively. This makes it possible for the user to write highly optimized code for these potentially expensive computations.

Value

If `states=FALSE` and `obs=FALSE` (the default), a list of `nsim` pomp objects is returned. Each has a simulated data set, together with the parameters used (in slot `params`) and the state trajectories also (in slot `states`). If `times` is specified, then the simulated observations will be at times `times`.

If `nsim=1`, then a single pomp object is returned (and not a singleton list).

If `states=TRUE` and `obs=FALSE`, simulated state trajectories are returned as a rank-3 array with dimensions `nvar` x `(ncol(params)*nsim)` x `ntimes`. Here, `nvar` is the number of state variables and `ntimes` the length of the argument `times`. The measurement process is not simulated in this case.

If `states=FALSE` and `obs=TRUE`, simulated observations are returned as a rank-3 array with dimensions `nobs` x `(ncol(params)*nsim)` x `ntimes`. Here, `nobs` is the number of observables.

If both `states=TRUE` and `obs=TRUE`, then a named list is returned. It contains the state trajectories and simulated observations as above.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pomp-class](#)

Examples

```
data(ou2)
x <- simulate(ou2, seed=3495485, nsim=10)
x <- simulate(ou2, seed=3495485, nsim=10, states=TRUE, obs=TRUE)
```

sir

Seasonal SIR model implemented using two stochastic simulation algorithms.

Description

`euler.sir` is a pomp object encoding a simple seasonal SIR model. Simulation is performed using an Euler multinomial approximation (AKA tau leap method). `gillespie.sir` has the same model implemented using Gillespie's algorithm.

Usage

```
data(euler.sir)
data(gillespie.sir)
```

See Also

[pomp-class](#) and the vignettes

Examples

```
data(euler.sir)
plot(euler.sir)
x <- simulate(euler.sir,nsim=10,seed=20348585)
plot(x[[1]])

data(gillespie.sir)
plot(gillespie.sir)
x <- simulate(gillespie.sir,nsim=1,seed=20348585)
plot(x)
```

sliceDesign	<i>Design matrices for likelihood slices.</i>
-------------	---

Description

sliceDesign generates a data-frame representing points taken along one or more slices through a point in a multidimensional space.

Usage

```
sliceDesign(center, ...)
```

Arguments

center	center is a named numeric vectors specifying the point through which the slice(s) is (are) to be taken.
...	Additional numeric vector arguments specify the slices.

Value

sliceDesign returns a data frame with one row per point along a slice. The column slice is a factor that tells which slice each point belongs to.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[profileDesign](#)

Examples

```
## A single 11-point slice through the point c(A=3,B=8,C=0) along the B direction.
x <- sliceDesign(center=c(A=3,B=8,C=0),B=seq(0,10,by=1))
dim(x)
plot(x)
## Two slices through the same point along the A and C directions.
x <- sliceDesign(c(A=3,B=8,C=0),A=seq(0,5,by=1),C=seq(0,5,length=11))
dim(x)
plot(x)
```

sobol	<i>Sobol' low-discrepancy sequence</i>
-------	--

Description

Generate a data-frame containing a Sobol' low-discrepancy sequence.

Usage

```
sobol(vars, n)
sobolDesign(lower, upper, nseq)
```

Arguments

vars	Named list of ranges of variables.
lower, upper	named numeric vectors giving the lower and upper bounds of the ranges, respectively.
n, nseq	Number of vectors requested.

Value

sobol	Returns a data frame with n 'observations' of the variables in vars.
sobolDesign	Returns a data frame with nseq 'observations' of the variables over the range specified.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

W. H. Press, S. A. Teukolsky, W. T. Vetterling, & B. P. Flannery, Numerical Recipes in C, Cambridge University Press, 1992

See Also

[sliceDesign](#), [profileDesign](#)

Examples

```
plot(sobol(vars=list(a=c(0,1),b=c(100,200)),100))
plot(sobolDesign(lower=c(a=0,b=100),upper=c(b=200,a=1),100))
```

spect	<i>Power spectrum computation for partially-observed Markov processes.</i>
-------	--

Description

spect estimates the power spectrum of time series data and model simulations and compares the results. It can be used to diagnose goodness of fit and/or as the basis for frequency-domain parameter estimation (spect.match).

spect.match tries to match the power spectrum of the model to that of the data. It calls an optimizer to adjust model parameters to minimize the discrepancy between simulated and actual data.

Usage

```
## S4 method for signature 'pomp'
spect(object, params, vars, kernel.width, nsim, seed = NULL,
      transform = identity,
      detrend = c("none","mean","linear","quadratic"),
      ...)
## S4 method for signature 'spect.pomp'
spect(object, params, vars, kernel.width, nsim, seed = NULL, transform,
      detrend, ...)
spect.match(object, start, est = character(0),
            vars, nsim, seed = NULL,
            kernel.width, transform = identity,
            detrend = c("none","mean","linear","quadratic"),
            weights, method = c("subplex","Nelder-Mead","SANN"),
            verbose = getOption("verbose"),
            eval.only = FALSE, fail.value = NA, ...)
```

Arguments

object	An object of class pomp.
params	optional named numeric vector of model parameters. By default, params=coef(object).
vars	optional; names of observed variables for which the power spectrum will be computed. This must be a subset of rownames(obs(object)). By default, the spectrum will be computed for all observables.
kernel.width	width parameter for the smoothing kernel used for calculating the estimate of the spectrum.
nsim	number of model simulations to be computed.

seed	optional; if non-NULL, the random number generator will be initialized with this seed for simulations. See simulate-pomp .
transform	function; this transformation will be applied to the observables prior to estimation of the spectrum, and prior to any detrending.
detrend	de-trending operation to perform. Options include no detrending, and subtraction of constant, linear, and quadratic trends from the data. Detrending is applied to each data series and to each model simulation independently.
weights	optional. The mismatch between model and data is measured by a weighted average of mismatch at each frequency. By default, all frequencies are weighted equally. weights can be specified either as a vector (which must have length equal to the number of frequencies) or as a function of frequency. If the latter, weights(freq) must return a nonnegative weight for each frequency.
start	named numeric vector; the initial guess of parameters.
est	character vector; the names of parameters to be estimated.
method	Optimization method. Choices are subplex and any of the methods used by optim .
verbose	logical; print diagnostic messages?
eval.only	logical; if TRUE, no optimization is attempted. Instead, the probe-mismatch value is simply evaluated at the start parameters.
fail.value	optional scalar; if non-NA, this value is substituted for non-finite values of the objective function.
...	Additional arguments. In the case of spect, these are currently ignored. In the case of spect.match, these are passed to optim or subplex in the control list.

Details

A call to `spect` results in the estimation of the power spectrum for the (transformed, detrended) data and `nsim` model simulations. The results of these computations are stored in an object of class [spect.pomp](#).

A call to `spect.match` results in an attempt to optimize the agreement between model and data spectrum over the parameters named in `est`. The results, including coefficients of the fitted model and power spectra of fitted model and data, are stored in an object of class [spect.matched.pomp](#).

Value

`spect` returns an object of class `spect.pomp`, which is derived from class [pomp](#) and therefore has all the slots of that class. In addition, `spect.pomp` objects have the following slots:

kernel.width width parameter of the smoothing kernel used.

transform transformation function used.

freq numeric vector of the frequencies at which the power spectrum is estimated.

datspec, simspec estimated power spectra for data and simulations, respectively.

pvals one-sided p-values: fraction of the simulated spectra that differ more from the mean simulated spectrum than does the data. The metric used is L^2 distance.

detrend detrending option used.

`spect.match` returns an object of class `spect.matched.pomp`, which is derived from class `{spect.pomp}` and therefore has all the slots of that class. In addition, `spect.matched.pomp` objects have the following slots:

est, weights, fail.value values of the corresponding arguments in the call to `spect.match`.

evals number of function and gradient evaluations by the optimizer. See [optim](#).

value Value of the objective function.

convergence, msg Convergence code and message from the optimizer. See [optim](#).

Author(s)

Daniel C. Reuman (d.reuman at imperial dot ac dot uk)

Cai GoGwilt

Aaron A. King (kingaa at umich dot edu)

References

D.C. Reuman, R.A. Desharnais, R.F. Costantino, O. Ahmad, J.E. Cohen (2006) Power spectra reveal the influence of stochasticity on nonlinear population dynamics. *Proceedings of the National Academy of Sciences* **103**, 18860-18865.

D.C. Reuman, R.F. Costantino, R.A. Desharnais, J.E. Cohen (2008) Color of environmental noise affects the nonlinear dynamics of cycling, stage-structured populations. *Ecology Letters*, **11**, 820-830.

See Also

[pomp-class](#), [pomp-methods](#), [probe](#), [probe.match](#)

Examples

```
data(ou2)
good <- spect(
  ou2,
  vars=c("y1", "y2"),
  kernel.width=3,
  detrend="mean",
  nsim=500
)
summary(good)
plot(good)

ou2.bad <- ou2
coef(ou2.bad, c("x1.0", "x2.0", "alpha.1", "alpha.4")) <- c(0, 0, 0.1, 0.2)
bad <- spect(
  ou2.bad,
  vars=c("y1", "y2"),
  kernel.width=3,
  detrend="mean",
```

```

      nsim=500
    )
summary(bad)
plot(bad)

```

traj.match

Trajectory matching

Description

Match trajectories to data.

Usage

```

## S4 method for signature 'pomp'
traj.match(object, start, est,
  method = c("Nelder-Mead", "sannbox", "subplex"),
  gr = NULL, eval.only = FALSE, ...)
## S4 method for signature 'traj.matched.pomp'
traj.match(object, start, est,
  method = c("Nelder-Mead", "sannbox", "subplex"),
  gr = NULL, eval.only = FALSE, ...)

```

Arguments

object	A pomp object.
start	initial guess for parameters.
est	character vector containing the names of parameters to be estimated.
method	Optimization method. Choices are subplex , “sannbox”, and any of the methods used by optim .
gr	Passed to optim .
eval.only	logical; if TRUE, no optimization is attempted and the log likelihood value is evaluated at the start parameters.
...	Arguments that will be passed to optim or subplex via their control lists.

Details

Trajectory matching is accomplished using [optim](#). The [trajectory](#) method is used for this, which in turn uses the skeleton slot of the pomp object. The quantity maximized is the likelihood of the data given the trajectory, as returned by [dmeasure](#).

Value

An object of class `traj.matched.pomp`. This class inherits from class `pomp` and contains the following additional slots:

evals number of function and gradient evaluations by the optimizer. See `optim`.

value value of the objective function. Larger values indicate better fit (i.e., `traj.match` attempts to maximize this quantity).

convergence, msg convergence code and message from the optimizer. See `optim`.

Available methods for objects of this type include `summary` and `logLik`. The other slots of this object can be accessed via the `$` operator.

See Also

`trajectory`, `pomp`, `optim`, `subplex`

Examples

```
data(ou2)
true.p <- c(
  alpha.1=0.9,alpha.2=0,alpha.3=-0.4,alpha.4=0.99,
  sigma.1=2,sigma.2=0.1,sigma.3=2,
  tau=1,
  x1.0=50,x2.0=-50
)
simdata <- simulate(ou2,nsim=1,params=true.p,seed=43553)
guess.p <- true.p
res <- traj.match(
  simdata,
  start=guess.p,
  est=c('alpha.1','alpha.3','alpha.4','x1.0','x2.0','tau'),
  maxit=2000,
  method="Nelder-Mead",
  reltol=1e-8
)

summary(res)

plot(range(time(res)),range(c(obs(res),states(res))),type='n',xlab="time",ylab="x,y")
points(y1~time,data=as(res,"data.frame"),col='blue')
points(y2~time,data=as(res,"data.frame"),col='red')
lines(x1~time,data=as(res,"data.frame"),col='blue')
lines(x2~time,data=as(res,"data.frame"),col='red')
```

trajectory

Compute trajectories of the deterministic skeleton.

Description

The method `trajectory` computes a trajectory of the deterministic skeleton of a Markov process. In the case of a discrete-time system, the deterministic skeleton is a map and a trajectory is obtained by iterating the map. In the case of a continuous-time system, the deterministic skeleton is a vectorfield; `trajectory` integrates the vectorfield to obtain a trajectory.

Usage

```
## S4 method for signature 'pomp'
trajectory(object, params, times, t0, ...)
```

Arguments

<code>object</code>	an object of class <code>pomp</code> .
<code>params</code>	a rank-2 array of parameters. Each column of <code>params</code> is a distinct parameter vector.
<code>times, t0</code>	<code>times</code> is a numeric vector specifying the times at which a trajectory is desired. <code>t0</code> specifies the start time (the time at which the initial conditions hold). The default for <code>times</code> is <code>times=time(object, t0=FALSE)</code> and <code>t0=timezero(object)</code> , respectively.
<code>...</code>	additional arguments are passed to the ODE integrator if the skeleton is a vectorfield and ignored if it is a map. See ode for a description of the additional arguments accepted.

Details

This function makes repeated calls to the user-supplied skeleton of the `pomp` object. For specifications on supplying this, see [pomp](#).

When the skeleton is a vectorfield, `trajectory` integrates it using [ode](#).

Unresolved issue: What is the behavior if `type="map"` and `times` are non-integral?

Value

Returns an array of dimensions `nvar x nreps x ntimes`. If `x` is the returned matrix, `x[i,j,k]` is the *i*-th component of the state vector at time `times[k]` given parameters `params[,j]`.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pomp](#), [traj.match](#), [ode](#)

Examples

```
data(euler.sir)
x <- trajectory(euler.sir)
plot(time(euler.sir),x["I",1,],type='l',xlab='time',ylab='I')
lines(time(euler.sir),x["cases",1,],col='red')

coef(euler.sir,c("gamma")) <- log(12)
x <- trajectory(euler.sir)
plot(time(euler.sir),x["I",1,],type='l',xlab='time',ylab='I')
lines(time(euler.sir),x["cases",1,],col='red')
```

verhulst

*Simple Verhulst-Pearl (logistic) model.***Description**

verhulst is a pomp object encoding a univariate stochastic logistic model with measurement error.

Usage

```
data(verhulst)
```

Details

The model is written as an Ito diffusion, $dn = rn(1 - n/K)dt + \sigma ndW$, where W is a Wiener process. It is implemented using the [euler.sim](#) plug-in.

See Also

[pomp-class](#) and the vignettes

Examples

```
data(verhulst)
plot(verhulst)
coef(verhulst)
params <- cbind(
  c(n.0=100,K=10000,r=0.2,sigma=0.4,tau=0.1),
  c(n.0=1000,K=11000,r=0.1,sigma=0.4,tau=0.1)
)
x <- simulate(verhulst,params=params,states=TRUE)
matplot(time(verhulst),t(x['n',,]),type='l')
y <- trajectory(verhulst,params=params)
matlines(time(verhulst),t(y['n',,]),type='l',lwd=2)
```

Index

*Topic **datasets**

- blowflies, [7](#)
- dacca, [9](#)
- gompertz, [12](#)
- LondonYorke, [12](#)
- ou2, [20](#)
- ricker, [46](#)
- rw2, [46](#)
- sir, [48](#)
- verhulst, [57](#)

*Topic **design**

- profileDesign, [45](#)
- sliceDesign, [49](#)
- sobol, [50](#)

*Topic **distribution**

- eulermultinom, [10](#)

*Topic **models**

- basic.probes, [5](#)
- mif, [13](#)
- mif-methods, [16](#)
- nlf, [18](#)
- pfilter, [21](#)
- pfilter-methods, [23](#)
- plugins, [24](#)
- pmcmc, [27](#)
- pmcmc-methods, [29](#)
- pomp, [30](#)
- pomp-methods, [37](#)
- pomp-package, [2](#)
- probe, [40](#)
- probed.pomp-methods, [43](#)
- simulate-pomp, [47](#)
- spect, [51](#)
- traj.match, [54](#)
- trajectory, [56](#)

*Topic **smooth**

- B-splines, [3](#)

*Topic **ts**

- basic.probes, [5](#)

- bsmc, [7](#)
- mif, [13](#)
- mif-methods, [16](#)
- nlf, [18](#)
- pfilter, [21](#)
- pfilter-methods, [23](#)
- pmcmc, [27](#)
- pmcmc-methods, [29](#)
- pomp, [30](#)
- pomp-methods, [37](#)
- pomp-package, [2](#)
- probe, [40](#)
- probed.pomp-methods, [43](#)
- simulate-pomp, [47](#)
- spect, [51](#)
- traj.match, [54](#)
- trajectory, [56](#)
- \$.pfilterd.pomp-method
(pfilter-methods), [23](#)
- \$.traj.matched.pomp-method
(traj.match), [54](#)
- \$.pfilterd.pomp (pfilter-methods), [23](#)
- \$.traj.matched.pomp (traj.match), [54](#)
- as.pomp-method (pomp-methods), [37](#)
- as.probed.pomp-method
(probed.pomp-methods), [43](#)
- as.data.frame.pomp, [36](#)
- as.data.frame.pomp (pomp-methods), [37](#)
- B-splines, [3](#)
- basic.probes, [5](#), [41](#), [42](#)
- blowflies, [7](#)
- blowflies1 (blowflies), [7](#)
- blowflies2 (blowflies), [7](#)
- bsmc, [3](#), [7](#)
- bsmc.pomp-method (bsmc), [7](#)
- bsmc-pomp (bsmc), [7](#)
- bspline.basis (B-splines), [3](#)

- coef, 17, 36
- coef, pomp-method (pomp-methods), 37
- coef-pomp (pomp-methods), 37
- coef<- (pomp-methods), 37
- coef<-, pomp-method (pomp-methods), 37
- coef<--pomp (pomp-methods), 37
- coef<-, 36
- coerce, pomp, data.frame-method (pomp-methods), 37
- coerce, probed.pomp, data.frame-method (probed.pomp-methods), 43
- compare.mif (mif-methods), 16
- compare.pmcmc (pmcmc-methods), 29
- continue (mif), 13
- continue, mif-method (mif), 13
- continue, pmcmc-method (pmcmc), 27
- continue-mif (mif), 13
- continue-pmcmc (pmcmc), 27
- conv.rec (mif-methods), 16
- conv.rec, mif-method (mif-methods), 16
- conv.rec, pmcmc-method (pmcmc-methods), 29
- conv.rec-mif (mif-methods), 16
- conv.rec-pmcmc (pmcmc-methods), 29
- dacca, 9
- data.array (pomp-methods), 37
- data.array, pomp-method (pomp-methods), 37
- data.array-pomp (pomp-methods), 37
- data.frame-pomp (pomp), 30
- deulermultinom (eulermultinom), 10
- discrete.time.sim (plugins), 24
- dmeasure, 40, 54
- dmeasure-pomp, 39
- dprior (pmcmc-methods), 29
- dprior, pmcmc-method (pmcmc-methods), 29
- dprior-pmcmc (pmcmc-methods), 29
- dprocess, 40
- dprocess-pomp, 39
- euler.sim, 34, 57
- euler.sim (plugins), 24
- euler.sir, 10
- euler.sir (sir), 48
- eulermultinom, 10, 26
- filter.mean (pfilter-methods), 23
- filter.mean, pfilterd.pomp-method (pfilter-methods), 23
- filter.mean, pmcmc-method (pmcmc-methods), 29
- filter.mean-pfilterd.pomp (pfilter-methods), 23
- filter.mean-pmcmc (pmcmc-methods), 29
- gillespie.sim, 34
- gillespie.sim (plugins), 24
- gillespie.sir (sir), 48
- gompertz, 12
- init.state, 40
- init.state-pomp, 33, 39
- kernel, 5
- logLik, mif-method (mif-methods), 16
- logLik, pfilterd.pomp-method (pfilter-methods), 23
- logLik, pmcmc-method (pmcmc-methods), 29
- logLik, probed.pomp-method (probed.pomp-methods), 43
- logLik, traj.matched.pomp-method (traj.match), 54
- logLik-mif (mif-methods), 16
- logLik-pfilterd.pomp (pfilter-methods), 23
- logLik-pmcmc (pmcmc-methods), 29
- logLik-probed.pomp (probed.pomp-methods), 43
- logLik-traj.matched.pomp (traj.match), 54
- LondonYorke, 12
- matrix-pomp (pomp), 30
- mean, 5
- mif, 3, 13, 17, 18, 34, 35
- mif, mif-method (mif), 13
- mif, pfilterd.pomp-method (mif), 13
- mif, pomp-method (mif), 13
- mif-methods, 16
- mif-methods, 16
- mif-mif (mif), 13
- mif-pfilterd.pomp (mif), 13
- mif-pomp (mif), 13
- nlf, 3, 18, 34, 35
- numeric-pomp (pomp), 30

- obs, 6, 36
- obs (pomp-methods), 37
- obs, pomp-method (pomp-methods), 37
- obs-pomp (pomp-methods), 37
- ode, 56
- onestep.dens, 34
- onestep.dens (plugins), 24
- onestep.sim, 34
- onestep.sim (plugins), 24
- optim, 19, 41, 42, 52–55
- ou2, 20
- paste, 4
- periodic.bspline.basis (B-splines), 3
- pfilter, 2, 3, 16, 18, 21, 24, 29, 30, 34, 35
- pfilter, pfilterd.pomp-method (pfilter), 21
- pfilter, pomp-method (pfilter), 21
- pfilter-methods, 23
- pfilter-pfilterd.pomp (pfilter), 21
- pfilter-pomp (pfilter), 21
- pfilterd.pomp, 22, 28
- pfilterd.pomp-class (pfilter), 21
- plot, mif-method (mif-methods), 16
- plot, pmcmc-method (pmcmc-methods), 29
- plot, pomp-method (pomp-methods), 37
- plot, probe.matched.pomp-method (probed.pomp-methods), 43
- plot, probed.pomp-method (probed.pomp-methods), 43
- plot, spect.matched.pomp-method (probed.pomp-methods), 43
- plot, spect.pomp-method (probed.pomp-methods), 43
- plot-mif (mif-methods), 16
- plot-pmcmc (pmcmc-methods), 29
- plot-pomp (pomp-methods), 37
- plot-probe.matched.pomp (probed.pomp-methods), 43
- plot-probed.pomp (probed.pomp-methods), 43
- plot-spect.pomp (probed.pomp-methods), 43
- plugins, 24, 31–34, 36
- pmcmc, 3, 27, 30
- pmcmc, pfilterd.pomp-method (pmcmc), 27
- pmcmc, pmcmc-method (pmcmc), 27
- pmcmc, pomp-method (pmcmc), 27
- pmcmc-class, 29
- pmcmc-methods, 29
- pmcmc-class (pmcmc), 27
- pmcmc-methods, 29
- pmcmc-pfilterd.pomp (pmcmc), 27
- pmcmc-pmcmc (pmcmc), 27
- pmcmc-pomp (pmcmc), 27
- pomp, 2, 3, 10, 16, 18, 20, 22, 26, 29, 30, 30, 33, 36, 40, 41, 52, 55, 56
- pomp, data.frame-method (pomp), 30
- pomp, matrix-method (pomp), 30
- pomp, numeric-method (pomp), 30
- pomp, pomp-method (pomp), 30
- pomp-class, 7, 12, 13, 16, 18, 29, 30, 46, 48, 57
- pomp-class, 7, 9, 23, 24, 40, 42, 48, 53
- pomp-methods, 7, 36, 37, 42, 53
- pomp-package, 2
- pomp-pomp (pomp), 30
- pred.mean (pfilter-methods), 23
- pred.mean, pfilterd.pomp-method (pfilter-methods), 23
- pred.mean-pfilterd.pomp (pfilter-methods), 23
- pred.var (pfilter-methods), 23
- pred.var, pfilterd.pomp-method (pfilter-methods), 23
- pred.var-pfilterd.pomp (pfilter-methods), 23
- print, pomp-method (pomp-methods), 37
- print-pomp (pomp-methods), 37
- probe, 6, 34, 35, 40, 44, 53
- probe, pomp-method (probe), 40
- probe, probed.pomp-method (probe), 40
- probe-pomp (probe), 40
- probe-probed.pomp (probe), 40
- probe.acf (basic.probes), 5
- probe.ccf (basic.probes), 5
- probe.marginal (basic.probes), 5
- probe.match, 3, 6, 34, 35, 42, 44, 53
- probe.match (probe), 40
- probe.match, pomp-method (probe), 40
- probe.match, probe.matched.pomp-method (probe), 40
- probe.match, probed.pomp-method (probe), 40
- probe.match-pomp (probe), 40
- probe.match-probe.matched.pomp (probe), 40

- probe.match-probed.pomp (probe), 40
- probe.matched.pomp, 41, 44
- probe.matched.pomp-class (probe), 40
- probe.matched.pomp-methods
 - (probed.pomp-methods), 43
- probe.mean (basic.probes), 5
- probe.median (basic.probes), 5
- probe.nlar (basic.probes), 5
- probe.period (basic.probes), 5
- probe.quantile (basic.probes), 5
- probe.sd (basic.probes), 5
- probe.var (basic.probes), 5
- probed.pomp, 44
- probed.pomp-class (probe), 40
- probed.pomp-methods, 43
- profileDesign, 45, 49, 50
- quantile, 5
- reulermultinom (eulermultinom), 10
- ricker, 46
- rmeasure, 40
- rmeasure-pomp, 39
- rprocess, 40
- rprocess-pomp, 39
- rw2, 46
- show, pomp-method (pomp-methods), 37
- show-pomp (pomp-methods), 37
- simulate, 2, 3, 34, 35, 40, 41
- simulate, pomp-method (simulate-pomp), 47
- simulate-pomp, 40
- simulate-pomp, 41, 47, 52
- sir, 48
- skeleton-pomp, 39
- sliceDesign, 49, 50
- sobol, 45, 50
- sobolDesign (sobol), 50
- spect, 51
- spect, pomp-method (spect), 51
- spect, spect.pomp-method (spect), 51
- spect-pomp (spect), 51
- spect-spect.pomp (spect), 51
- spect.match, 3
- spect.match (spect), 51
- spect.match, pomp-method (spect), 51
- spect.match, spect.pomp-method (spect), 51
- spect.match-pomp (spect), 51
- spect.match-spect.pomp (spect), 51
- spect.matched.pomp, 52
- spect.matched.pomp-class (spect), 51
- spect.matched.pomp-methods
 - (probed.pomp-methods), 43
- spect.pomp, 52
- spect.pomp-class (spect), 51
- spect.pomp-methods
 - (probed.pomp-methods), 43
- sprintf, 4
- states, 36
- states (pomp-methods), 37
- states, pomp-method (pomp-methods), 37
- states-pomp (pomp-methods), 37
- subplex, 19, 41, 52, 54, 55
- summary, probe.matched.pomp-method
 - (probed.pomp-methods), 43
- summary, probed.pomp-method
 - (probed.pomp-methods), 43
- summary, spect.matched.pomp-method
 - (probed.pomp-methods), 43
- summary, spect.pomp-method
 - (probed.pomp-methods), 43
- summary, traj.matched.pomp-method
 - (traj.match), 54
- summary-probe.matched.pomp
 - (probed.pomp-methods), 43
- summary-probed.pomp
 - (probed.pomp-methods), 43
- summary-spect.matched.pomp
 - (probed.pomp-methods), 43
- summary-spect.pomp
 - (probed.pomp-methods), 43
- summary-traj.matched.pomp (traj.match), 54
- time, 36
- time, pomp-method (pomp-methods), 37
- time-pomp (pomp-methods), 37
- time<- (pomp-methods), 37
- time<-, pomp-method (pomp-methods), 37
- time<--pomp (pomp-methods), 37
- time<-, 36
- timezero, 36
- timezero (pomp-methods), 37
- timezero, pomp-method (pomp-methods), 37
- timezero-pomp (pomp-methods), 37
- timezero<- (pomp-methods), 37

timezero<-,pomp-method (pomp-methods),
37
timezero<--pomp (pomp-methods), 37
timezero<-, 36
traj.match, 3, 36, 54, 56
traj.match,pomp-method (traj.match), 54
traj.match,traj.matched.pomp-method
(traj.match), 54
traj.match-pomp (traj.match), 54
traj.match-traj.matched.pomp
(traj.match), 54
traj.matched.pomp-class (traj.match), 54
trajectory, 3, 36, 54, 55, 56
trajectory,pomp-method (trajectory), 56
trajectory-pomp (trajectory), 56

verhulst, 57

window, 36
window,pomp-method (pomp-methods), 37
window-pomp (pomp-methods), 37