Package 'pomp'

March 21, 2014

Type Package

Title Statistical inference for partially observed Markov processes

Version 0.49-1

Date 2014-03-20

URL <http://pomp.r-forge.r-project.org>

Description Inference methods for partially-observed Markov processes

Depends R(>= 2.14.1), stats, graphics, methods, mvtnorm, subplex, nloptr, deSolve

License GPL(>= 2)

LazyData true

BuildVignettes true

MailingList
Subscribe to pomp-announce@r-forge.r-project.org for announcements by going to http://lists.r-forge.r-project.org/mailman/listinfo/pomp-announce.

Collate aaa.R authors.R generics.R version.R eulermultinom.R
plugins.R parmat.R logmeanexp.R slice-design.R profile-design.R sobol.R bsplines.R sannbox.R
pomp-fun.R pomp-class.R pomp.R pomp-methods.R
rmeasure-pomp.R rprocess-pomp.R init-state-pomp.R
dmeasure-pomp.R dprocess-pomp.R skeleton-pomp.R dprior-pomp.R rprior-pomp.R
simulate-pomp.R trajectory-pomp.R plot-pomp.R
pfilter.R pfilter-methods.R minim.R traj-match.R bsmc.R
mif-class.R particles-mif.R mif.R mif-methods.R compare-mif.R
pmcmc.R pmcmc-methods.R compare-pmcmc.R nlf-funcs.R nlf-guts.R nlf-objfun.R nlf.R
probe.R probe-match.R basic-probes.R spect.R spect-match.R abc.R abc-methods.R
builder.R example.R

R **topics documented:**

---

pomp-package   *Partially-observed Markov processes*

Description

The **pomp** package provides facilities for inference on time series data using partially-observed Markov processes (AKA state-space models or nonlinear stochastic dynamical systems). One can use **pomp** to fit nonlinear, non-Gaussian dynamic models to time-series data. The first step in using **pomp** is to encode one's model and data in an object of class pomp. One does this via a call to pomp, which involves specifying the process and measurement components of the model in one or more of a variety of ways. Details on this are given in the documentation for the pomp function and examples are given in the 'intro_to_pomp' vignette.

Currently, **pomp** provides algorithms for (i) simulation of stochastic dynamical systems (see simulate), (ii) particle filtering (AKA sequential Monte Carlo or sequential importance sampling), see pfilter), (iii) the iterated filtering method of Ionides et al. (2006), see mif, (iv) the nonlinear forecasting algorithm of Kendall et al. (2005), see nlf, (v) the particle MCMC approach of Andrieu et al. (2010), see pmcmc, (vi) basic trajectory matching, see traj.match, (vi) the probe-matching method of Wood (2010) and Kendall et al. (1999), see probe.match, (vii) a spectral probe-matching method (Reuman et al., 2006), see spect.match. See the package website http://pomp.r-forge.r-project.org for these references. The package also provides various tools for plotting and extracting information on models and data as well as an API for algorithm development. Future support for additional algorithms in envisioned, and implementations of the Bayesian sequential Monte Carlo approach of Liu \& West. Much of the work in **pomp** has been done under the auspices of a working group of the National Center for Ecological Analysis and Synthesis (NCEAS), "Inference for Mechanistic Models".

The package is provided under the GNU Public License (GPL). Contributions are welcome, as are comments, suggestions for improvements, and bug reports. See the package website http://pomp.r-forge.r-project.org for more information, access to the package mailing list, links to the authors' websites, and references to the literature.

Classes

pomp makes extensive use of S4 classes. The basic class, pomp, encodes a partially-observed Markov process together with a uni- or multi-variate data set and (possibly) parameters.

Vignettes

The vignette 'Introduction to pomp' illustrates the facilities of the package using familiar stochastic processes. Run vignette("intro_to_pomp") or look at the HTML documentation to view the vignette. Methods for accelerating your codes are discussed in the 'Advanced topics in pomp' vignette; run vignette("advanced_topics_in_pomp") to view it.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pomp](), [pfilter](), [simulate](), [trajectory](), [mif](), [nlf](), [probe.match](), [traj.match](), [bsmc](), [pmcmc]()

---

abc   *The ABC algorithm*

Description

The approximate Bayesian computation (ABC) algorithm for estimating the parameters of a partially-observed Markov process.

Usage

```
## S4 method for signature 'pomp'
abc(object, Nabc = 1, start, pars,
    rw.sd, probes, scale, epsilon,
    verbose = getOption("verbose"), transform = FALSE, ...)
## S4 method for signature 'probed.pomp'
abc(object, probes,
    verbose = getOption("verbose"), transform = FALSE, ...)
## S4 method for signature 'abc'
abc(object, Nabc, start, pars,
    rw.sd, probes, scale, epsilon,
    verbose = getOption("verbose"), transform, ...)
## S4 method for signature 'abc'
continue(object, Nabc = 1, ...)
```

Arguments

object	An object of class pomp.
Nabc	The number of ABC iterations to perform.
start	named numeric vector; the starting guess of the parameters.
pars	optional character vector naming the ordinary parameters to be estimated. Every parameter named in pars must have a positive random-walk standard deviation specified in rw.sd. Leaving pars unspecified is equivalent to setting it equal to the names of all parameters with a positive value of rw.sd.
rw.sd	numeric vector with names; used to parameterize a Gaussian random walk MCMC proposal. The random walk is only applied to parameters named in pars. The algorithm requires that the random walk be nontrivial, so each element in rw.sd[pars] must be positive. The following must be satisfied: names(rw.sd) must be a subset of names(start), rw.sd must be non-negative (zeros are simply ignored), the name of every positive element of rw.sd must be in pars.
probes	
scale	

```
epsilon
```

verbose	logical; if TRUE, print progress reports.
transform	logical; if TRUE, optimization is performed on the transformed scale.
...	Additional arguments. These are currently ignored.

Value

An object of class abc. This class inherits from class [probed.pomp](#) and contains the following additional slots:

pars, Nabc, dprior, hyperparams, transform, scale, epsilon These slots hold the values of the corresponding arguments of the call to abc.

random.walk.sd a named numeric vector containing the random-walk variances used to parameterize a Gaussian random walk MCMC proposal.

probes, conv.rec

Re-running ABC Iterations

To re-run a sequence of ABC iterations, one can use the abc method on a abc object. By default, the same parameters used for the original ABC run are re-used (except for tol, max.fail, and verbose, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing ABC Iterations

One can continue a series of ABC iterations from where one left off using the continue method. A call to abc to perform Nabc=m iterations followed by a call to continue to perform Nabc=n iterations will produce precisely the same effect as a single call to abc to perform Nabc=m+n iterations. By default, all the algorithmic parameters are the same as used in the original call to abc. Additional arguments will override the defaults.

Details

TO APPEAR.

Author(s)

Edward L. Ionides <ionides at umich dot edu>, Aaron A. King <kingaa at umich dot edu>

References

T. Toni and M. P. H. Stumpf, Simulation-based model selection for dynamical systems in systems and population biology, Bioinformatics 26:104–110, 2010.

T. Toni, D. Welch, N. Strelkowa, A. Ipsen, and M. P. H. Stumpf, Approximate Bayesian computation scheme for parameter inference and model selection in dynamical systems Journal of the Royal Society, Interface 6:187–202, 2009.

See Also

[abc-methods](#), [pomp](#), [probe](#). See the "intro_to_pomp" vignette for an example

---

abc-methods   *Methods of the "abc" class*

Description

Methods of the "abc" class.

Usage

```
## S4 method for signature 'abc'
conv.rec(object, pars, ...)
## S4 method for signature 'abc'
plot(x, y, pars, scatter = FALSE, ...)
## S4 method for signature 'abc'
dprior(object, params, log = FALSE, ...)
```

Arguments

object, x	The abc object.
pars	Names of parameters.
y	Ignored.
scatter	WHAT DOES THIS DO?
params	Named vector of parameters.
log	if TRUE, log probabilities are returned.
...	Further arguments (either ignored or passed to underlying functions).

Methods

conv.rec conv.rec(object, pars = NULL) returns the columns of the convergence-record matrix corresponding to the names in pars. By default, all rows are returned.

plot Plots a series of diagnostic plots.

Author(s)

Edward L. Ionides <ionides at umich dot edu>, Aaron A. King <kingaa at umich dot edu>

See Also

[abc](#), [pomp](#)

B-splines *B-spline bases*

Description

These functions generate B-spline basis functions. `bspline.basis` gives a basis of spline functions. `periodic.bspline.basis` gives a basis of periodic spline functions.

Usage

```
bspline.basis(x, nbasis, degree = 3, names = NULL)
periodic.bspline.basis(x, nbasis, degree = 3, period = 1, names = NULL)
```

Arguments

x	Vector at which the spline functions are to be evaluated.
nbasis	The number of basis functions to return.
degree	Degree of requested B-splines.
period	The period of the requested periodic B-splines.
names	optional; the names to be given to the basis functions. These will be the column-names of the matrix returned. If the names are specified as a format string (e.g., "basis%d"), [sprintf](#) will be used to generate the names from the column number. If a single non-format string is specified, the names will be generated by [paste](#)-ing name to the column number. One can also specify each column name explicitly by giving a length-`nbasis` string vector. By default, no column-names are given.

Details

Direct access to the underlying C routines is available. See the header file "pomp.h" for details.

Value

bspline.basis	Returns a matrix with `length(x)` rows and `nbasis` columns. Each column contains the values one of the spline basis functions.
periodic.bspline.basis	
	Returns a matrix with `length(x)` rows and `nbasis` columns. The basis functions returned are periodic with period `period`.

Author(s)

Aaron A. King <kingaa at umich dot edu>

Examples

```
x <- seq(0,2,by=0.01)
y <- bspline.basis(x,degree=3,nbasis=9,names="basis")
matplot(x,y,type='l',ylim=c(0,1.1))
lines(x,apply(y,1,sum),lwd=2)

x <- seq(-1,2,by=0.01)
y <- periodic.bspline.basis(x,nbasis=5,names="spline%d")
matplot(x,y,type='l')
```

---

basic.probes   *Some probes for partially-observed Markov processes*

Description

Several simple and configurable probes are provided in the package. These can be used directly and as examples for building custom probes.

Usage

```
probe.mean(var, trim = 0, transform = identity, na.rm = TRUE)
probe.median(var, na.rm = TRUE)
probe.var(var, transform = identity, na.rm = TRUE)
probe.sd(var, transform = identity, na.rm = TRUE)
probe.marginal(var, ref, order = 3, diff = 1, transform = identity)
probe.nlar(var, lags, powers, transform = identity)
probe.acf(var, lags, type = c("covariance", "correlation"),
          transform = identity)
probe.ccf(vars, lags, type = c("covariance", "correlation"),
          transform = identity)
probe.period(var, kernel.width, transform = identity)
probe.quantile(var, prob, transform = identity)
```

Arguments

var, vars	character; the name(s) of the observed variable(s).
trim	the fraction of observations to be trimmed (see [mean](#)).
transform	transformation to be applied to the data before the probe is computed.
na.rm	if TRUE, remove all NA observations prior to computing the probe.
kernel.width	width of modified Daniell smoothing kernel to be used in power-spectrum computation: see [kernel](#).
prob	a single probability; the quantile to compute: see [quantile](#).
lags	In probe.ccf, a vector of lags between time series. Positive lags correspond to x advanced relative to y; negative lags, to the reverse. In probe.nlar, a vector of lags present in the nonlinear autoregressive model that will be fit to the actual and simulated data. See Details, below, for a precise description.

powers	the powers of each term (corresponding to lags) in the the nonlinear autoregressive model that will be fit to the actual and simulated data. See Details, below, for a precise description.
type	Compute autocorrelation or autocovariance?
ref	empirical reference distribution. Simulated data will be regressed against the values of ref, sorted and, optionally, differenced. The resulting regression coefficients capture information about the shape of the marginal distribution. A good choice for ref is the data itself.
order	order of polynomial regression.
diff	order of differencing to perform.
...	Additional arguments to be passed through to the probe computation.

Details

Each of these functions is relatively simple. See the source code for a complete understanding of what each does.

probe.mean, probe.median, probe.var, probe.sd return functions that compute the mean, median, variance, and standard deviation of variable var, respectively.

probe.period returns a function that estimates the period of the Fourier component of the var series with largest power.

probe.marginal returns a function that regresses the marginal distribution of variable var against the reference distribution ref. If diff>0, the data and the reference distribution are first differenced diff times and centered. Polynomial regression of order order is used. This probe returns order regression coefficients (the intercept is zero).

probe.nlar returns a function that fit a nonlinear (polynomial) autoregressive model to the univariate series (variable var). Specifically, a model of the form $y_t = \sum \beta_k y_{t-\tau_k}^{p_k} + \epsilon_t$ will be fit, where τ_k are the lags and p_k are the powers. The data are first centered. This function returns the regression coefficients, β_k.

probe.acf returns a function that, if type=="covariance", computes the autocovariance of variable var at lags lags; if type=="correlation", computes the autocorrelation of variable var at lags lags.

probe.ccf returns a function that, if type=="covariance", computes the cross covariance of the two variables named in vars at lags lags; if type=="correlation", computes the cross correlation.

probe.quantile returns a function that estimates the prob-th quantile of variable var.

Value

A call to any one of these functions returns a probe function, suitable for use in probe or probe.match. That is, the function returned by each of these takes a data array (such as comes from a call to obs) as input and returns a single numerical value.

Author(s)

Daniel C. Reuman (d.reuman at imperial dot ac dot uk)

Aaron A. King (kingaa at umich dot edu)

References

B. E. Kendall, C. J. Briggs, W. M. Murdoch, P. Turchin, S. P. Ellner, E. McCauley, R. M. Nisbet, S. N. Wood Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches, Ecology, 80:1789–1805, 1999.

S. N. Wood Statistical inference for noisy nonlinear ecological dynamic systems, Nature, 466: 1102–1104, 2010.

See Also

pomp-class, pomp-methods

---

blowflies	*Model for Nicholson's blowflies.*

Description

`blowflies1` and `blowflies2` are pomp objects encoding stochastic delay-difference models.

Details

The data are from "population I", a control culture in one of A. J. Nicholson's experiments with the Australian sheep-blowfly *Lucilia cuprina*. The experiment is described on pp. 163–4 of Nicholson (1957). Unlimited quantities of larval food were provided; the adult food supply (ground liver) was constant at 0.4g per day. The data were taken from the table provided by Brillinger et al. (1980).

The models are discrete delay equations:

$$R(t+1) \sim \text{Poisson}(PN(t-\tau)\exp\left(-N(t-\tau)/N_0\right)e(t+1)\Delta t)$$

$$S(t+1) \sim \text{binomial}(N(t),\exp\left(-\delta\epsilon(t+1)\Delta t\right))$$

$$N(t) = R(t) + S(t)$$

where $e(t)$ and $\epsilon(t)$ are Gamma-distributed i.i.d. random variables with mean 1 and variances $\sigma_p^2/\Delta t$, $\sigma_d^2/\Delta t$, respectively. `blowflies1` has a timestep (Δt) of 1 day, and `blowflies2` has a timestep of 2 days. The process model in `blowflies1` thus corresponds exactly to that studied by Wood (2010). The measurement model in both cases is taken to be

$$y(t) \sim \text{negbin}(N(t), 1/\sigma_y^2)$$

, i.e., the observations are assumed to be negative-binomially distributed with mean $N(t)$ and variance $N(t) + (\sigma_y N(t))^2$.

Do

```
file.show(system.file("examples","blowflies.R",package="pomp"))
```

to view the code that constructs these pomp objects.

References

A. J. Nicholson (1957) The self-adjustment of populations to change. Cold Spring Harbor Symposia on Quantitative Biology, **22**, 153–173.

Y. Xia and H. Tong (2011) Feature Matching in Time Series Modeling. *Statistical Science* **26**, 21–46.

E. L. Ionides (2011) Discussion of "Feature Matching in Time Series Modeling" by Y. Xia and H. Tong. *Statistical Science* **26**, 49–52.

S. N. Wood (2010) Statistical inference for noisy nonlinear ecological dynamic systems. *Nature* **466**, 1102–1104.

W. S. C. Gurney, S. P. Blythe, and R. M. Nisbet (1980) Nicholson's blowflies revisited. *Nature* **287**, 17–21.

D. R. Brillinger, J. Guckenheimer, P. Guttorp and G. Oster (1980) Empirical modelling of population time series: The case of age and density dependent rates. in G. Oster (ed.), Some Questions in Mathematical Biology, vol. 13, pp. 65–90. American Mathematical Society, Providence.

See Also

pomp-class and the vignettes

Examples

```
pompExample(blowflies)
plot(blowflies1)
plot(blowflies2)
```

---

bsmc   *Liu and West Bayesian Particle Filter*

Description

Generates draws from the posterior distribution for the parameters using the Liu and West algorithm. bsmc gives draws from the posterior.

Usage

```
## S4 method for signature 'pomp'
bsmc(object, params, Np, est, smooth = 0.1,
     ntries = 1, tol = 1e-17, lower = -Inf, upper = Inf, seed = NULL,
     verbose = getOption("verbose"), max.fail = 0,
     transform = FALSE, ...)
```

Arguments

object	An object of class pomp or inheriting class pomp.
params, Np	Specifications for the prior distribution of particles. See details below.
est	Names of the rows of params that are to be estimated. No updates will be made to the other parameters. If est is not specified, all parameters for which there is variation in params will be estimated.
smooth	Kernel density smoothing parameters. The compensating shrinkage factor will be sqrt(1-smooth^2). Thus, smooth=0 means that no noise will be added to parameters. Generally, the value of smooth should be chosen close to 0 (i.e., shrink~0.1).
ntries	Number of draws from rprocess per particle used to estimate the expected value of the state process at time t+1 given the state and parameters at time t.
tol	Particles with log likelihood below tol are considered to be "lost". A filtering failure occurs when, at some time point, all particles are lost. When all particles are lost, the conditional log likelihood at that time point is set to be log(tol).
lower, upper	optional; lower and upper bounds on the priors. This is useful in case there are box constraints satisfied by the priors. The posterior is guaranteed to lie within these bounds.
seed	optional; an object specifying if and how the random number generator should be initialized ('seeded'). If seed is an integer, it is passed to set.seed prior to any simulation and is returned as the "seed" element of the return list. By default, the state of the random number generator is not changed and the value of .Random.seed on the call is stored in the "seed" element of the return list.
verbose	logical; if TRUE, print diagnostic messages.
max.fail	The maximum number of filtering failures allowed. If the number of filtering failures exceeds this number, execution will terminate with an error.
transform	logical; if TRUE, the algorithm operates on the transformed scale.
...	currently ignored.

Details

There are two ways to specify the prior distribution of particles. If params is unspecified or is a named vector, Np draws are made from the prior distribution, as specified by [rprior](). Alternatively, params can be specified as an npars x Np matrix (with rownames).

Value

An object of class "bsmcd.pomp". The "params" slot of this object will hold the parameter posterior medians. The slots of this class include:

post	A matrix containing draws from the approximate posterior distribution.
prior	A matrix containing draws from the prior distribution (identical to params on call).
eff.sample.size	
	A vector containing the effective number of particles at each time point.

smooth	The smoothing parameter used (see above).
seed	The state of the random number generator at the time `bsmc` was called. If the argument `seed` was specified, this is a copy; if not, this is the internal state of the random number generator at the time of call.
nfail	The number of filtering failures encountered.
cond.log.evidence	
	A vector containing the conditional log evidence scores at each time point.
log.evidence	The estimated log evidence.
weights	The resampling weights for each particle.

Author(s)

Michael Lavine (lavine at math dot umass dot edu), Matthew Ferrari (mferrari at psu dot edu), Aaron A. King

See Also

[pomp-class](#)

Examples

```
## See the vignettes for examples.
```

---

dacca   *Model of cholera transmission for historic Bengal.*

Description

dacca is a pomp object containing census and cholera mortality data from the Dacca district of the former British province of Bengal over the years 1891 to 1940 together with a stochastic differential equation transmission model. The model is that of King et al. (2008). The parameters are the MLE for the SIRS model with seasonal reservoir.

Data are provided courtesy of Dr. Menno J. Bouma, London School of Tropical Medicine and Hygiene.

Details

dacca is a pomp object containing the model, data, and MLE parameters. Parameters that naturally range over the positive reals are log-transformed; parameters that range over the unit interval are logit-transformed; parameters that are naturally unbounded or take integer values are not transformed.

References

King, A. A., Ionides, E. L., Pascual, M., and Bouma, M. J. Inapparent infections and cholera dynamics. Nature 454:877-880 (2008)

See Also

euler.sir, pomp

Examples

```
pompExample(dacca)
plot(dacca)
#MLEs on the natural scale
coef(dacca)
#MLEs on the transformed scale
coef(dacca,transform=TRUE)
plot(simulate(dacca))
# now change 'eps' and simulate again
coef(dacca,"eps") <- 1
plot(simulate(dacca))
```

---

eulermultinom   *Euler-multinomial death process*

Description

Density and random-deviate generation for the Euler-multinomial death process with parameters size, rate, and dt.

Usage

```
reulermultinom(n = 1, size, rate, dt)
deulermultinom(x, size, rate, dt, log = FALSE)
rgammawn(n = 1, sigma, dt)
```

Arguments

n	integer; number of random variates to generate.
size	scalar integer; number of individuals at risk.
rate	numeric vector of hazard rates.
sigma	numeric scalar; intensity of the Gamma white noise process.
dt	numeric scalar; duration of Euler step.
x	matrix or vector containing number of individuals that have succumbed to each death process.
log	logical; if TRUE, return logarithm(s) of probabilities.

Details

If N individuals face constant hazards of death in k ways at rates $r_1, r_2, \ldots, r_k$, then in an interval of duration Δt, the number of individuals remaining alive and dying in each way is multinomially distributed:

$$(N - \sum_{i=1}^{k} \Delta n_i, \Delta n_1, \ldots, \Delta n_k) \sim \text{multinomial}(N; p_0, p_1, \ldots, p_k),$$

where Δn_i is the number of individuals dying in way i over the interval, the probability of remaining alive is $p_0 = \exp(-\sum_i r_i \Delta t)$, and the probability of dying in way j is

$$p_j = \frac{r_j}{\sum_i r_i}(1 - \exp(-\sum_i r_i \Delta t)).$$

In this case, we can say that

$$(\Delta n_1, \ldots, \Delta n_k) \sim \text{eulermultinom}(N, r, \Delta t),$$

where $r = (r_1, \ldots, r_k)$. Draw m random samples from this distribution by doing

```
dn <- reulermultinom(n=m,size=N,rate=r,dt=dt),
```

where r is the vector of rates. Evaluate the probability that $x = (x_1, \ldots, x_k)$ are the numbers of individuals who have died in each of the k ways over the interval Δt =dt, by doing

```
deulermultinom(x=x,size=N,rate=r,dt=dt).
```

Breto & Ionides discuss how an infinitesimally overdispersed death process can be constructed by compounding a binomial process with a Gamma white noise process. The Euler approximation of the resulting process can be obtained as follows. Let the increments of the equidispersed process be given by

```
reulermultinom(size=N,rate=r,dt=dt).
```

In this expression, replace the rate r with $r\Delta W/\Delta t$, where $\Delta W \tilde{} Gamma(dt/\sigma^2, \sigma^2)$ is the increment of an integrated Gamma white noise process with intensity σ. That is, ΔW has mean Δt and variance $\sigma^2 \Delta t$. The resulting process is overdispersed and converges (as Δt goes to zero) to a well-defined process. The following lines of R code accomplish this:

```
dW <- rgammawn(sigma=sigma,dt=dt)
```

```
dn <- reulermultinom(size=N,rate=r,dt=dW)
```

or

```
dn <- reulermultinom(size=N,rate=r*dW/dt,dt=dt).
```

He et al. use such overdispersed death processes in modeling measles.

For all of the functions described here, direct access to the underlying C routines is available: see the header file "pomp.h", included with the package.

Value

reulermultinom Returns a length(rate) by n matrix. Each column is a different random draw. Each row contains the numbers of individuals succumbed to the corresponding process.

deulermultinom Returns a vector (of length equal to the number of columns of x) containing the probabilities of observing each column of x given the specified parameters (size, rate, dt).

rgammawn Returns a vector of length n containing random increments of the integrated Gamma white noise process with intensity sigma.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

C. Bret\'o & E. L. Ionides, Compound Markov counting processes and their applications to modeling infinitesimally over-dispersed systems. Stoch. Proc. Appl., 121:2571–2591, 2011.

D. He, E. L. Ionides, & A. A. King, Plug-and-play inference for disease dynamics: measles in large and small populations as a case study. J. R. Soc. Interface, 7:271–283, 2010.

Examples

```
print(dn <- reulermultinom(5,size=100,rate=c(a=1,b=2,c=3),dt=0.1))
deulermultinom(x=dn,size=100,rate=c(1,2,3),dt=0.1)
## an Euler-multinomial with overdispersed transitions:
dt <- 0.1
dW <- rgammawn(sigma=0.1,dt=dt)
print(dn <- reulermultinom(5,size=100,rate=c(a=1,b=2,c=3),dt=dW))
```

---

gompertz *Gompertz model with log-normal observations.*

Description

gompertz is a pomp object encoding a stochastic Gompertz population model with log-normal measurement error.

Details

The state process is $X_{t+1} = K^{1-S}X_t^S\epsilon_t$, where $S = e^{-r}$ and the ϵ_t are i.i.d. lognormal random deviates with variance σ^2. The observed variables Y_t are distributed as $\mathrm{lognormal}(\log X_t, \tau)$. Parameters include the per-capita growth rate r, the carrying capacity K, the process noise s.d. σ, the measurement error s.d. τ, and the initial condition X_0. The pomp object includes parameter transformations that log-transform the parameters for estimation purposes.

See Also

pomp-class and the introductory vignette vignette("intro_to_pomp").

Examples

```
pompExample(gompertz)
plot(gompertz)
coef(gompertz)
coef(gompertz,transform=TRUE)
```

---

logmeanexp   *The log-mean-exp trick*

Description

`logmeanexp` computes the log-mean-exp of a set of numbers.

Usage

```
logmeanexp(x, se = FALSE)
```

Arguments

x	numeric
se	logical; give approximate standard error?

Value

`log(mean(exp(x)))` computed so as to avoid over- or underflow. If `se = FALSE`, the approximate standard error is returned as well.

Author(s)

Aaron A. King <kingaa at umich dot edu>

Examples

```
## generate a bifurcation diagram for the Ricker map
pompExample(ricker)
ll <- replicate(n=5,logLik(pfilter(ricker,Np=1000)))
## an estimate of the log likelihood:
logmeanexp(ll)
logmeanexp(ll,se=TRUE)
```

LondonYorke   *Historical childhood disease incidence data*

Description

LondonYorke is a data-frame containing the monthly number of reported cases of chickenpox, measles, and mumps from two American cities (Baltimore and New York) in the mid-20th century (1928–1972).

Usage

```
data(LondonYorke)
```

References

W. P. London and J. A. Yorke, Recurrent Outbreaks of Measles, Chickenpox and Mumps: I. Seasonal Variation in Contact Rates, American Journal of Epidemiology, 98:453–468, 1973.

See Also

[pomp-class](#) and the vignettes

Examples

```
data(LondonYorke)

plot(cases~time,data=LondonYorke,subset=disease=="measles",type='n',main="measles",bty='l')
lines(cases~time,data=LondonYorke,subset=disease=="measles"&town=="Baltimore",col="red")
lines(cases~time,data=LondonYorke,subset=disease=="measles"&town=="New York",col="blue")
legend("topright",legend=c("Baltimore","New York"),lty=1,col=c("red","blue"),bty='n')

plot(
    cases~time,
    data=LondonYorke,
    subset=disease=="chickenpox"&town=="New York",
    type='l',col="blue",main="chickenpox, New York",
    bty='l'
    )

plot(
    cases~time,
    data=LondonYorke,
    subset=disease=="mumps"&town=="New York",
    type='l',col="blue",main="mumps, New York",
    bty='l'
    )
```

---

mif   *The MIF algorithm*

Description

The MIF algorithm for estimating the parameters of a partially-observed Markov process.

Usage

```
mif(object, ...)
## S4 method for signature 'pomp'
mif(object, Nmif = 1, start, pars, ivps = character(0),
    particles, rw.sd, Np, ic.lag, var.factor,
    cooling.type, cooling.fraction, cooling.factor,
    method = c("mif","unweighted","fp","mif2"),
    tol = 1e-17, max.fail = Inf,
    verbose = getOption("verbose"), transform = FALSE, ...)
## S4 method for signature 'pfilterd.pomp'
mif(object, Nmif = 1, Np, tol, ...)
## S4 method for signature 'mif'
mif(object, Nmif, start, ivps,
    particles, rw.sd, Np, ic.lag, var.factor,
    cooling.type, cooling.fraction,
    method, tol, transform, ...)
## S4 method for signature 'mif'
continue(object, Nmif = 1, ...)
```

Arguments

object	An object of class pomp.
Nmif	The number of MIF iterations to perform.
start	named numerical vector; the starting guess of the parameters.
pars	optional character vector naming the ordinary parameters to be estimated. Every parameter named in pars must have a positive random-walk standard deviation specified in rw.sd. Leaving pars unspecified is equivalent to setting it equal to the names of all parameters with a positive value of rw.sd that are not ivps.
ivps	optional character vector naming the initial-value parameters (IVPs) to be estimated. Every parameter named in ivps must have a positive random-walk standard deviation specified in rw.sd. If pars is empty, i.e., only IVPs are to be estimated, see below ""Using MIF to estimate initial-value parameters only"".
particles	Function of prototype particles(Np,center,sd,...) which sets up the starting particle matrix by drawing a sample of size Np from the starting particle distribution centered at center and of width sd. If particles is not supplied by the user, the default behavior is to draw the particles from a multivariate normal distribution with mean center and standard deviation sd.

rw.sd   numeric vector with names; the intensity of the random walk to be applied to parameters. The random walk is only applied to parameters named in `pars` (i.e., not to those named in `ivps`). The algorithm requires that the random walk be nontrivial, so each element in `rw.sd[pars]` must be positive. `rw.sd` is also used to scale the initial-value parameters (via the `particles` function). Therefore, each element of `rw.sd[ivps]` must be positive. The following must be satisfied: `names(rw.sd)` must be a subset of `names(start)`, `rw.sd` must be non-negative (zeros are simply ignored), the name of every positive element of `rw.sd` must be in either `pars` or `ivps`.

Np   the number of particles to use in filtering. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timestep, one may specify Np either as a vector of positive integers (of length `length(time(object,t0=TRUE))`) or as a function taking a positive integer argument. In the latter case, `Np(k)` must be a single positive integer, representing the number of particles to be used at the k-th timestep: `Np(0)` is the number of particles to use going from `timezero(object)` to `time(object)[1]`, `Np(1)`, from `timezero(object)` to `time(object)[1]`, and so on, while when `T=length(time(object,t0=TRUE))`, `Np(T)` is the number of particles to sample at the end of the time-series.

ic.lag   a positive integer; the timepoint for fixed-lag smoothing of initial-value parameters. The mif update for initial-value parameters consists of replacing them by their filtering mean at time `times[ic.lag]`, where `times=time(object)`. It makes no sense to set `ic.lag>length(times)`; if it is so set, `ic.lag` is set to `length(times)` with a warning. For `method='mif2'`, the default is `ic.lag=length(times)`.

var.factor   a positive number; the scaling coefficient relating the width of the starting particle distribution to `rw.sd`. In particular, the width of the distribution of particles at the start of the first MIF iteration will be `random.walk.sd*var.factor`.

cooling.type, cooling.fraction, cooling.factor

     specifications for the cooling schedule, i.e., the manner in which the intensity of the parameter perturbations is reduced with successive filtering iterations. `cooling.type` specifies the nature of the cooling schedule. When `cooling.type="geometric"`, on the n-th MIF iteration, the relative perturbation intensity is `cooling.fraction^(n/50)`. When `cooling.type="hyperbolic"`, on the n-th MIF iteration, the relative perturbation intensity is `(s+1)(s+n)`, where `(s+1)/(s+50)=cooling.fraction`. `cooling.fraction` is the relative magnitude of the parameter perturbations after 50 MIF iterations. `cooling.factor` is now deprecated: to achieve the old behavior, use `cooling.type="geometric"` and `cooling.fraction=(cooling.factor)^50`.

method   `method` sets the update rule used in the algorithm. `method="mif"` uses the iterated filtering update rule (Ionides 2006, 2011); `method="unweighted"` updates the parameter to the unweighted average of the filtering means of the parameters at each time; `method="fp"` updates the parameter to the filtering mean at the end of the time series. `method="mif2"` [EXPLAIN]

tol   See the description under pfilter.

max.fail   See the description under pfilter.

verbose   logical; if TRUE, print progress reports.

| transform | logical; if `TRUE`, optimization is performed on the transformed scale. |
| `...` | additional arguments that override the defaults. |

Re-running MIF Iterations

To re-run a sequence of MIF iterations, one can use the `mif` method on a `mif` object. By default, the same parameters used for the original MIF run are re-used (except for `weighted`, `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing MIF Iterations

One can resume a series of MIF iterations from where one left off using the `continue` method. A call to `mif` to perform `Nmif=m` iterations followed by a call to `continue` to perform `Nmif=n` iterations will produce precisely the same effect as a single call to `mif` to perform `Nmif=m+n` iterations. By default, all the algorithmic parameters are the same as used in the original call to `mif`. Additional arguments will override the defaults.

Using MIF to estimate initial-value parameters only

One can use MIF's fixed-lag smoothing to estimate only initial value parameters (IVPs). In this case, `pars` is left empty and the IVPs to be estimated are named in `ivps`. If `theta` is the current parameter vector, then at each MIF iteration, `Np` particles are drawn from a distribution centered at `theta` and with width proportional to `var.factor*rw.sd`, a particle filtering operation is performed, and `theta` is replaced by the filtering mean at `time(object)[ic.lag]`. Note the implication that, when `mif` is used in this way on a time series any longer than `ic.lag`, unnecessary work is done. If the time series in `object` is longer than `ic.lag`, consider replacing `object` with `window(object,end=ic.lag)`.

Details

If `particles` is not specified, the default behavior is to draw the particles from a multivariate normal distribution. **It is the user's responsibility to ensure that, if the optional `particles` argument is given, that the `particles` function satisfies the following conditions:**

`particles` has at least the following arguments: `Np`, `center`, `sd`, and `...`. `Np` may be assumed to be a positive integer; `center` and `sd` will be named vectors of the same length. Additional arguments may be specified; these will be filled with the elements of the `userdata` slot of the underlying pomp object (see [pomp-class](#)).

`particles` returns a `length(center)` x `Np` matrix with rownames matching the names of `center` and `sd`. Each column represents a distinct particle.

The center of the particle distribution returned by `particles` should be `center`. The width of the particle distribution should vary monotonically with `sd`. In particular, when `sd=0`, the `particles` should return matrices with `Np` identical columns, each given by the parameters specified in `center`.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

E. L. Ionides, C. Bret\'o, & A. A. King, Inference for nonlinear dynamical systems, Proc. Natl. Acad. Sci. U.S.A., 103:18438–18443, 2006.

E. L. Ionides, A. Bhadra, Y. Atchad\'e, & A. A. King, Iterated filtering, Annals of Statistics, 39:1776–1802, 2011.

A. A. King, E. L. Ionides, M. Pascual, and M. J. Bouma, Inapparent infections and cholera dynamics, Nature, 454:877–880, 2008.

See Also

mif-methods, pomp, pomp-class, pfilter. See the "intro_to_pomp" vignette for examples.

---

mif-methods   *Methods of the "mif" class*

Description

Methods of the "mif" class.

Usage

```
## S4 method for signature 'mif'
logLik(object, ...)
## S4 method for signature 'mif'
conv.rec(object, pars, transform = FALSE, ...)
## S4 method for signature 'mif'
plot(x, y = NULL, ...)
compare.mif(z)
```

Arguments

object	The mif object.
pars	Names of parameters.
x	The mif object.
y	Ignored.
z	A mif object or list of mif objects.
transform	optional logical; should the parameter transformations be applied? See coef for details.
...	Further arguments (either ignored or passed to underlying functions).

Methods

conv.rec conv.rec(object, pars = NULL) returns the columns of the convergence-record matrix corresponding to the names in pars. By default, all rows are returned.

logLik Returns the value in the loglik slot.

mif Re-runs the MIF iterations. See the documentation for [mif](#).

compare.mif Given a mif object or a list of mif objects, compare.mif produces a set of diagnostic plots.

plot Plots a series of diagnostic plots. When x is a mif object, plot(x) is equivalent to compare.mif(list(x)).

predvarplot predvarplot(object, pars = NULL, mean = FALSE, ...) produces a plot of the scaled prediction variances for each parameter. This can be used to diagnose a good value of the mif parameters var.factor and ic.lag. If used in this way, one should run mif with Nmif=1 first. Additional arguments in ... will be passed to the actual plotting function.

print Prints a summary of the mif object.

show Displays the mif object.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

E. L. Ionides, C. Bret\'o, & A. A. King, Inference for nonlinear dynamical systems, Proc. Natl. Acad. Sci. U.S.A., 103:18438–18443, 2006.

A. A. King, E. L. Ionides, M. Pascual, and M. J. Bouma, Inapparent infections and cholera dynamics, Nature, 454:877–880, 2008.

See Also

[mif](#), [pomp](#), [pomp-class](#), [pfilter](#)

---

nlf   *Fit Model to Data Using Nonlinear Forecasting (NLF)*

Description

Calls an optimizer to maximize the nonlinear forecasting (NLF) goodness of fit, by simulating data from a model, fitting a nonlinear autoregressive model to the simulated time series (which may be multivariate) and using the fitted model to predict some or all variables in the data time series. NLF is an 'indirect inference' method using a quasi-likelihood as the objective function.

Usage

```
nlf(object, start, est, lags, period = NA, tensor = FALSE,
    nconverge=1000, nasymp=1000, seed = 1066,
    transform = identity,
    nrbf = 4, method = "subplex", skip.se = FALSE,
    verbose = FALSE, gr = NULL,
    bootstrap=FALSE, bootsamp = NULL,
    lql.frac = 0.1, se.par.frac = 0.1, eval.only = FALSE,
    transform.params = FALSE, ...)
```

Arguments

object	A pomp object, with the data and model to fit to it.
start	Named numeric vector with guessed parameters.
est	Vector containing the names or indices of parameters to be estimated.
lags	A vector specifying the lags to use when constructing the nonlinear autoregressive prediction model. The first lag is the prediction interval.
period	numeric; period=NA means the model is nonseasonal. period>0 is the period of seasonal forcing in 'real time'.
tensor	logical; if FALSE, the fitted model is a generalized additive model with time mod period as one of the predictors, i.e., a gam with time-varying intercept. If TRUE, the fitted model is a gam with lagged state variables as predictors and time-periodic coefficients, constructed using tensor products of basis functions of state variables with basis functions of time.
nconverge	Number of convergence timesteps to be discarded from the model simulation.
nasymp	Number of asymptotic timesteps to be recorded from the model simulation.
seed	Integer specifying the random number seed to use. When fitting, it is usually best to always run the simulations with the same sequence of random numbers, which is accomplished by setting seed to an integer. If you want a truly random simulation, set seed=NULL.
transform	optional function. If specified, forecasting is performed using data and model simulations transformed by this function. By default, transform is the identity function, i.e., no transformation is performed. The main purpose of transform is to achieve approximately multivariate normal forecasting errors. If data are univariate, transform should take a scalar and return a scalar. If data are multivariate, transform should assume a vector input and return a vector of the same length.
nrbf	A scalar specifying the number of radial basis functions to be used at each lag.
method	Optimization method. Choices are subplex and any of the methods used by optim.
skip.se	Logical; if TRUE, skip the computation of standard errors.
verbose	Logical; if TRUE, the negative log quasilikelihood and parameter values are printed at each iteration of the optimizer.
gr	optional; passed to optim if optim is used.

bootstrap	Logical; if TRUE the indices in bootsamp will determine which of the conditional likelihood values be used in computing the quasi-loglikelihood.
bootsamp	Vector of integers; used to have the quasi-loglikelihood evaluated using a bootstrap re-sampling of the data set.
lql.frac	target fractional change in log quasi-likelihood for quadratic standard error estimate
se.par.frac	initial parameter-change fraction for quadratic standard error estimate
eval.only	logical; if TRUE, no optimization is attempted and the quasi-loglikelihood value is evaluated at the start parameters.
transform.params	
	logical; if TRUE, parameters are optimized on the transformed scale.
...	Arguments that will be passed to optim or subplex in the control list.

Details

This is functionally a wrapper for nlf.objfun, which does the statistical heavy lifting and should be consulted for details.

Value

A list corresponding to the output from the optimizer, except that the full parameter vector is returned (not just the ones fitted), the log quasilikelihood (LQL) (*not* -LQL) is reported, xstart is included, and asymptotic Wald standard errors based on M-estimator theory are returned for each fitted parameter.

Author(s)

Stephen P. Ellner <spe2 at cornell dot edu> and Bruce E. Kendall <kendall at bren dot ucsb dot edu>

References

The following papers describe and motivate the NLF approach to model fitting:

Ellner, S. P., Bailey, B. A., Bobashev, G. V., Gallant, A. R., Grenfell, B. T. and Nychka D. W. (1998) Noise and nonlinearity in measles epidemics: combining mechanistic and statistical approaches to population modeling. *American Naturalist* **151**, 425–440.

Kendall, B. E., Briggs, C. J., Murdoch, W. W., Turchin, P., Ellner, S. P., McCauley, E., Nisbet, R. M. and Wood S. N. (1999) Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches. *Ecology* **80**, 1789–1805. Available online at http://www2.bren.ucsb.edu/~kendall/pubs/1999Ecology.pdf

Kendall, B. E., Ellner, S. P., McCauley, E., Wood, S. N., Briggs, C. J., Murdoch, W. W. and Turchin, P. (2005) Population cycles in the pine looper moth (*Bupalus piniarius*): dynamical tests of mechanistic hypotheses. *Ecological Monographs* **75**, 259–276. Available online at http://repositories.cdlib.org/postprints/818/

---

ou2   *Two-dimensional discrete-time Ornstein-Uhlenbeck process*

Description

ou2 is a pomp object encoding a bivariate discrete-time Ornstein-Uhlenbeck process.

Details

If the state process is $X(t) = (x_1(t), x_2(t))$, then

$$X(t + 1) = \alpha X(t) + \sigma \epsilon(t),$$

where α and σ are 2x2 matrices, σ is lower-triangular, and $\epsilon(t)$ is standard bivariate normal. The observation process is $Y(t) = (y_1(t), y_2(t))$, where $y_i(t) \sim \mathrm{normal}(x_i(t), \tau)$. The functions rprocess, dprocess, rmeasure, dmeasure, and skeleton are implemented using compiled C code for computational speed: see the source code for details. This object is demonstrated in the vignette "Advanced topics in pomp".

See Also

pomp and the vignettes

Examples

```
pompExample(ou2)
plot(ou2)
coef(ou2)
x <- simulate(ou2)
plot(x)
pf <- pfilter(ou2,Np=1000)
logLik(pf)
```

---

parmat   *Create a matrix of parameters*

Description

parmat is a utility that makes a vector of parameters suitable for use in **pomp** functions.

Usage

```
parmat(params, nrep = 1)
```

Arguments

params	named numeric vector or matrix of parameters.
nrep	number of replicates (columns) desired.

Value

parmat returns a matrix consisting of nrep copies of params.

Author(s)

Aaron A. King <kingaa at umich dot edu>

Examples

```
## generate a bifurcation diagram for the Ricker map
pompExample(ricker)
p <- parmat(coef(ricker),nrep=500)
p["r",] <- exp(seq(from=1.5,to=4,length=500))
x <- trajectory(ricker,times=seq(from=1000,to=2000,by=1),params=p)
matplot(p["r",],x["N",,],pch='.',col='black',xlab="log(r)",ylab="N",log='x')
```

---

pfilter   *Particle filter*

Description

Run a plain vanilla particle filter. Resampling is performed at each observation.

Usage

```
## S4 method for signature 'pomp'
pfilter(object, params, Np, tol = 1e-17,
    max.fail = Inf, pred.mean = FALSE, pred.var = FALSE,
    filter.mean = FALSE, save.states = FALSE,
    save.params = FALSE, seed = NULL,
    verbose = getOption("verbose"), ...)
## S4 method for signature 'pfilterd.pomp'
pfilter(object, params, Np, tol, ...)
```

Arguments

object	An object of class pomp or inheriting class pomp.
params	A npars x Np numeric matrix containing the parameters corresponding to the initial state values in xstart. This must have a 'rownames' attribute. If it desired that all particles should share the same parameter values, one one may supply params as a named numeric vector.
Np	the number of particles to use. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timesteps, one may specify Np either as a vector of positive integers (of length length(time(object,t0=TRUE))) or as a function taking a positive integer argument. In the latter case, Np(k) must be a single positive integer, representing the number of particles to be

used at the k-th timestep: Np(0) is the number of particles to use going from
timezero(object) to time(object)[1], Np(1), from timezero(object) to
time(object)[1], and so on, while when T=length(time(object,t0=TRUE)),
Np(T) is the number of particles to sample at the end of the time-series. When
object is of class mif, this is by default the same number of particles used in
the mif iterations.

tol   positive numeric scalar; particles with likelihood less than tol are considered to
               be "lost". A filtering failure occurs when, at some time point, all particles are
               lost. When all particles are lost, the conditional likelihood at that time point is
               set to tol.

max.fail   integer; the maximum number of filtering failures allowed. If the number of
               filtering failures exceeds this number, execution will terminate with an error. By
               default, max.fail is set to infinity, so no error can be triggered.

pred.mean   logical; if TRUE, the prediction means are calculated for the state variables and
               parameters.

pred.var   logical; if TRUE, the prediction variances are calculated for the state variables
               and parameters.

filter.mean   logical; if TRUE, the filtering means are calculated for the state variables and
               parameters.

save.states, save.params

               logical. If save.states=TRUE, the state-vector for each particle at each time
               is saved in the saved.states slot of the returned pfilterd.pomp object. If
               save.params=TRUE, the parameter-vector for each particle at each time is saved
               in the saved.params slot of the returned pfilterd.pomp object.

seed   optional; an object specifying if and how the random number generator should
               be initialized ('seeded'). If seed is an integer, it is passed to set.seed prior
               to any simulation and is returned as the "seed" element of the return list. By
               default, the state of the random number generator is not changed and the value
               of .Random.seed on the call is stored in the "seed" element of the return list.

verbose   logical; if TRUE, progress information is reported as pfilter works.

...   By default, when pfilter pfilter is run on a pfilterd.pomp object, the set-
               tings in the original call are re-used. This default behavior can be overridden by
               changing the settings (see Examples below).

Value

An object of class pfilterd.pomp. This class inherits from class pomp and contains the following
additional slots:

pred.mean, pred.var, filter.mean matrices of prediction means, variances, and filter means, re-
spectively. In each of these, the rows correspond to states and parameters (if appropriate), in
that order, the columns to successive observations in the time series contained in object.

eff.sample.size numeric vector containing the effective number of particles at each time point.

cond.loglik numeric vector containing the conditional log likelihoods at each time point.

saved.states If pfilter was called with save.states=TRUE, this is the list of state-vectors at each
time point, for each particle. It is a length-ntimes list of nvars-by-Np arrays. In particular,
saved.states[[t]][,i] can be considered a sample from $f[X_t|y_{1:t}]$.

saved.params If `pfilter` was called with `save.params=TRUE`, this is the list of parameter-vectors at each time point, for each particle. It is a length-`ntimes` list of `npars`-by-`Np` arrays. In particular, `saved.params[[t]][,i]` is the parameter portion of the i-th particle at time t.

seed the state of the random number generator at the time `pfilter` was called. If the argument `seed` was specified, this is a copy; if not, this is the internal state of the random number generator at the time of call.

Np, tol, nfail the number of particles used, failure tolerance, and number of filtering failures, respectively.

loglik the estimated log-likelihood.

These can be accessed using the `$` operator as if the returned object were a list. In addition, `logLik` returns the log likelihood. Note that if the argument `params` is a named vector, then these parameters are included in the `params` slot of the returned `pfilterd.pomp` object. That is `coef(pfilter(obj,params=theta))==theta` if `theta` is a named vector of parameters.

Author(s)

Aaron A. King `<kingaa at umich dot edu>`

References

M. S. Arulampalam, S. Maskell, N. Gordon, & T. Clapp. A Tutorial on Particle Filters for Online Nonlinear, Non-Gaussian Bayesian Tracking. IEEE Trans. Sig. Proc. 50:174–188, 2002.

See Also

pomp-class

Examples

```
pompExample(gompertz)
pf <- pfilter(gompertz,Np=1000) ## use 1000 particles
logLik(pf)
cond.logLik(pf) ## conditional log-likelihoods
eff.sample.size(pf)
logLik(pfilter(gompertz))        ## run it again with 1000 particles
logLik(pfilter(gompertz,Np=2000))    ## run it with 2000 particles and get the filtering means

## See the vignettes for more examples.
```

---

pfilter-methods	*Methods of the "pfilterd.pomp" class*

Description

Methods of the "pfilterd.pomp" class.

Usage

```
## S4 method for signature 'pfilterd.pomp'
logLik(object, ...)
## S4 method for signature 'pfilterd.pomp'
pred.mean(object, pars, ...)
## S4 method for signature 'pfilterd.pomp'
pred.var(object, pars, ...)
## S4 method for signature 'pfilterd.pomp'
filter.mean(object, pars, ...)
## S4 method for signature 'pfilterd.pomp'
eff.sample.size(object, ...)
## S4 method for signature 'pfilterd.pomp'
cond.logLik(object, ...)
## S4 method for signature 'pfilterd.pomp'
as(object, class)
## S4 method for signature 'pfilterd.pomp,data.frame'
coerce(from, to = "data.frame", strict = TRUE)
## S3 method for class 'pfilterd.pomp'
as.data.frame(x, row.names, optional, ...)
```

Arguments

`object, x`	An object of class `pfilterd.pomp` or inheriting class `pfilterd.pomp`.
`pars`	Names of parameters.
`class`	character; name of the class to which `object` should be coerced.
`from, to`	the classes between which coercion should be performed.
`strict`	ignored.
`row.names, optional`	ignored.
`...`	Additional arguments unused at present.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pfilter](#), [pomp-class](#)

plugins	*Plug-ins for dynamical models based on stochastic Euler algorithms*

Description

Plug-in facilities for implementing discrete-time Markov processes and continuous-time Markov processes using the Euler algorithm. These can be used in the rprocess and dprocess slots of pomp.

Usage

```
onestep.sim(step.fun, PACKAGE)
euler.sim(step.fun, delta.t, PACKAGE)
discrete.time.sim(step.fun, delta.t = 1, PACKAGE)
gillespie.sim(rate.fun, v, d, PACKAGE)
onestep.dens(dens.fun, PACKAGE)
```

Arguments

step.fun	This can be either an R function or the name of a compiled, dynamically loaded native function containing the model simulator. It should be written to take a single Euler step from a single point in state space. If it is an R function, it should be of the form step.fun(x,t,params,delta.t,...). Here, x is a named numeric vector containing the value of the state process at time t, params is a named numeric vector containing parameters, and delta.t is the length of the Euler time-step. If step.fun is the name of a native function, it must be of type "pomp_onestep_sim" as defined in the header "pomp.h", which is included with the **pomp** package. For details on how to write such codes, see Details.
rate.fun	This can be either an R function or the name of a compiled, dynamically loaded native function that computes the transition rates. If it is an R function, it should be of the form rate.fun(j,x,t,params,...). Here, j is the number of the event, x is a named numeric vector containing the value of the state process at time t and params is a named numeric vector containing parameters. If rate.fun is a native function, it must be of type "pomp_ssa_rate_fn" as defined in the header "pomp.h", which is included with the package. For details on how to write such codes, see Details.
v, d	Matrices that specify the continuous-time Markov process in terms of its elementary events. Each should have dimensions nvar x nevent, where nvar is the number of state variables and nevent is the number of elementary events. v describes the changes that occur in each elementary event: it will usually comprise the values 1, -1, and 0 according to whether a state variable is incremented, decremented, or unchanged in an elementary event. d is a binary matrix that describes the dependencies of elementary event rates on state variables: d[i,j] will have value 1 if event rate j must be updated as a result of a change in state variable i and 0 otherwise

dens.fun   This can be either an R function or a compiled, dynamically loaded native func-
                   tion containing the model transition log probability density function. If it is
                   an R function, it should be of the form dens.fun(x1,x2,t1,t2,params,...).
                   Here, x1 and x2 are named numeric vectors containing the values of the state
                   process at times t1 and t2, params is a named numeric vector containing pa-
                   rameters. If dens.fun is the name of a native function, it should be of type
                   "pomp_onestep_pdf" as defined in the header "pomp.h", which is included with
                   the **pomp** package. This function should return the log likelihood of a transition
                   from x1 at time t1 to x2 at time t2, assuming that no intervening transitions
                   have occurred. For details on how to write such codes, see Details.

delta.t   Size of Euler time-steps.

PACKAGE   an optional argument that specifies to which dynamically loaded library we re-
                   strict the search for the native routines. If this is "base", we search in the R
                   executable itself.

Details

onestep.sim is the appropriate choice when it is possible to simulate the change in state from
one time to another, regardless of how large the interval between them is. To use onestep.sim,
you must write a function step.fun that will advance the state process from one arbitrary time to
another. euler.sim is appropriate when one cannot do this but can compute the change in state
via a sequence of smaller steps. This is desirable, for example, if one is simulating a continuous
time process but is willing to approximate it using an Euler approach. discrete.time.sim is
appropriate when the process evolves in discrete time. In this case, by default, the intervals between
observations are integers.

To use euler.sim or discrete.time.sim, you must write a function step.fun that will take a
single Euler step, of size at most delta.t. euler.sim and discrete.time.sim will create simu-
lators that take as many steps as needed to get from one time to another. See below for information
on how euler.sim chooses the actual step size it uses.

gillespie.sim allows exact simulation of a continuous-time, discrete-state Markov process using
Gillespie's algorithm. This is an "event-driven" approach: correspondingly, to use gillespie.sim,
you must write a function rate.fun that computes the rates of each elementary event and specify
two matrices (d,v) that describe, respectively, the dependencies of each rate and the consequences
of each event.

onestep.dens will generate a suitable dprocess function when one can compute the likelihood
of a given state transition simply by knowing the states at two times under the assumption that
the state has not changed between the times. This is typically possible, for instance, when the
rprocess function is implemented using onestep.sim, euler.sim, or discrete.time.sim. [NB:
currently, there are no high-level algorithms in **pomp** that use dprocess. This function is provided
for completeness only, and with an eye toward future development.]

If step.fun is written as an R function, it must have at least the arguments x, t, params, delta.t,
and .... On a call to this function, x will be a named vector of state variables, t a scalar time, and
params a named vector of parameters. The length of the Euler step will be delta.t. If the argument
covars is included and a covariate table has been included in the pomp object, then on a call to this
function, covars will be filled with the values, at time t, of the covariates. This is accomplished via
interpolation of the covariate table. Additional arguments may be given: these will be filled by the
correspondingly-named elements in the userdata slot of the pomp object (see [pomp](#)). If step.fun

is written in a native language, it must be a function of type "pomp_onestep_sim" as specified in the header "pomp.h" included with the package (see the directory "include" in the installed package directory).

If `rate.fun` is written as an R function, it must have at least the arguments `j`, `x`, `t`, `params`, and `....` Here, `j` is the an integer that indicates which specific elementary event we desire the rate of. `x` is a named vector containing the value of the state process at time `t`, and `params` is a named vector containing parameters. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t`, of the covariates. This is accomplished via interpolation of the covariate table. If `rate.fun` is a native function, it must be of type "pomp_ssa_rate_fn" as defined in the header "pomp.h", which is included with the package.

In writing `dens.fun`, you must assume that no state transitions have occurred between `t1` and `t2`. If `dens.fun` is written as an R function, it must have at least the arguments `x1`, `x2`, `t1`, `t2`, `params`, and `....` On a call to this function, `x1` and `x2` will be named vectors of state variables at times `t1` and `t2`, respectively. The named vector `params` contains the parameters. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t1`, of the covariates. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t1`, of the covariates. This is accomplished via interpolation of the covariate table. As above, any additional arguments will be filled by the correspondingly-named elements in the `userdata` slot of the `pomp` object (see [pomp](#)). If `dens.fun` is written in a native language, it must be a function of type "pomp_onestep_pdf" as defined in the header "pomp.h" included with the package (see the directory "include" in the installed package directory).

Value

`onestep.sim`, `euler.sim`, `discrete.time.sim`, and `gillespie.sim` each return functions suitable for use as the argument `rprocess` argument in [pomp](#).

`onestep.dens` returns a function suitable for use as the argument `dprocess` in [pomp](#).

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[eulermultinom](#), [pomp](#)

Examples

```
## examples showing how to use these functions
## are provided in the vignette "intro_to_pomp"
## Not run:
vignette("intro_to_pomp")

## End(Not run)
```

---

pmcmc   *The PMCMC algorithm*

Description

The Particle MCMC algorithm for estimating the parameters of a partially-observed Markov process.

Usage

```
## S4 method for signature 'pomp'
pmcmc(object, Nmcmc = 1, start, pars,
    rw.sd, Np, tol = 1e-17, max.fail = 0,
    verbose = getOption("verbose"), transform = FALSE, ...)
## S4 method for signature 'pfilterd.pomp'
pmcmc(object, Nmcmc = 1, Np, tol, ...)
## S4 method for signature 'pmcmc'
pmcmc(object, Nmcmc, start, pars,
    rw.sd, Np, tol, max.fail = 0,
    verbose = getOption("verbose"), transform, ...)
## S4 method for signature 'pmcmc'
continue(object, Nmcmc = 1, ...)
```

Arguments

object	An object of class pomp.
Nmcmc	The number of PMCMC iterations to perform.
start	named numeric vector; the starting guess of the parameters.
pars	optional character vector naming the ordinary parameters to be estimated. Every parameter named in pars must have a positive random-walk standard deviation specified in rw.sd. Leaving pars unspecified is equivalent to setting it equal to the names of all parameters with a positive value of rw.sd.
rw.sd	numeric vector with names; used to parameterize a Gaussian random walk MCMC proposal. The random walk is only applied to parameters named in pars. The algorithm requires that the random walk be nontrivial, so each element in rw.sd[pars] must be positive. The following must be satisfied: names(rw.sd) must be a subset of names(start), rw.sd must be non-negative (zeros are simply ignored), the name of every positive element of rw.sd must be in pars.
Np	a positive integer; the number of particles to use in each filtering operation.
tol	numeric scalar; particles with log likelihood below tol are considered to be "lost". A filtering failure occurs when, at some time point, all particles are lost.
max.fail	integer; maximum number of filtering failures permitted. If the number of failures exceeds this number, execution will terminate with an error.
verbose	logical; if TRUE, print progress reports.
transform	logical; if TRUE, optimization is performed on the transformed scale.
...	Additional arguments. These are currently ignored.

Value

An object of class pmcmc. This class inherits from class [`pfilterd.pomp`](pfilterd.pomp) and contains the following additional slots:

pars, Nmcmc, dprior, hyperparams These slots hold the values of the corresponding arguments of the call to pmcmc.

random.walk.sd a named numeric vector containing the random-walk variances used to parameterize a Gaussian random walk MCMC proposal.

log.prior a numeric value containing the log of the prior density evaluated at the parameter vector in the params slot.

Re-running PMCMC Iterations

To re-run a sequence of PMCMC iterations, one can use the pmcmc method on a pmcmc object. By default, the same parameters used for the original PMCMC run are re-used (except for tol, max.fail, and verbose, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing PMCMC Iterations

One can continue a series of PMCMC iterations from where one left off using the continue method. A call to pmcmc to perform Nmcmc=m iterations followed by a call to continue to perform Nmcmc=n iterations will produce precisely the same effect as a single call to pmcmc to perform Nmcmc=m+n iterations. By default, all the algorithmic parameters are the same as used in the original call to pmcmc. Additional arguments will override the defaults.

Details

pmcmc implements an MCMC algorithm in which the true likelihood of the data is replaced by an unbiased estimate computed by a particle filter. This gives an asymptotically correct Bayesian procedure for parameter estimation (Andrieu and Roberts, 2009). An extension to give a correct Bayesian posterior distribution of unobserved state variables (Andrieu et al, 2010) has not yet been implemented.

Author(s)

Edward L. Ionides <ionides at umich dot edu>, Aaron A. King <kingaa at umich dot edu>

References

C. Andrieu, A. Doucet and R. Holenstein, Particle Markov chain Monte Carlo methods, J. R. Stat. Soc. B, to appear, 2010.

C. Andrieu and G.O. Roberts, The pseudo-marginal approach for efficient computation, Ann. Stat. 37:697-725, 2009.

See Also

pmcmc-class, pmcmc-methods, pomp, pomp-class, pfilter. See the "intro_to_pomp" vignette
for an example [CURRENTLY, ONLY DEMONSTRATING THE MIF ALGORITHM, WHICH IS
ALGORITHMICALLY VERY SIMILAR TO PMCMC SINCE THEY BOTH DEPEND CRITI-
CALLY ON A PARTICLE FILTERING STEP].

---

pmcmc-methods   *Methods of the "pmcmc" class*

Description

Methods of the "pmcmc" class.

Usage

```
## S4 method for signature 'pmcmc'
logLik(object, ...)
## S4 method for signature 'pmcmc'
conv.rec(object, pars, ...)
## S4 method for signature 'pmcmc'
filter.mean(object, pars, ...)
## S4 method for signature 'pmcmc'
plot(x, y = NULL, ...)
compare.pmcmc(z)
```

Arguments

object, x	The pmcmc object.
pars	Names of parameters.
y	Ignored.
z	A pmcmc object or list of pmcmc objects.
...	Further arguments (either ignored or passed to underlying functions).

Methods

conv.rec conv.rec(object, pars = NULL) returns the columns of the convergence-record matrix
corresponding to the names in pars. By default, all rows are returned.

logLik Returns the value in the loglik slot.

pmcmc Re-runs the PMCMC iterations. See the documentation for pmcmc.

compare.pmcmc Given a pmcmc object or a list of pmcmc objects, compare.pmcmc produces a set
of diagnostic plots.

plot Plots a series of diagnostic plots. When x is a pmcmc object, plot(x) is equivalent to compare.pmcmc(list(x)).

filter.mean filter.mean(object, pars = NULL) returns the rows of the filtering-mean matrix
corresponding to the names in pars. By default, all rows are returned.

print Prints a summary of the pmcmc object.

show Displays the pmcmc object.

pfilter See [pfilter](#).

Author(s)

Edward L. Ionides <ionides at umich dot edu>, Aaron A. King <kingaa at umich dot edu>

References

C. Andrieu, A. Doucet and R. Holenstein, Particle Markov chain Monte Carlo methods, J. Roy. Stat. Soc B, to appear, 2010.

C. Andrieu and G.O. Roberts, The pseudo-marginal approach for efficient computation, Ann Stat 37:697-725, 2009.

See Also

[pmcmc](#), [pomp](#), [pomp-class](#), [pfilter](#)

---

pomp   *Partially-observed Markov process object.*

Description

Create a new pomp object to hold a partially-observed Markov process model together with a uni- or multi-variate time series.

Usage

```
## S4 method for signature 'data.frame'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
        measurement.model,
      skeleton = NULL, skeleton.type = c("map","vectorfield"), skelmap.delta.t = 1,
       initializer, rprior, dprior, params, covar, tcovar,
       obsnames, statenames, paramnames, covarnames, zeronames,
       PACKAGE, parameter.transform, parameter.inv.transform)
## S4 method for signature 'numeric'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
        measurement.model,
      skeleton = NULL, skeleton.type = c("map","vectorfield"), skelmap.delta.t = 1,
       initializer, rprior, dprior, params, covar, tcovar,
       obsnames, statenames, paramnames, covarnames, zeronames,
       PACKAGE, parameter.transform, parameter.inv.transform)
## S4 method for signature 'matrix'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
        measurement.model,
      skeleton = NULL, skeleton.type = c("map","vectorfield"), skelmap.delta.t = 1,
```

```
        initializer, rprior, dprior, params, covar, tcovar,
        obsnames, statenames, paramnames, covarnames, zeronames,
        PACKAGE, parameter.transform, parameter.inv.transform)
## S4 method for signature 'pomp'
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
        measurement.model,
        skeleton, skeleton.type, skelmap.delta.t,
        initializer, rprior, dprior, params, covar, tcovar,
        obsnames, statenames, paramnames, covarnames, zeronames,
        PACKAGE, parameter.transform, parameter.inv.transform)
```

Arguments

data, times
: The time series data and times at which observations are made. data can be specified as a vector, a matrix, a data-frame, or a pomp object.. If data is a numeric vector, times must be a numeric vector of the same length. If data is a matrix, it should have dimensions nobs x ntimes, where nobs is the number of observed variables and ntimes is the number of times at which observations were made (i.e., each column is a distinct observation of the nobs variables). In this case, times must be given as a numeric vector (of length ntimes). If data is a data-frame, times must name of the column of observation times. Note that, in this case, data is a data-frame, it will be internally coerced to an array with storage-mode 'double'. Note that the times must be numeric and strictly increasing.

t0
: The zero-time. This must be no later than the time of the first observation, times[1]. The stochastic dynamical system is initialized at time t0.

rprocess
: optional function; a function of prototype rprocess(xstart,times,params,...) that simulates from the unobserved process. The easiest way to specify rprocess is to use one of the plugins provided as part of the **pomp** package. See below for details.

dprocess
: optional function; a function of prototype dprocess(x,times,params,log,...) that evaluates the likelihood of a sequence of consecutive state transitions. The easiest way to specify dprocess is to use one of the plugins provided as part of the **pomp** package. It is not typically necessary (or even feasible) to define dprocess. See below for details.

rmeasure
: optional; the measurement model simulator. This can be specified in one of three ways: (1) as a function of prototype rmeasure(x,t,params,...) that makes a draw from the observation process given states x, time t, and parameters params. (2) as the name of a native (compiled) routine with prototype "pomp_measure_model_simulator" as defined in the header file "examples/pomp.h". In the above cases, if the measurement model depends on covariates, the optional argument covars will be filled with interpolated values at each call. (3) using the formula-based measurement.model facility (see below).

dmeasure
: optional; the measurement model probability density function. This can be specified in one of three ways: (1) as a function of prototype dmeasure(y,x,t,params,log,...) that computes the p.d.f. of y given x, t, and params. (2) as the name of a native (compiled) routine with prototype "pomp_measure_model_density" as defined

in the header file "examples/pomp.h". In the above cases, if the measurement model depends on covariates, the optional argument covars will be filled with interpolated values at each call. (3) using the formula-based measurement.model facility (see below). As might be expected, if log=TRUE, this function should return the log likelihood.

measurement.model

optional; a formula or list of formulae, specifying the measurement model. These formulae are parsed internally and used to generate rmeasure and dmeasure functions. If measurement.model is given it overrides any specification of rmeasure or dmeasure. See below for an example. **NB:** it will typically be possible to acclerate measurement model computations by writing dmeasure and/or rmeasure functions directly.

skeleton, skeleton.type, skelmap.delta.t

The function skeleton specifies the deterministic skeleton of the unobserved Markov process. If we are dealing with a discrete-time Markov process, its deterministic skeleton is a map: indicate this by specifying skeleton.type="map". If we are dealing with a continuous-time Markov process, its deterministic skeleton is a vectorfield: indicate this by specifying skeleton.type="vectorfield".

The skeleton function can be specified in one of two ways: (1) as an R function of prototype skeleton(x,t,params,...) that evaluates the deterministic skeleton at state x and time t given the parameters params, or (2) as the name of a native (compiled) routine with prototype "pomp_skeleton" as defined in the header file "pomp.h". If the deterministic skeleton depends on covariates, the optional argument covars will be filled with interpolated values of the covariates at the time t.

With a discrete-time skeleton, the default assumption is that time advances 1 unit per iteration of the map; to change this, set skelmap.delta.t to the appropriate time-step.

initializer   optional function of prototype initializer(params,t0,...) that yields initial conditions for the state process when given a vector, params, of parameters. By default (i.e., if it is unspecified when pomp is called), the initializer assumes any parameters in params the names of which end in ".0" are initial values. These are simply copied over as initial conditions when init.state is called (see init.state-pomp). The names of the state variables are the same as the corresponding initial value parameters, but with the ".0" dropped.

rprior   optional; function drawing a sample from a prior distribution on parameters. This can be specified in one of two ways: (1) as an R function of prototype rprior(params,...), (2) as the name of a native (compiled) routine with prototype "pomp_rprior" as defined in the header file "pomp.h".

dprior   optional; function evaluating the prior distribution. This can be specified in one of two ways: (1) as an R function of prototype dprior(params,log=FALSE,...), (2) as the name of a native (compiled) routine with prototype "pomp_dprior" as defined in the header file "pomp.h".

params   optional named numeric vector of parameters.

covar, tcovar An optional table of covariates: covar is the table (with one column per variable) and tcovar the corresponding times (one entry per row of covar). covar

can be specified as either a matrix or a data frame. In either case the columns are taken to be distinct covariates. If covar is a data frame, tcovar can be either the name or the index of the time variable. If a covariate table is supplied, then the value of each of the covariates is interpolated as needed, i.e., whenever rprocess, dprocess, rmeasure, dmeasure, skeleton, or init.state is evaluated. The resulting interpolated values are passed to the corresponding functions as a numeric vector named covars.

obsnames, statenames, paramnames, covarnames

Optional character vectors specifying the names of observables, state variables, parameters, or covariates, respectively. These are only used in the event that one or more of the basic functions (rprocess, dprocess, rmeasure, dmeasure, skeleton) are defined using native routines. In that case, these name vectors are matched against the corresponding names and the indices of the names are passed to the native routines. Using this facility allows one to write one or more of rprocess, dprocess, rmeasure, dmeasure, skeleton in native code in a way that does not depend on the order of states, parameters, and covariates at run time. See the "Advanced topics in pomp" vignette for more on this topic and examples.

zeronames   Optional character vector specifying the names of accumulator variables. See the "Advanced topics in pomp" vignette for a discussion of this.

PACKAGE   An optional string giving the name of the dynamically loaded library in which any native routines are to be found.

parameter.transform, parameter.inv.transform

Optional functions specifying parameter transformations. These functions must have arguments params and .... parameter.transform should transform parameters from the user's scale to the scale that rprocess, dprocess, rmeasure, dmeasure, skeleton, and initializer will use internally. parameter.inv.transform should be the inverse of parameter.transform. Note that it is the user's responsibility to make sure this holds. If obj is the constructed pomp object, and coef(obj) is non-empty, a simple check of this is x <- coef(obj,transform=TRUE); obj1 <- obj; c and identical(coef(obj1,transform=TRUE),x). By default, both functions are the identity transformation. See the "introduction_to_pomp" vignette for an example.

...   Any additional arguments given to pomp will be stored in the pomp object and passed as arguments to each of the functions rprocess, dprocess, rmeasure, dmeasure, and initializer whenever they are evaluated.

Details

It is not typically necessary (or desirable, or even feasible) to define all of the functions rprocess**,** dprocess**,** rmeasure**,** dmeasure**, and** skeleton **in any given problem. Each algorithm makes use of a different subset of these functions.** In general, the specification of process-model codes rprocess and/or dprocess can be somewhat nontrivial: for this reason, plugins have been developed to streamline this process for the user. Currently, if one's process model evolves in discrete time or one is willing to make such an approximation (e.g., via an Euler approximation), then the euler.sim, discrete.time.sim, or onestep.sim plugin for rprocess and onestep.dens plugin for dprocess are available. For exact simulation of certain continuous-

time Markov chains, an implementation of Gillespie's algorithm is available (see `gillespie.sim`). To use the plugins, consult the help documentation (?`plugins`) and the vignettes.

It is anticipated that, in specific cases, it will be possible to obtain increased computational efficiency by writing custom versions of `rprocess` and/or `dprocess`. See the "Advanced topics in pomp" vignette for a discussion of this. If such custom versions are desired, the following describes how each of these functions should be written in this case.

`rprocess` In general, the specification of `rprocess` can be somewhat nontrivial: for this reason, `plugins` have been developed to streamline this process for the user. Currently, if one's process model evolves in discrete time or one is willing to make such an approximation (e.g., via an Euler approximation), then the `euler.sim` or `onestep.sim` plugin is available. For exact simulation of certain continuous-time Markov chains, an implementation of Gillespie's algorithm is available (see `gillespie.sim`). To use the plugins, consult the help documentation (?`plugins`) and the vignettes.

If the plugins are not used `rprocess` must have at least the following arguments: `xstart`, `times`, `params`, and `...`. It can also take additional arguments. It is guaranteed that these will be filled with the corresponding elements the user has included as additional arguments in the construction of the `pomp` object.

In calls to `rprocess`, `xstart` can be assumed to be a rank-2 array (matrix) with rows corresponding to state variables and columns corresponding to independent realizations of the process. `params` will similarly be a rank-2 array with rows corresponding to parameters and columns corresponding to independent realizations. The columns of `params` correspond to those of `xstart`; in particular, they will agree in number. Both `xstart` and `params` will have rownames, which are available for use by the user.

`rprocess` must return a rank-3 array with rownames. Suppose `x` is the array returned. Then `dim(x)=c(nvars,nreps,ntimes)`, where `nvars` (=`nrow(xstart)`) is the number of state variables, `nreps` (=`ncol(xstart)`) is the number of independent realizations simulated, and `ntimes` is the length of the vector `times`. `x[,j,k]` is the value of the state process in the j-th realization at time `times[k]`. In particular, `x[,,1]` must be identical to `xstart`. The rownames of `x` must correspond to those of `xstart`.

At present, the following methods make use of `rprocess`:

- `simulate`
- `pfilter`
- `mif`
- `nlf`
- `probe`
- `probe.match`

`dprocess` In general, the specification of `dprocess` can be somewhat nontrivial: for this reason, `plugins` have been developed to streamline this process for the user. Currently, if one's process model evolves in discrete time or one is willing to make such an approximation (e.g., via an Euler approximation), then the `onestep.dens` plugin for `dprocess` is available. To use the plugins, consult the help documentation (?`plugins`) and the vignettes.

If the plugins are not used, `dprocess` must have at least the following arguments: `x`, `times`, `params`, `log`, and `...`. It may take additional arguments. It is guaranteed that these will be filled with the corresponding elements the user has included as additional arguments in the construction of the `pomp` object.

In calls to dprocess, x may be assumed to be an nvars x nreps x ntimes array, where these terms have the same meanings as above. params will be a rank-2 array with rows corresponding to individual parameters and columns corresponding to independent realizations. The columns of params correspond to those of x; in particular, they will agree in number. Both x and params will have rownames, available for use by the user.

dprocess must return a rank-2 array (matrix). Suppose d is the array returned. Then dim(d)=c(nreps,ntimes-1). d[j,k] is the probability density of the transition from state x[,j,k-1] at time times[k-1] to state x[,j,k] at time times[k]. If log=TRUE, then the log of the pdf is returned.

In writing this function, you may assume that the transitions are consecutive. It should be quite clear that, but for this assumption, it would be quite difficult in general to write the transition probabilities. In fact, from one perspective, the algorithms in **pomp** are designed to overcome just this difficulty.

At present, no methods in pomp make use of dprocess**.**

The measurement-model, deterministic skeleton, and initializer components are easily specified without the use of plugins. The following is a guide to writing these components.

rmeasure if provided, must take at least the arguments x, t, params, and .... It may take additional arguments, which will be filled with user-specified data as above. x may be assumed to be a named numeric vector of length nvars, (which has the same meanings as above). t is a scalar quantity, the time at which the measurement is made. params may be assumed to be a named numeric vector of length npars.

rmeasure must return a named numeric vector. If y is the returned vector, then length(y)=nobs, where nobs is the number of observable variables.

At present, the following methods make use of rmeasure:

- simulate
- nlf
- probe
- probe.match

dmeasure if provided, must take at least the arguments y, x, t, params, log, and .... y may be assumed to be a named numeric vector of length nobs containing (actual or simulated) values of the observed variables; x will be a named numeric vector of length nvar containing state variables params, a named numeric vector containing parameters; and t, a scalar, the corresponding observation time. It may take additional arguments which will be filled with user-specified data as above. dmeasure must return a single numeric value, the pdf of y given x at time t. If log=TRUE, then the log of the pdf is returned.

At present, the following methods make use of dmeasure:

- pfilter
- mif

skeleton If skeleton is an R function, it must have at least the arguments x, t, params, and .... x is a numeric vector containing the coordinates of a point in state space at which evaluation of the skeleton is desired. t is a numeric value giving the time at which evaluation of the skeleton is desired. Of course, these will be irrelevant in the case of an autonomous skeleton. params is a numeric vector holding the parameters. The optional argument covars is a numeric vector containing the values of the covariates at the time t. covars will have one value for each column of the covariate table specified when the pomp object was created. covars

is constructed from the covariate table (see covar, below) by interpolation. skeleton may take additional arguments, which will be filled, as above, with user-specified data. skeleton must return a numeric vector of the same length as x. The return value is interpreted as the vectorfield (if the dynamical system is continuous) or the value of the map (if the dynamical system is discrete), at the point x at time t.

If skeleton is the name of a native routine, this routine must be of prototype "pomp_skeleton" as defined in the header "pomp.h" (see the "include" directory in the installed package directory).

At present, the following methods make use of skeleton:

- [trajectory](#)
- [traj.match](#)

initializer if provided, must have at least the arguments params, t0, and .... params is a named numeric vector of parameters. t0 will be the time at which initial conditions are desired. initializer must return a named numeric vector of initial states.

Value

An object of class pomp. If data is an object of class [pomp](#), then by default the returned pomp object is identical to data. If additional arguments are given, these override the defaults.

Warning

Some error checking is done by pomp, but complete error checking is impossible. If the user-specified functions do not conform to the above specifications (see Details), then the results may be invalid. In particular, if both rmeasure and dmeasure are specified, the user should verify that these two functions correspond to the same model and if skeleton is specified, the user is responsible for verifying that it corresponds to the true deterministic skeleton of the model. Each **pomp**-package algorithm uses some subset of the five basic components (rprocess, dprocess, rmeasure, dmeasure, skeleton). If an algorithm requires a component that was not given in the construction of the pomp object, an error is generated.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pomp-methods](#), [plugins](#), [time](#), [time<-](#), [timezero](#), [timezero<-](#), [coef](#), [coef<-](#), [obs](#), [states](#), [window](#), [as.data.frame.pomp](#)

Examples

```
## For examples, see the vignettes, the data()-loadable
## example \code{pomp} objects, and the provided example files.
## Not run:
vignette("intro_to_pomp")
vignette("advanced_topics_in_pomp")
pompExample()
pomp.home <- system.file("examples",package="pomp")
```

```
pomp.examples <- list.files(pomp.home)
file.show(
          file.path(pomp.home,pomp.examples),
          header=paste("======",pomp.examples,"=======")
         )

## End(Not run)
```

---

pomp-methods   *Methods of the "pomp" class*

Description

Methods of the pomp class.

Usage

```
## S3 method for class 'pomp'
as.data.frame(x, row.names, optional, ...)
## S4 method for signature 'pomp'
coef(object, pars, transform = FALSE, ...)
## S4 replacement method for signature 'pomp'
coef(object, pars, transform = FALSE, ...) <- value
## S4 method for signature 'pomp'
obs(object, vars, ...)
## S4 method for signature 'pomp'
data.array(object, vars, ...)
## S4 method for signature 'pomp'
partrans(object, params, dir = c("forward","inverse"), ...)
## S4 method for signature 'pomp'
plot(x, y, variables, panel = lines,
     nc = NULL, yax.flip = FALSE,
     mar = c(0, 5.1, 0, if (yax.flip) 5.1 else 2.1),
     oma = c(6, 0, 5, 0), axes = TRUE, ...)
## S4 method for signature 'pomp'
print(x, ...)
## S4 method for signature 'pomp'
show(object)
## S4 method for signature 'pomp'
states(object, vars, ...)
## S4 method for signature 'pomp'
time(x, t0 = FALSE, ...)
## S4 replacement method for signature 'pomp'
time(object, t0 = FALSE, ...) <- value
## S4 method for signature 'pomp'
timezero(object, ...)
## S4 replacement method for signature 'pomp'
```

```
timezero(object, ...) <- value
## S4 method for signature 'pomp'
window(x, start, end, ...)
## S4 method for signature 'pomp'
as(object, class)
## S4 method for signature 'pomp,data.frame'
coerce(from, to = "data.frame", strict = TRUE)
```

Arguments

`object, x`	The pomp object.
`pars`	optional character; names of parameters to be retrieved or set.
`vars`	optional character; names of observed variables to be retrieved.
`transform`	optional logical; should the parameter transformations be applied?
`value`	numeric; values to be assigned.
`params`	a vector or matrix of parameters to be transformed.
`dir`	direction of the transformation. `dir="forward"` applies the transformation from the "natural" scale to the "internal" scale. This is the transformation specified by the `parameter.transform` argument to pomp; it is stored in the 'par.trans' slot of `object`. `dir="inverse"` applies the inverse transformation (stored in the 'par.untrans' slot).
`t0`	logical; if TRUE on a call to `time`, the zero time is prepended to the time vector; if TRUE on a call to `time<-`, the first element in `value` is taken to be the initial time.
`start, end`	start and end times of the window.
`class`	character; name of the class to which `object` should be coerced.
`from, to`	the classes between which coercion should be performed.
`strict`	ignored.
`y`	ignored.
`variables`	optional character; names of variables to plot.
`panel`	a function of prototype `panel(x, col, bg, pch, type, ...)` which gives the action to be carried out in each panel of the display.
`nc`	the number of columns to use. Defaults to 1 for up to 4 series, otherwise to 2.
`yax.flip`	logical; if TRUE, the y-axis (ticks and numbering) should flip from side 2 (left) to 4 (right) from series to series.
`mar, oma`	the 'par' settings for 'mar' and 'oma' to use. Modify with care!
`axes`	logical; indicates if x- and y- axes should be drawn.
`row.names, optional`	ignored.
`...`	Further arguments (either ignored or passed to underlying functions).

Details

coef coef(object) returns the contents of the params slot of object. coef(object,pars) re-
turns only those parameters named in pars. coef(object,transform=TRUE) returns parameter.inv.transform(coe
where parameter.inv.transform is the user parameter inverse transformation function spec-
ified when object was created. Likewise, coef(object,pars,transform=TRUE) returns
parameter.inv.transform(coef(object))[pars].

coef<- Assigns values to the params slot of the pomp object. coef(object) <- value has
the effect of replacing the parameters of object with value. If coef(object) exists, then
coef(object,pars) <- value replaces those parameters of object named in pars with the
elements of value; the names of value are ignored. If some of the names in pars do not al-
ready name parameters in coef(object), then they are concatenated. If coef(object) does
not exist, then coef(object,pars) <- value assigns value to the parameters of object;
in this case, the names of object will be pars and the names of value will be ignored.
coef(object,transform=TRUE) <- value assigns parameter.transform(value) to the
params slot of object. Here, parameter.transform is the parameter transformation func-
tion specified when object was created. coef(object,pars,transform=TRUE) <- value
first, discards any names the value may have, sets names(value) <- pars, and then replaces
the elements of object's params slot parameter.transform(value). In this case, if some
of the names in pars do not already name parameters in coef(object,transform=TRUE),
then they are concatenated.

obs, data.array These functions are synonymous. obs(object) returns the array of observations.
obs(object,vars) gives just the observations of variables named vars. vars may specify
the variables by position or by name.

states states(object) returns the array of states. states(object,vars) gives just the state
variables named in vars. vars may specify the variables by position or by name.

time time(object) returns the vector of observation times. time(object,t0=TRUE) returns the
vector of observation times with the zero-time t0 prepended.

time<- time(object) <- value replaces the observation times slot (times) of object with
value. time(object,t0=TRUE) <- value has the same effect, but the first element in value
is taken to be the initial time. The second and subsequent elements of value are taken to be
the observation times. Those data and states (if they exist) corresponding to the new times are
retained.

timezero, timezero<- timezero(object) returns the zero-time t0. timezero(object) <- value
sets the zero-time to value.

window window(x,start=t1,end=t2 returns a new pomp object, identical to x but with only the
data in the window between times t1 and t2 (inclusive). By default, start is the time of the
first observation and end is the time of the last.

show Displays the pomp object.

plot Plots the data and state trajectories (if the latter exist). Additional arguments are passed to the
low-level plotting routine.

print Prints the pomp object in a nice way.

as, coerce The coerce method should typically not be used directly. It is defined by setAs as a
method to be used by as. A pomp object can be coerced to a data frame via as(object,"data.frame").
The data frame contains the times, the data, and the state trajectories, if they exist.

rprocess simulates the process model. See `rprocess-pomp`.

dprocess evaluates the process model density. See `dprocess-pomp`.

rmeasure simulates the measurement model. See `rmeasure-pomp`.

dmeasure evaluates the measurement-model density. See `dmeasure-pomp`.

skeleton evaluates the deterministic skeleton (be it a vector field or a map). See `skeleton-pomp`.

init.state returns a vector of initialial conditions. See `init.state-pomp`.

simulate `simulate` can be used to simulate state and observation trajectories. See documentation under `simulate-pomp`.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

`pomp`, `pomp-class`, `rprocess`, `dprocess`, `rmeasure`, `dmeasure`, `init.state`, `simulate`

---

pompBuilder	*Write, compile, link, and build a pomp object using native codes*

Description

`pompBuilder` is an EXPERIMENTAL facility for producing compiled pomp objects.

Usage

```
pompBuilder(data, times, t0, name, statenames, paramnames,
            tcovar, covar,
            rmeasure, dmeasure, step.fn, step.fn.delta.t,
            skeleton, skeleton.type, skelmap.delta.t = 1,
            parameter.transform, parameter.inv.transform,
            ..., link = TRUE, save = FALSE)
```

Arguments

data, times, t0

The data, times, and zero-time. See `pomp` for more information. `data` must be a data-frame.

name   character; the stem of the name for the files that will be produced.

statenames, paramnames

names of state-variables and parameters, respectively

tcovar, covar   optional. `covar` is a data-frame containing covariates (variables in columns, timepoints in rows); `tcovar` is the name of the column containing time.

rmeasure, dmeasure

C codes implementing the measurement model

step.fn, step.fn.delta.t

>   step.fn is a C code that implements an Euler step function. The Euler time-step is step.fn.delta.t, which should be a positive number.

skeleton, skeleton.type, skelmap.delta.t

>   skeleton is a C code that implements the deterministic skeleton. As in pomp, skeleton.type indicates whether the skeleton is a map (discrete-time) or vectorfield (continuous-time). If the former, skelmap.delta.t is the time-step of the map.

parameter.transform, parameter.inv.transform

>   optional C codes that implement parameter transformations. parameter.transform maps parameters from the estimation scale to the natural scale; parameter.inv.transformation maps them from the natural scale to the estimation scale.

...   additional arguments are passed to pomp

link   logical; if TRUE, the resulting code will be linked after compilation.

save   logical; if TRUE, the resulting C code will be saved in the current working directory. The shared-object library will also be created in this directory. If FALSE, the C code and shared-object library will be created in the temporary directory.

Details

pompBuilder assumes that files can be written to the current working directory and that dynamic shared objects can be compiled and linked, i.e., that R CMD SHLIB will work. This will not typically be the case in out-of-the-box Windows installations.

Value

The constructed pomp object. A side-effect is the writing and compilation of a C code into a dynamics shared object. These files will reside in the current working directory (see getwd). If pompBuilder has been called with link=FALSE, this dynamic shared object must be linked (see dyn.load) before the pomp object can be used.

Parallel computing caution: note that, since a side-effect of running pompBuilder is the writing and compilation of a file, caution must be used when running parallel computations. Simultaneous execution of pompBuilder on nodes sharing a common disk will cause all nodes to attempt to write the same file at the same time; the results will be undefined. If multiple nodes are to work on the same pomp object built by pompBuilder, first run pompBuilder once (perhaps with link=FALSE) to write the C file and compile it into a dynamic shared object file. The pomp object returned by pompBuilder can then be shared with all worker nodes. Note, however, that all worker nodes must dynamically load (using dyn.load) the dynamic shared object file before anything can be done with the pomp object.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

pomp and the demos.

Examples

```
## see the demos
```

---

pompExample   *Pre-built examples of pomp objects.*

Description

pompExample loads pre-built example pomp objects.

Usage

```
pompExample(example, envir = .GlobalEnv)
```

Arguments

example	example to load given as a name or literal character string. Evoked without an argument, pompExample lists all available examples.
envir	the environment into which the objects should be loaded.

Value

pompExample has the side effect of creating one or more pomp objects in the global workspace.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[blowflies](), [dacca](), [gompertz](), [ou2](), [ricker](), [rw2](), [euler.sir](), [gillespie.sir](), [bbs](), [verhulst]()

Examples

```
pompExample()
pompExample(euler.sir)
pompExample("gompertz")
```

---

probe   *Probe a partially-observed Markov process.*

Description

probe applies one or more "probes" to time series data and model simulations and compares the results. It can be used to diagnose goodness of fit and/or as the basis for "probe-matching", a generalized method-of-moments approach to parameter estimation. probe.match calls an optimizer to adjust model parameters to do probe-matching, i.e., to minimize the discrepancy between simulated and actual data. This discrepancy is measured using the "synthetic likelihood" as defined by Wood (2010). probe.match.objfun constructs an objective function for probe-matching suitable for use in optim-like optimizers.

Usage

```
## S4 method for signature 'pomp'
probe(object, probes, params, nsim, seed = NULL, ...)
## S4 method for signature 'probed.pomp'
probe(object, probes, params, nsim, seed, ...)
## S4 method for signature 'pomp'
probe.match.objfun(object, params, est, probes, nsim,
            seed = NULL, fail.value = NA, transform = FALSE, ...)
## S4 method for signature 'probed.pomp'
probe.match.objfun(object, probes, nsim, seed, ...)
## S4 method for signature 'pomp'
probe.match(object, start, est = character(0),
            probes, nsim, seed = NULL,
            method = c("subplex","Nelder-Mead","SANN","BFGS",
                       "sannbox","nloptr"),
            verbose = getOption("verbose"),
            fail.value = NA, transform = FALSE, ...)
## S4 method for signature 'probed.pomp'
probe.match(object, probes, nsim, seed,
            ..., verbose = getOption("verbose"))
## S4 method for signature 'probe.matched.pomp'
probe.match(object, est, probes,
            nsim, seed, transform, fail.value, ...,
            verbose = getOption("verbose"))
```

Arguments

object   An object of class pomp.

probes   A single probe or a list of one or more probes. A probe is simply a scalar- or
                vector-valued function of one argument that can be applied to the data array of
                a pomp. A vector-valued probe must always return a vector of the same size. A
                number of basic examples are provided with the package (see basic.probes).

params	optional named numeric vector of model parameters. By default, `params=coef(object)`.
nsim	The number of model simulations to be computed.
seed	optional; if non-NULL, the random number generator will be initialized with this seed for simulations. See [simulate-pomp](#).
start	named numeric vector; the initial guess of parameters.
est	character vector; the names of parameters to be estimated.
method	Optimization method. Choices refer to algorithms used in [optim](#), [subplex](#), and [nloptr](#).
verbose	logical; print diagnostic messages?
fail.value	optional numeric scalar; if non-NA, this value is substituted for non-finite values of the objective function. It should be a large number (i.e., bigger than any legitimate values the objective function is likely to take).
transform	logical; if TRUE, optimization is performed on the transformed scale.
...	Additional arguments. In the case of `probe`, these are currently ignored. In the case of `probe.match`, these are passed to the optimizer (one of [optim](#), [subplex](#), [nloptr](#), or [sannbox](#)). These are passed via the optimizer's `control` list (in the case of optim, subplex, and sannbox) or the `opts` list (in the case of nloptr).

Details

A call to `probe` results in the evaluation of the probe(s) in `probes` on the data. Additionally, `nsim` simulated data sets are generated (via a call to [simulate](#)) and the probe(s) are applied to each of these. The results of the probe computations on real and simulated data are stored in an object of class `probed.pomp`.

A call to `probe.match` results in an attempt to optimize the agreement between model and data, as measured by the specified probes, over the parameters named in `est`. The results, including coefficients of the fitted model and values of the probes for data and fitted-model simulations, are stored in an object of class [probe.matched.pomp](#).

The objective function minimized by `probe.match` — in a form suitable for use with [optim](#)-like optimizers — is created by a call to `probe.match.objfun`. Specifically, `probe.match.objfun` will return a function that takes a single numeric-vector argument that is assumed to cotain the parameters named in `est`, in that order. This function will return the negative synthetic log likelihood for the probes specified.

Value

`probe` returns an object of class `probed.pomp`. `probed.pomp` is derived from the [pomp](#) class and therefore have all the slots of pomp. In addition, a `probed.pomp` class has the following slots:

probes list of the probes applied.

datvals, simvals values of each of the probes applied to the real and simulated data, respectively.

quantiles fraction of simulations with probe values less than the value of the probe of the data.

pvals two-sided p-values: fraction of the `simvals` that deviate more extremely from the mean of the `simvals` than does `datavals`.

synth.loglik the log synthetic likelihood (Wood 2010). This is the likelihood assuming that the probes are multivariate-normally distributed.

probe.match returns an object of class probe.matched.pomp, which is derived from class probed.pomp. probe.matched.pomp objects therefore have all the slots above plus the following:

est, transform, fail.value values of the corresponding arguments in the call to probe.match.

value value of the objective function at the optimum.

evals number of function and gradient evaluations by the optimizer. See optim.

convergence, msg Convergence code and message from the optimizer. See optim and nloptr.

probe.match.objfun returns a function suitable for use as an objective function in an optim-like optimizer.

Author(s)

Daniel C. Reuman, Aaron A. King

References

B. E. Kendall, C. J. Briggs, W. M. Murdoch, P. Turchin, S. P. Ellner, E. McCauley, R. M. Nisbet, S. N. Wood Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches, Ecology, 80:1789–1805, 1999.

S. N. Wood Statistical inference for noisy nonlinear ecological dynamic systems, Nature, 466: 1102–1104, 2010.

See Also

pomp-class, pomp-methods, basic.probes, probe.match

Examples

```
pompExample(ou2)
good <- probe(
             ou2,
             probes=list(
               y1.mean=probe.mean(var="y1"),
               y2.mean=probe.mean(var="y2"),
               y1.sd=probe.sd(var="y1"),
               y2.sd=probe.sd(var="y2"),
               extra=function(x)max(x["y1",])
               ),
             nsim=500
             )
summary(good)
plot(good)

bad <- probe(
             ou2,
             params=c(alpha.1=0.1,alpha.4=0.2,x1.0=0,x2.0=0,
                      alpha.2=-0.5,alpha.3=0.3,
```

```
                      sigma.1=3,sigma.2=-0.5,sigma.3=2,
                      tau=1),
             probes=list(
                y1.mean=probe.mean(var="y1"),
                y2.mean=probe.mean(var="y2"),
                y1.sd=probe.sd(var="y1"),
                y2.sd=probe.sd(var="y2"),
                extra=function(x)range(x["y1",])
                ),
             nsim=500
             )
summary(bad)
plot(bad)
```

---

probed.pomp-methods   *Methods of the "probed.pomp", "probe.matched.pomp", "spect.pomp",*
                           and "spect.matched.pomp" classes

Description

Methods of the probed.pomp, probe.matched.pomp, spect.pomp, and spect.matched.pomp classes

Usage

```
## S4 method for signature 'probed.pomp'
summary(object, ...)
## S4 method for signature 'probed.pomp'
plot(x, y, ...)
## S4 method for signature 'probe.matched.pomp'
summary(object, ...)
## S4 method for signature 'probe.matched.pomp'
plot(x, y, ...)
## S4 method for signature 'spect.pomp'
summary(object, ...)
## S4 method for signature 'probed.pomp'
logLik(object, ...)
## S4 method for signature 'spect.pomp'
plot(x, y, max.plots.per.page = 4,
          plot.data = TRUE,
          quantiles = c(.025, .25, .5, .75, .975),
          quantile.styles = list(lwd=1, lty=1, col="gray70"),
          data.styles = list(lwd=2, lty=2, col="black"))
## S4 method for signature 'spect.matched.pomp'
summary(object, ...)
## S4 method for signature 'spect.matched.pomp'
plot(x, y, ...)
## S4 method for signature 'probed.pomp'
as(object, class)
```

Arguments

`object, x`	the object to be summarized or plotted.
`y`	ignored.
`max.plots.per.page`	
	maximum number of plots per page
`plot.data`	plot the data spectrum?
`quantiles`	quantiles to plot
`quantile.styles`	
	plot style parameters for the quantiles
`data.styles`	plot style parameters for the data spectrum
`class`	character; name of the class to which `object` should be coerced.
`...`	Further arguments (either ignored or passed to underlying functions).

Methods

plot displays diagnostic plots.

summary displays summary information.

logLik returns the synthetic likelihood for the probes. NB: in general, this is not the same as the likelihood.

as when a 'probed.pomp' is coerced to a 'data.frame', the first row gives the probes applied to the data; the rest of the rows give the probes evaluated on simulated data. The rownames of the result can be used to distinguish these.

In addition, slots of this object can be accessed via the $ operator.

Author(s)

Daniel C. Reuman, Aaron A. King

See Also

[probe](), [probed.pomp](), [probe.matched.pomp](), [probe.match]()

---

profileDesign	*Design matrices for likelihood profile calculations.*

Description

`profileDesign` generates a data-frame where each row can be used as the starting point for a profile likelihood calculation.

Usage

```
profileDesign(..., lower, upper, nprof,
              stringsAsFactors = default.stringsAsFactors())
```

Arguments

`...`	Specifies the parameters over which to profile.
`lower, upper`	Named numeric vectors, specifying the range over which the other parameters are to be sampled.
`nprof`	The number of starts per profile point.
`stringsAsFactors`	should character vectors be converted to factors?

Value

`profileDesign` returns a data frame with `nprof` points per profile point. The other parameters in `vars` are sampled using `sobol`.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[sobol](#)

Examples

```
## A one-parameter profile design:
x <- profileDesign(p=1:10,lower=c(a=0,b=0),upper=c(a=1,b=5),nprof=20)
dim(x)
plot(x)
## A two-parameter profile design:
x <- profileDesign(p=1:10,q=3:5,lower=c(a=0,b=0),upper=c(b=5,a=1),nprof=20)
dim(x)
plot(x)
```

---

ricker	*Ricker model with Poisson observations.*

Description

`ricker` is a pomp object encoding a stochastic Ricker model with Poisson measurement error.

Details

The state process is $N_{t+1} = r N_t \exp(-N_t + e_t)$, where the e_t are i.i.d. normal random deviates with zero mean and variance σ^2. The observed variables y_t are distributed as $\mathrm{Poisson}(\phi N_t)$.

See Also

[pomp-class](#) and the vignettes

Examples

```
pompExample(ricker)
plot(ricker)
coef(ricker)
```

---

rw2   *Two-dimensional random-walk process*

Description

rw2 is a pomp object encoding a 2-D normal random walk.

Details

The random-walk process is fully but noisily observed.

See Also

pomp-class and the vignettes

Examples

```
pompExample(rw2)
plot(rw2)
x <- simulate(rw2,nsim=10,seed=20348585L,params=c(x1.0=0,x2.0=0,s1=1,s2=3,tau=1))
plot(x[[1]])
```

---

sannbox   *Simulated annealing with box constraints.*

Description

sannbox is a straightforward implementation of simulated annealing with box constraints.

Usage

```
sannbox(par, fn, control = list(), ...)
```

Arguments

par	Initial values for the parameters to be optimized over.
fn	A function to be minimized, with first argument the vector of parameters over which minimization is to take place. It should return a scalar result.
control	A named list of control parameters. See 'Details'.
...	ignored.

Details

The control argument is a list that can supply any of the following components:

trace Non-negative integer. If positive, tracing information on the progress of the optimization is produced. Higher values may produce more tracing information.

fnscale An overall scaling to be applied to the value of fn during optimization. If negative, turns the problem into a maximization problem. Optimization is performed on fn(par)/fnscale.

parscale A vector of scaling values for the parameters. Optimization is performed on par/parscale and these should be comparable in the sense that a unit change in any element produces about a unit change in the scaled value.

maxit The total number of function evaluations: there is no other stopping criterion. Defaults to 10000.

temp starting temperature for the cooling schedule. Defaults to 1.

tmax number of function evaluations at each temperature. Defaults to 10.

candidate.dist function to randomly select a new candidate parameter vector. This should be a function with three arguments, the first being the current parameter vector, the second the temperature, and the third the parameter scaling. By default, candidate.dist is function(par,temp,scale) rnorm(n=length(par),mean=par,sd=scale*temp).

sched cooling schedule. A function of a three arguments giving the temperature as a function of iteration number and the control parameters temp and tmax. By default, sched is function(k,temp,tmax) temp/log((

Alternatively, one can supply a numeric vector of temperatures. This must be of length at least maxit.

Value

sannbox returns a list with components:

counts two-element integer vector. The first number gives the number of calls made to fn. The second number is provided for compatibility with [optim](#) and will always be NA.

convergence provided for compatibility with [optim](#); will always be 0.

final.params last tried value of par.

final.value value of fn corresponding to final.params.

par best tried value of par.

value value of fn corresponding to par.

Author(s)

Daniel Reuman, Imperial College London and Aaron A. King <kingaa at umich dot edu>

See Also

[traj.match](#), [probe.match](#).

---

simulate-pomp   *Running simulations of a partially-observed Markov process*

Description

`simulate` can be used to generate simulated data sets and/or to simulate the state process.

Usage

```
## S4 method for signature 'pomp'
simulate(object, nsim = 1, seed = NULL, params,
         states = FALSE, obs = FALSE, times, t0,
         as.data.frame = FALSE, ...)
```

Arguments

`object`	An object of class pomp.
`nsim`	The number of simulations to perform. Note that the number of replicates will be `nsim` times `ncol(xstart)`.
`seed`	optional; if set, the pseudorandom number generator (RNG) will be initialized with `seed`. the random seed to use. The RNG will be restored to its original state afterward.
`params`	either a named numeric vector or a numeric matrix with rownames. The parameters to use in simulating the model. If `params` is not given, then the contents of the `params` slot of `object` will be used, if they exist.
`states`	Do we want the state trajectories?
`obs`	Do we want data-frames of the simulated observations?
`times, t0`	`times` specifies the times at which simulated observations will be made. `t0` specifies the start time (the time at which the initial conditions hold). The default for `times` is is `times=time(object,t0=FALSE)` and `t0=timezero(object)`, respectively.
`as.data.frame`	logical; if TRUE, return the result as a data-frame.
`...`	further arguments that are currently ignored.

Details

Simulation of the state process and of the measurement process are each accomplished by a single call to the user-supplied `rprocess` and `rmeasure` functions, respectively. This makes it possible for the user to write highly optimized code for these potentially expensive computations.

Value

If states=FALSE and obs=FALSE (the default), a list of nsim pomp objects is returned. Each has a simulated data set, together with the parameters used (in slot params) and the state trajectories also (in slot states). If times is specified, then the simulated observations will be at times times.

If nsim=1, then a single pomp object is returned (and not a singleton list).

If states=TRUE and obs=FALSE, simulated state trajectories are returned as a rank-3 array with dimensions nvar x (ncol(params)*nsim) x ntimes. Here, nvar is the number of state variables and ntimes the length of the argument times. The measurement process is not simulated in this case.

If states=FALSE and obs=TRUE, simulated observations are returned as a rank-3 array with dimensions nobs x (ncol(params)*nsim) x ntimes. Here, nobs is the number of observables.

If both states=TRUE and obs=TRUE, then a named list is returned. It contains the state trajectories and simulated observations as above.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

pomp-class

Examples

```
pompExample(ou2)
x <- simulate(ou2,seed=3495485,nsim=10)
x <- simulate(ou2,seed=3495485,nsim=10,states=TRUE,obs=TRUE)
```

---

sir   *SIR models.*

Description

euler.sir is a pomp object encoding a simple seasonal SIR model. Simulation is performed using an Euler multinomial approximation. gillespie.sir has the same model implemented using Gillespie's algorithm. bbs is a nonseasonal SIR model together with data from a 1978 outbreak of influenza in a British boarding school.

Details

This example is discussed extensively in the "Introduction to **pomp**" and "Advanced topics in **pomp**" vignettes.

The codes that construct these pomp objects can be found in the "examples" directory in the installed package. Do system.file("examples",package="pomp")) to find this directory. For the basic rprocess, dmeasure, rmeasure, and skeleton functions, these codes use compiled native routines

built into the package's library. View "src/sir.c" in the package source or `file.show("examples/sir.c")`
from an R session to view these codes.

The boarding school influenza outbreak is described in Anonymous (1978).

References

Anonymous (1978). Influenza in a boarding school. British Medical Journal 1:587

See Also

[pomp-class](#) and the vignettes

Examples

```
pompExample(euler.sir)
plot(euler.sir)
plot(simulate(euler.sir,seed=20348585))
coef(euler.sir)

pompExample(gillespie.sir)
plot(gillespie.sir)
plot(simulate(gillespie.sir,seed=20348585))
coef(gillespie.sir)

pompExample(bbs)
plot(bbs)
coef(bbs)
```

---

sliceDesign   *Design matrices for likelihood slices.*

Description

sliceDesign generates a data-frame representing points taken along one or more slices through a
point in a multidimensional space.

Usage

```
sliceDesign(center, ...)
```

Arguments

center	center is a named numeric vector specifying the point through which the slice(s) is (are) to be taken.
...	Additional numeric vector arguments specify the slices.

Value

sliceDesign returns a data frame with one row per point along a slice. The column slice is a factor that tells which slice each point belongs to.

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[profileDesign](#)

Examples

```
## A single 11-point slice through the point c(A=3,B=8,C=0) along the B direction.
x <- sliceDesign(center=c(A=3,B=8,C=0),B=seq(0,10,by=1))
dim(x)
plot(x)
## Two slices through the same point along the A and C directions.
x <- sliceDesign(c(A=3,B=8,C=0),A=seq(0,5,by=1),C=seq(0,5,length=11))
dim(x)
plot(x)
```

---

sobol	*Sobol' low-discrepancy sequence*

Description

Generate a data-frame containing a Sobol' low-discrepancy sequence.

Usage

```
sobol(vars, n)
sobolDesign(lower, upper, nseq)
```

Arguments

vars	Named list of ranges of variables.
lower,upper	named numeric vectors giving the lower and upper bounds of the ranges, respectively.
n,nseq	Number of vectors requested.

Value

sobol	Returns a data frame with n 'observations' of the variables in vars.
sobolDesign	Returns a data frame with nseq 'observations' of the variables over the range specified.

Author(s)

Aaron A. King <kingaa at umich dot edu>

References

W. H. Press, S. A. Teukolsky, W. T. Vetterling, \& B. P. Flannery, Numerical Recipes in C, Cambridge University Press, 1992

See Also

sliceDesign, profileDesign

Examples

```
plot(sobol(vars=list(a=c(0,1),b=c(100,200)),100))
plot(sobolDesign(lower=c(a=0,b=100),upper=c(b=200,a=1),100))
```

---

spect   *Power spectrum computation for partially-observed Markov processes.*

Description

spect estimates the power spectrum of time series data and model simulations and compares the results. It can be used to diagnose goodness of fit and/or as the basis for frequency-domain parameter estimation (spect.match).

spect.match tries to match the power spectrum of the model to that of the data. It calls an optimizer to adjust model parameters to minimize the discrepancy between simulated and actual data.

Usage

```
  ## S4 method for signature 'pomp'
spect(object, params, vars, kernel.width, nsim, seed = NULL,
      transform = identity,
      detrend = c("none","mean","linear","quadratic"),
      ...)
  ## S4 method for signature 'spect.pomp'
spect(object, params, vars, kernel.width, nsim, seed = NULL, transform,
      detrend, ...)
  spect.match(object, start, est = character(0),
              vars, nsim, seed = NULL,
              kernel.width, transform = identity,
              detrend = c("none","mean","linear","quadratic"),
              weights, method = c("subplex","Nelder-Mead","SANN"),
              verbose = getOption("verbose"),
              eval.only = FALSE, fail.value = NA, ...)
```

Arguments

object	An object of class pomp.
params	optional named numeric vector of model parameters. By default, params=coef(object).
vars	optional; names of observed variables for which the power spectrum will be computed. This must be a subset of rownames(obs(object)). By default, the spectrum will be computed for all observables.
kernel.width	width parameter for the smoothing kernel used for calculating the estimate of the spectrum.
nsim	number of model simulations to be computed.
seed	optional; if non-NULL, the random number generator will be initialized with this seed for simulations. See [simulate-pomp](#).
transform	function; this transformation will be applied to the observables prior to estimation of the spectrum, and prior to any detrending.
detrend	de-trending operation to perform. Options include no detrending, and subtraction of constant, linear, and quadratic trends from the data. Detrending is applied to each data series and to each model simulation independently.
weights	optional. The mismatch between model and data is measured by a weighted average of mismatch at each frequency. By default, all frequencies are weighted equally. weights can be specified either as a vector (which must have length equal to the number of frequencies) or as a function of frequency. If the latter, weights(freq) must return a nonnegative weight for each frequency.
start	named numeric vector; the initial guess of parameters.
est	character vector; the names of parameters to be estimated.
method	Optimization method. Choices are [subplex](#) and any of the methods used by [optim](#).
verbose	logical; print diagnostic messages?
eval.only	logical; if TRUE, no optimization is attempted. Instead, the probe-mismatch value is simply evaluated at the start parameters.
fail.value	optional scalar; if non-NA, this value is substituted for non-finite values of the objective function.
...	Additional arguments. In the case of spect, these are currently ignored. In the case of spect.match, these are passed to optim or subplex in the control list.

Details

A call to spect results in the estimation of the power spectrum for the (transformed, detrended) data and nsim model simulations. The results of these computations are stored in an object of class [spect.pomp](#).

A call to spect.match results in an attempt to optimize the agreement between model and data spectrum over the parameters named in est. The results, including coefficients of the fitted model and power spectra of fitted model and data, are stored in an object of class [spect.matched.pomp](#).

Value

spect returns an object of class spect.pomp, which is derived from class pomp and therefore has all the slots of that class. In addition, spect.pomp objects have the following slots:

kernel.width width parameter of the smoothing kernel used.

transform transformation function used.

freq numeric vector of the frequencies at which the power spectrum is estimated.

datspec, simspec estimated power spectra for data and simulations, respectively.

pvals one-sided p-values: fraction of the simulated spectra that differ more from the mean simulated spectrum than does the data. The metric used is L^2 distance.

detrend detrending option used.

spect.match returns an object of class spect.matched.pomp, which is derived from class {spect.pomp} and therefore has all the slots of that class. In addition, spect.matched.pomp objects have the following slots:

est, weights, fail.value values of the corresponding arguments in the call to spect.match.

evals number of function and gradient evaluations by the optimizer. See optim.

value Value of the objective function.

convergence, msg Convergence code and message from the optimizer. See optim.

Author(s)

Daniel C. Reuman, Cai GoGwilt, Aaron A. King

References

D.C. Reuman, R.A. Desharnais, R.F. Costantino, O. Ahmad, J.E. Cohen (2006) Power spectra reveal the influence of stochasticity on nonlinear population dynamics. *Proceedings of the National Academy of Sciences* **103**, 18860-18865.

D.C. Reuman, R.F. Costantino, R.A. Desharnais, J.E. Cohen (2008) Color of environmental noise affects the nonlinear dynamics of cycling, stage-structured populations. *Ecology Letters*, **11**, 820-830.

See Also

pomp-class, pomp-methods, probe, probe.match

Examples

```
pompExample(ou2)
good <- spect(
              ou2,
              vars=c("y1","y2"),
              kernel.width=3,
              detrend="mean",
              nsim=500
              )
```

```
summary(good)
plot(good)

ou2.bad <- ou2
coef(ou2.bad,c("x1.0","x2.0","alpha.1","alpha.4")) <- c(0,0,0.1,0.2)
bad <- spect(
              ou2.bad,
              vars=c("y1","y2"),
              kernel.width=3,
              detrend="mean",
              nsim=500
              )
summary(bad)
plot(bad)
```

---

traj.match   *Trajectory matching*

Description

Facilities for matching trajectories to data. Trajectory matching is equivalent to maximum likeli-hood estimation under the assumption that process noise is entirely absent, i.e., that all stochasticity is measurement error.

Usage

```
  ## S4 method for signature 'pomp'
traj.match(object, start, est = character(0),
           method = c("Nelder-Mead","subplex","SANN","BFGS",
                      "sannbox","nloptr"),
           transform = FALSE, ...)
  ## S4 method for signature 'traj.matched.pomp'
traj.match(object, est, transform, ...)
  ## S4 method for signature 'pomp'
traj.match.objfun(object, params, est, transform = FALSE, ...)
```

Arguments

object	A [pomp](#) object. If `object` has no `skeleton` slot, an error will be generated.
start	named numeric vector containing an initial guess for parameters. By default `start=coef(object)` if the latter exists.
params	optional named numeric vector of parameters. This should contain all parameters needed by the `skeleton` and `dmeasure` slots of `object`. In particular, any parameters that are to be treated as fixed should be present here. Parameter values given in `params` for parameters named in `est` will be ignored. By default, `params=coef(object)` if the latter exists.

est   character vector containing the names of parameters to be estimated. In the case
                   of `traj.match.objfun`, the objective function that is constructed will assume
                   that its argument contains the parameters in this order.

method   Optimization method. Choices are [subplex], "sannbox", and any of the methods
                   used by [optim].

transform   logical; if TRUE, optimization is performed on the transformed scale.

...   Extra arguments that will be passed either to the optimizer ([optim], [subplex],
                   [nloptr], or [sannbox], via their `control` (optim, subplex, sannbox) or `opts`
                   (nloptr) lists) or to the ODE integrator. In `traj.match`, extra arguments will be
                   passed to the optimizer. In `traj.match.objfun`, extra arguments are passed to
                   [trajectory]. If extra arguments are needed by both optimizer and [trajectory],
                   construct an objective function first using `traj.match.objfun`, then give this
                   objective function to the optimizer.

Details

In **pomp**, trajectory matching is the term used for maximizing the likelihood of the data under the
assumption that there is no process noise. Specifically, `traj.match` calls an optimizer ([optim],
[subplex], and [sannbox] are the currently supported options) to minimize an objective function. For
any value of the model parameters, this objective function is calculated by

1. computing the deterministic trajectory of the model given the parameters. This is the trajectory
   returned by [trajectory], which relies on the model's deterministic skeleton as specified in the
   construction of the `pomp` object `object`.

2. evaluating the negative log likelihood of the data under the measurement model given the
   deterministic trajectory and the model parameters. This is accomplished via the model's
   `dmeasure` slot. The negative log likelihood is the objective function's value.

The objective function itself — in a form suitable for use with [optim]-like optimizers — is created
by a call to `traj.match.objfun`. Specifically, `traj.match.objfun` will return a function that
takes a single numeric-vector argument that is assumed to cotain the parameters named in `est`, in
that order.

Value

`traj.match` returns an object of class `traj.matched.pomp`. This class inherits from class [pomp]
and contains the following additional slots:

transform, est the values of these arguments on the call to `traj.match`.

evals number of function and gradient evaluations by the optimizer. See [optim].

value value of the objective function. Larger values indicate better fit (i.e., `traj.match` attempts
   to maximize this quantity.

convergence, msg convergence code and message from the optimizer. See [optim].

Available methods for objects of this type include `summary` and `logLik`. The other slots of this
object can be accessed via the `$` operator.

`traj.match.objfun` returns a function suitable for use as an objective function in an [optim]-like
optimizer.

See Also

trajectory, pomp, optim, subplex

Examples

```
pompExample(ou2)
true.p <- c(
    alpha.1=0.9,alpha.2=0,alpha.3=-0.4,alpha.4=0.99,
    sigma.1=2,sigma.2=0.1,sigma.3=2,
    tau=1,
            x1.0=50,x2.0=-50
    )
simdata <- simulate(ou2,nsim=1,params=true.p,seed=43553)
guess.p <- true.p
res <- traj.match(
  simdata,
  start=guess.p,
  est=c('alpha.1','alpha.3','alpha.4','x1.0','x2.0','tau'),
  maxit=2000,
  method="Nelder-Mead",
  reltol=1e-8
  )

summary(res)

plot(range(time(res)),range(c(obs(res),states(res))),type='n',xlab="time",ylab="x,y")
points(y1~time,data=as(res,"data.frame"),col='blue')
points(y2~time,data=as(res,"data.frame"),col='red')
lines(x1~time,data=as(res,"data.frame"),col='blue')
lines(x2~time,data=as(res,"data.frame"),col='red')

pompExample(ricker)
ofun <- traj.match.objfun(ricker,est=c("r","phi"),transform=TRUE)
optim(fn=ofun,par=c(2,0),method="BFGS")

pompExample(bbs)
## some options are passed to the ODE integrator
ofun <- traj.match.objfun(bbs,est=c("beta","gamma"),transform=TRUE,hmax=0.001,rtol=1e-6)
optim(fn=ofun,par=c(0,-1),method="Nelder-Mead",control=list(reltol=1e-10))
```

---

trajectory   *Compute trajectories of the deterministic skeleton.*

Description

The method `trajectory` computes a trajectory of the deterministic skeleton of a Markov process. In the case of a discrete-time system, the deterministic skeleton is a map and a trajectory is obtained by iterating the map. In the case of a continuous-time system, the deterministic skeleton is a vector-field; `trajectory` integrates the vectorfield to obtain a trajectory.

Usage

```
## S4 method for signature 'pomp'
trajectory(object, params, times, t0, as.data.frame = FALSE, ...)
```

Arguments

object	an object of class pomp.
params	a rank-2 array of parameters. Each column of params is a distinct parameter vector.
times, t0	times is a numeric vector specifying the times at which a trajectory is desired. t0 specifies the start time (the time at which the initial conditions hold). The default for times is times=time(object,t0=FALSE) and t0=timezero(object), respectively.
as.data.frame	logical; if TRUE, return the result as a data-frame.
...	additional arguments are passed to the ODE integrator if the skeleton is a vectorfield and ignored if it is a map. See [ode] for a description of the additional arguments accepted.

Details

This function makes repeated calls to the user-supplied skeleton of the pomp object. For specifications on supplying this, see [pomp].

When the skeleton is a vectorfield, trajectory integrates it using [ode].

When the skeleton is a map, trajectory iterates it. By default, time is advanced 1 unit per iteration. The user can change this behavior by specifying the desired timestep using the argument skelmap.delta.t in the construction of the pomp object.

Value

Returns an array of dimensions nvar x nreps x ntimes. If x is the returned matrix, x[i,j,k] is the i-th component of the state vector at time times[k] given parameters params[,j].

Author(s)

Aaron A. King <kingaa at umich dot edu>

See Also

[pomp], [traj.match], [ode]

Examples

```
pompExample(euler.sir)
x <- trajectory(euler.sir)
plot(time(euler.sir),x["I",1,],type='l',xlab='time',ylab='I')
lines(time(euler.sir),x["cases",1,],col='red')

coef(euler.sir,c("gamma")) <- log(12)
```

```
x <- trajectory(euler.sir)
plot(time(euler.sir),x["I",1,],type='l',xlab='time',ylab='I')
lines(time(euler.sir),x["cases",1,],col='red')

x <- trajectory(euler.sir,as.data.frame=TRUE)
```

---

verhulst   *Simple Verhulst-Pearl (logistic) model.*

Description

verhulst is a pomp object encoding a univariate stochastic logistic model with measurement error.

Details

The model is written as an Ito diffusion, $dn = rn(1 - n/K)dt + \sigma n dW$, where W is a Wiener process. It is implemented using the euler.sim plug-in.

See Also

pomp-class and the vignettes

Examples

```
pompExample(verhulst)
plot(verhulst)
coef(verhulst)
params <- cbind(
                c(n.0=100,K=10000,r=0.2,sigma=0.4,tau=0.1),
                c(n.0=1000,K=11000,r=0.1,sigma=0.4,tau=0.1)
                )
x <- simulate(verhulst,params=params,states=TRUE)
matplot(time(verhulst),t(x['n',,]),type='l')
y <- trajectory(verhulst,params=params)
matlines(time(verhulst),t(y['n',,]),type='l',lwd=2)
```

Index

70